

Stat 442 Lecture Notes

Spring Semester, 2025

Last Updated: May 12, 2025

1 Wednesday 2/5: Welcome, what is statistical learning, basics of supervised learning

1.1 What is statistical learning?

Classical statistics is built on the scientific method. It assumes that we start with a hypothesis, and then we go out and we collect data that can be used to test this hypothesis. Importantly, the statistical methods that we will use to test this hypothesis are pre-specified; we choose them before we ever look at our data.

For example, maybe we come up with the hypothesis that college students who sleep more than 7 hours per night have higher average GPAs than those who sleep less than 7 hours per night. We then go out and take a random sample of students, and ask them all if they sleep more than 7 hours per night, and also ask for their GPA. We then use a t-test to test if there is a statistically significant difference in GPA between students who sleep more or less than 7 hours per night.

Statistical learning is different. While there is no single, satisfying definition of statistical learning, it refers generally to a collection of tools used to make sense of complex data. Unlike classical statistics, where a hypothesis is defined before data collection, statistical learning often involves discovering patterns in data without an initial hypothesis.

In the context of our motivating example, suppose that we collect a large survey of students at a college. We are generally interested in what factors impact the GPA of a college student. We search through the data, try different models, and end up realizing that “hours of sleep per night” seems to be an important factor in determining GPA: more important than, say, major, whether or not you are a varsity athlete, etc. However, we also realize that the relationship between hours of sleep per night and GPA does not look linear. We decide to turn “hours of sleep per night” into a binary variable: yes/no do you sleep more than 7 hours per night. We have ended up with the same “model” as before, but this was a much more exploratory process. It started with a general task of *prediction* (as opposed to a specific model or hypothesis), and used *variable selection* and *feature engineering* to settle on a final model. I would consider this to be an example of statistical learning.

In recent decades, as datasets have become larger and computers have become more powerful, there have been a proliferation of complex methods for analyzing data that are often labeled as machine learning methods. In general, however, there is no clear dividing line between statistics and machine learning. Linear regression and logistic regression, which you all studied in detail in Stat 346, are certainly examples of machine learning methods: they help us learn from data. In fact, they are both methods for *supervised learning*, which is one of two primary types of statistical learning.¹

- In *supervised learning*, we start with a response variable of interest and a set of potential predictor variables (also known as explanatory variables). Our general goal is to use the predictor variables to explain or predict the response variable.
- In *unsupervised learning*, we just start with a collection of variables: there is no specific response variable. Our goal is to find structure or patterns in the data.

In this course, we will spend around 9 weeks on supervised learning, and only 2 weeks at the end on unsupervised learning. However, I do not want to give you the impression that unsupervised learning is less important! Unsupervised learning is extremely important, and is in fact the back-bone of recent generative AI technology. We start with supervised learning because it has a more natural connection to classical statistics and builds more naturally on your previous work. Hopefully, some of the concepts that you learn during our supervised learning unit will help you be equipped to delve more deeply into unsupervised learning in the future if you choose.

There is one more type of learning that has become really important in recent years, partly due to the proliferation of generative AI. This is the field of *semi-supervised learning*, in which we have a (potentially small) set of *labeled data* (data with predictors + response) and a (potentially much larger) set of *unlabeled data* (data that is missing the response). While we will not cover this in lecture, it would make a great *final project topic*. Throughout the semester, I will try to flag these topics as they come up!

¹I would be remiss if I did not mention *reinforcement learning*, which is a very important topic in machine learning but does not relate as directly to drawing insights from a given dataset.

1.2 Supervised learning

1.2.1 Introduction

As mentioned, you all already know at least two methods for supervised learning: linear regression and logistic regression. In general, these represent one from each of two classes of supervised learning methods.

- Regression: we use this term for any method that is used to predict a quantitative response variable.
- Classification: we use this term for any method that is used to predict a categorical response variable.

As a reminder, a quantitative variable can be either continuous or discrete, but it must be a number. A categorical variable can be either ordinal (the categories have orders) or unordered. In this class, we will not cover any methods for considering ordinal response variables: if this topic interests you, it is another *potential final project topic*.

Linear regression and logistic regression are basic examples of supervised learning methods. In the past few decades, much more complex algorithms have been developed, and have exploded in popularity. There are a few reasons for this proliferation in complex methods.

- Datasets have gotten bigger. As we will learn in this class, complex models are often only appropriate when applied to enough data. Thus, it is natural that the “big data” era has gone hand-in-hand with the development of complex statistical learning methods.
- Computers have gotten better, faster, and cheaper, which has enabled the fitting of complex models.
- Free, open-source software programs like R and python have allowed researchers to develop algorithms and then easily share them with non-experts. This has allowed many methods to become popular.

Given the huge number of statistical learning methods, we could spend each class this semester learning a new method, such that at the end you are left with a huge cookbook of methods to choose from for your future data analysis needs. In fact, this class will be a little bit like this. But, there is a problem with this approach of treating statistical learning like a cookbook. New machine learning methods get published every day, and the state-of-the-art is constantly changing. If your goal is to learn a list of methods or algorithms, you will find yourself continually behind and out-of-date! Thus, the goal of this class is not really to teach you a list of methods. Instead, the goal is to teach you the fundamental concepts and overarching themes that go into the development of the methods, such that you can compare methods, recognize their limitations, and learn new methods on your own in the future. This goal helps explain two features of this course:

- We will spend a fair amount of time on older, less state-of-the-art methods. We will use these methods to illustrate general principles and themes, even if they may not be the methods that you will use in practice in your future.
- You will all be doing a *teaching presentation* this semester. You will be responsible for doing independent reading to learn about a method or a concept. You will need to synthesize what you learn, and give a 20 minute presentation to the class explaining the method. This skill mirrors the way you will all interact with machine learning in the future: you will be continually learning and synthesizing new methods.

1.2.2 Notation for datasets and random variables

The backbone of statistical learning is a matrix of predictors $\mathbf{X} \in \mathbb{R}^{n \times p}$ and a vector of length n that stores the responses \mathbf{y} . We let n denote the number of *observations* (rows of our dataset) and let p be the number of *variables* (columns of our dataset, usually not including the response): this notation is quite standard in statistics, but isn't always particularly precise.

Your textbook (ISL) denotes pieces of the dataset \mathbf{X} and \mathbf{y} as follows. We observe x_i for $i = 1, \dots, n$, where $x_i = (x_{i1}, \dots, x_{ip})^\top$ is a p -vector for each observation. We use $\mathbf{x}_1, \dots, \mathbf{x}_p$ to denote the columns of \mathbf{X} ; each is an n -vector representing its own variable. Your textbook says that it will use bold lowercase letters anytime the vector has length n and use unbold lowercase letters otherwise: I am sure that I will start messing this up soon, but I will try to be consistent.

The other important distinction is between a random variable and its observed realization. We will use capital, unbold letters to denote random variables. In our example from Section 1.1, we might let X_1 denote the number of hours that a randomly selected student sleeps per night and let Y denote their GPA. Once we observe a particular n students for our dataset, we let their sleep values be $\mathbf{x}_1 = (x_{11}, x_{12}, \dots, x_{1n})^\top$ and we let their GPA values be $\mathbf{y} = (y_1, \dots, y_n)^\top$.

Consider a dataset that contains information on n pets. The response variable is the weight of the pet \mathbf{y} , and we have one single predictor variable: the type of pet. This is a categorical variable that takes on values {cat, dog, hamster, rabbit}. I could represent \mathbf{X} in this case as a matrix with one column; where the column stores the actual values {dog, dog,

cat, dog, hamster, cat, dog, ...}. But, as you all know from Stat 346, it is likely going to be convenient for fitting to instead let \mathbf{X} be a matrix with four columns, that store binary variables indicating “is this a dog”, “is this a cat”, etc. This simple example shows us how notation can get complicated: do we say that $p = 4$ or $p = 1$ for this dataset? Practically speaking, the $p = 4$ is likely more relevant! But it can depend!

We can also get into problems when counting n sometimes. Suppose we take measurements on 920 students, who are spread out between 8 schools. The schools are randomly assigned to either have required yoga PE class, or to have traditional PE class. While we have 920 students, the students within a given school are not *independent* of one another: there may be factors other than the yoga class that are making their test scores more similar to one another. On the other hand, the schools underwent random assignment and can be safely assumed to be independent of one another. Thus, is the sample size n the 920 students or the 8 schools? It turns out that for this *clustered* or *hierarchical* data, we have a notion of *effective sample size* that is somewhere between 920 and 8, and depends on how correlated the students are within the schools.

We should also mention that data does not always come to you in a form where it looks like $\mathbf{X} \in \mathbb{R}^{n \times p}$ and $\mathbf{y} \in \mathbb{R}^n$. Consider the task of image classification. You get a set of 500 images, which are stored as 400×400 pixel images, where each pixel records a numerical value between 0 and 255 for red, green, and blue. Each image is associated with a label in $\{\text{cat}, \text{dog}\}$, and you observe y_1, \dots, y_{500} . In this case, the RGB values in the pixels in the images are your predictor values, but they don't come in a simple $n \times p$ matrix. We need to reshape the data so that each image i has an associated vector $x_i \in \mathbb{R}^{400 \times 400 \times 3}$ that stores all three color values for all 400×400 pixels. It will also likely be convenient to code y_i such that a 1 denotes a dog and a 0 denotes a cat, such that y becomes numerical. When there are more than 2 classes, such as $\{\text{cat}, \text{dog}, \text{fish}\}$, we will need more than one column to store our response variables numerically.

Who knew that counting n and p could be so complicated! In this class, just think of it as the number of rows and columns (minus the response columns) of the data, and remember these subtleties in case they ever come up.

In our supervised learning unit, we will start with *regression*. Thus, for the rest of today, we assume that Y is numerical.

1.3 Regression

In a regression problem, we assume that our response variable Y is related to a set of predictors $X = (X_1, \dots, X_p)$ via some model that can be written as

$$Y = f(X) + \epsilon. \quad (1)$$

The function $f()$ is a fixed but unknown function that represents the systematic information that X provides about Y , and ϵ is a random error term, which should satisfy the condition that $E[\epsilon] = 0$ and that ϵ is independent of X . These conditions ensure that $f()$ is capturing all parts of Y that can be explained by X , and ϵ captures the parts of Y that cannot be explained by X . The task in regression is to come up with a good approximation for the unknown function $f()$. More specifically, we want to find a function $\hat{f}()$ such that, $Y \approx \hat{f}(X)$.

We talked really briefly about two methods that you might use to find a good \hat{f} , that sort of represent opposite ends of the spectrum in terms of how “wiggly” they are. We talked about:

- Linear regression, which you have all seen before.
- KNN, which many of you have not seen before. We will go into a lot more detail on KNN next class.

1.3.1 Two possible goals

There are two reasons why we might want to come up with a function $\hat{f}()$ such that, $Y \approx \hat{f}(X)$.

1. We might want to be able to come up with predictions $\hat{Y} = \hat{f}(X)$ for future realizations of data where Y itself is not observed.
2. We might want to understand which predictors in X are most associated with Y in order to draw a scientific conclusion, or at least take a step towards drawing a scientific conclusion.

When beginning a regression problem, it is important to clearly define which of these two reasons is more important to you. This will aid you in choosing the best strategy for approaching the problem!

As we will see throughout the semester, these two goals can sometimes be at odds with one another! If we only care about prediction, we do not need our function $\hat{f}()$ to be interpretable: we can use an extremely complex model (known as a “black box”), and we will be happy as long as $\hat{Y} \approx Y$. On the other hand, if we want to draw scientific

conclusions, then we might wish to use a simple, interpretable model for $\hat{f}()$ that lends itself to rigorous inference. We should also note that it is not always an either-or situation. Sometimes, we want to make good predictions and also understand which variables are contributing significantly to the predictions. Managing the tradeoff, or lack thereof, between these two goals will be one of our fundamental themes for the semester.

We will pick up with everything else next class!

2 Monday, Feb 10: The bias variance tradeoff, as illustrated with KNN and linear regression

Like last class, we will focus on a really simple regression setting. We observe a single response $\mathbf{y} \in \mathbb{R}^n$ and a single predictor $\mathbf{x} \in \mathbb{R}^n$. We assume that our observations are i.i.d. realizations of random variables (X, Y) , where

$$Y = f(X) + \epsilon.$$

We assume that $E[\epsilon] = 0$ and $\epsilon \perp\!\!\!\perp X$, but the function $f()$ is unknown. Our goal is to find a function \hat{f} such that $Y \approx \hat{f}(X)$. In this unit, we will talk about different algorithms for finding such a function \hat{f} !

We will assume that we are looking for \hat{f} that makes the squared error loss:

$$\left(Y - \hat{f}(X)\right)^2$$

small on average. Today we are going to talk a lot about this goal. And we are going to talk about how well KNN and linear regression each achieve this goal.

For today, we are using squared error loss everywhere. But please remember that this is not the only way to measure model accuracy! The exact math of the bias-variance decomposition that we will discuss today holds only for squared error loss, but similar concepts apply for other loss functions.

Here is the agenda for today:

1. Training error vs. expected test error.
2. The bias-variance decomposition of the expected test error.
3. What is the ideal function for minimizing the expected test error?
4. How do KNN and Linear Regression each approximate this ideal function?
5. R Demo of KNN and Linear regression and the bias variance tradeoff.

We might move a little bit fast, but HW1 is going to give you a chance to practice all of these concepts.

2.1 Training error vs. test error

If we use our dataset $\mathbf{y} \in \mathbb{R}^n$ and $\mathbf{x} \in \mathbb{R}^n$ to *train* a function \hat{f} , then our *training error* is:

$$MSE_{\text{train}} = \frac{1}{n} \sum_{i=1}^n \left(y_i - \hat{f}(x_i)\right)^2.$$

For a fixed function \hat{f} that we already trained, if we observe some test points $x_i^{\text{test}}, y_i^{\text{test}}$ for $i = 1, \dots, n_{\text{test}}$, we can compute our test error as

$$MSE_{\text{test}} = \frac{1}{n_{\text{test}}} \sum_{i=1}^{n_{\text{test}}} \left(y_i^{\text{test}} - \hat{f}(x_i^{\text{test}})\right)^2,$$

where the important thing is that the points $x_i^{\text{test}}, y_i^{\text{test}}$ were not used to compute the function \hat{f} .

While in practice we usually do compute MSE_{test} over some fixed test set that has size n_{test} , we don't want our notion of model performance to be specific to the test set that we happened to see. We are likely studying MSE_{test} to estimate how big our error will be on a random new datapoint. In this case, we might be computing this for the particular function \hat{f} that we already computed. In this case, our expected prediction error conditional on our particular function \hat{f} that we already computed is:

$$E_{X^{\text{test}}, Y^{\text{test}}} \left[\left(Y^{\text{test}} - \hat{f}(X^{\text{test}})\right)^2 \right],$$

where the function \hat{f} is not treated as random.

If we want to evaluate the potential of an algorithm or a procedure to make good predictions, then we want to take into account the randomness in our training set. We want to acknowledge that we could have seen a different training set that would have given us a different function \hat{f} . With this in mind, in the next section we will define a more complex notion of expected prediction error. This more complex notion will let us come up with a very beautiful decomposition!

2.2 The bias-variance tradeoff

Suppose that we train our model \hat{f} on a dataset $\mathbf{X}^{\text{train}}, \mathbf{y}^{\text{train}}$. Since this dataset is random, the function $\hat{f}()$ is also random. Now, suppose that we would like to know about the expected prediction error for a new datapoint with $X = x^{\text{test}}$ and unknown $Y = y^{\text{test}}$. The prediction error itself is $(y^{\text{test}} - \hat{f}(x^{\text{test}}))^2$, but to find the *expected* prediction error we need to take the expected value over multiple sources of randomness.

For our purposes, let's treat x^{test} as fixed. We want to find the average prediction error that we would obtain if we repeatedly (1) sampled training sets of size n from the population, (2) refit the function \hat{f} using this training data, and (3) evaluated the squared prediction error for a datapoint drawn with $X = x^{\text{test}}$. So, we want to evaluate:

$$E_{\mathbf{X}^{\text{train}}, \mathbf{y}^{\text{train}}, Y | X = x^{\text{test}}} \left[(Y - \hat{f}(X))^2 \mid X = x^{\text{test}} \right] = E_{\mathbf{X}^{\text{train}}, \mathbf{y}^{\text{train}}, Y | X = x^{\text{test}}} \left[(Y - \hat{f}(x^{\text{test}}))^2 \mid X = x^{\text{test}} \right].$$

For brevity, I am not going to keep writing the subscripts in the expected values. I am going to assume that we are treating x^{test} as fixed but taking the expected value over all other randomness. But remember that it is always good to know what exactly you are taking the expected value over!

To break this quantity down, we are going to do a clever trick that involves adding and subtracting 0 twice. This is a very common proof technique! After we do this, we expand our big squared quantity by thinking of it as $(a + b + c)^2$, and we apply linearity of expectation to put each term in its own expected value.

$$\begin{aligned} E \left[(Y - \hat{f}(x^{\text{test}}))^2 \right] &= E \left[(Y - \textcolor{red}{f}(x^{\text{test}}) + \textcolor{red}{f}(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})] + E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}}))^2 \right] \\ &= E \left[(Y - f(x^{\text{test}}))^2 \right] + E \left[(f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})])^2 \right] + E \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}}))^2 \right] \\ &\quad + 2E \left[(Y - f(x^{\text{test}})) (f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) \right] \\ &\quad + 2E \left[(Y - f(x^{\text{test}})) (E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] \\ &\quad + 2E \left[(f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) (E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right]. \end{aligned}$$

We will now argue that each of the cross terms is 0. We start with the orange term, and note that:

$$\begin{aligned} E \left[(Y - f(x^{\text{test}})) (f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) \right] &= (f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) E[Y - f(x^{\text{test}})] \\ &= (f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) (E[Y] - f(x^{\text{test}})) \\ &= (f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) (f(x^{\text{test}}) - f(x^{\text{test}})) = 0. \end{aligned}$$

The first two equalities follow because the quantities $f(x^{\text{test}})$ and $E[\hat{f}(x^{\text{test}})]$ are not random, since x^{test} is a constant. The last equality follows since, because we are conditioning on $X = x^{\text{test}}$, $Y = f(x^{\text{test}}) + \epsilon$, and so by our assumption that $E[\epsilon] = 0$, $E[Y] = f(x^{\text{test}})$.

We now consider the purple term. We note that $Y - f(x^{\text{test}})$ is independent of the training dataset, since $f()$ is non-random and Y is a new realization. On the other hand, $(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}}))$ is a function only of the training data, since x^{test} is non-random. Thus, this term has the form $E[ab]$ where a and b are independent. So we can write

it as:

$$\begin{aligned} E \left[(Y - f(x^{\text{test}})) (E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] &= E[Y - f(x^{\text{test}})] E \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] \\ &= E[Y - f(x^{\text{test}})] (E[\hat{f}(x^{\text{test}})] - E[\hat{f}(x^{\text{test}})]) = 0. \end{aligned}$$

Finally, we consider the green term.

$$\begin{aligned} E \left[(f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) (E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] &= (f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) E \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] \\ &= (f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})]) (E[\hat{f}(x^{\text{test}})] - E[\hat{f}(x^{\text{test}})]) = 0. \end{aligned}$$

Now that we know that all three cross terms are 0, we know that

$$E \left[(Y - \hat{f}(x^{\text{test}}))^2 \right] = E \left[(Y - f(x^{\text{test}}))^2 \right] + E \left[(f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})])^2 \right] + E \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}}))^2 \right].$$

Our next goal is to understand each of these terms!

1. The first term can be written as:

$$E \left[(Y - f(x^{\text{test}}))^2 \right] = E \left[(f(x^{\text{test}}) + \epsilon - f(x^{\text{test}}))^2 \right] = E[\epsilon^2] = \text{Var}(\epsilon).$$

This is our irreducible error term! It comes from the inherent noise in our data that cannot be explained by X .

2. In the second term, $f(x^{\text{test}})$ and $E[\hat{f}(x^{\text{test}})]$ are not random, so we do not need the expected value! This just becomes

$$(f(x^{\text{test}}) - E[\hat{f}(x^{\text{test}})])^2,$$

which we call the squared bias. Is $\hat{f}(x^{\text{test}})$ equal to $f(x^{\text{test}})$ on average? If so, then \hat{f} is unbiased and this term will disappear.

3. Finally, in the third term:

$$E \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}}))^2 \right] = \text{Var} \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] + E \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}}))^2 \right].$$

First, note that $\text{Var} \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] = \text{Var} \left[\hat{f}(x^{\text{test}}) \right]$, because $E[\hat{f}(x^{\text{test}})]$ is not random. Then, note that $E \left[(E[\hat{f}(x^{\text{test}})] - \hat{f}(x^{\text{test}})) \right] = E[\hat{f}(x^{\text{test}})] - E[\hat{f}(x^{\text{test}})] = 0$. So, this third term is simply the variance (with respect to the training set) of $\hat{f}(x^{\text{test}})$.

This was a lot of work, but we can finally say that, at a given point x^{test} ,

$$E \left[(Y - \hat{f}(x^{\text{test}}))^2 \right] = \text{Var}(\epsilon) + \text{Bias}(\hat{f}(x^{\text{test}}))^2 + \text{Var}(\hat{f}(x^{\text{test}})).$$

This is the bias variance decomposition, which will be important throughout the semester!

The $\text{Var}(\epsilon)$ term is our irreducible error. No matter how good we are at estimating f , this is just the amount of noise in our data, so we will always have this error. On the other hand, we can reduce the bias and the variance terms if we pick a better statistical learning algorithm for estimating f .

Without giving too much away: very simple models tend to have high bias but low variance. Very complex (wiggy) models have low bias (they are never constrained!) but very high variance (they overfit to the training data). To minimize our expected prediction error, we are always looking for functions that hit the sweet spot of complexity! This sweet-spot is usually at the minimum of a U-shaped curve for test-set error! You will make U-shaped curves on your homework this week!

2.3 The ideal function \hat{f}

Let's briefly forget the training data! If our goal was to minimize

$$E_{X,Y} \left[\left(Y - \hat{f}(X) \right)^2 \right]$$

and we had all of the resources in the world, what would we choose for \hat{f} ?

Well, under the law of total expectation, we can write this as:

$$E_X \left[E_{Y|X} \left[(Y - \hat{f}(X))^2 \mid X \right] \right].$$

In the inner expected value, X is no longer random, and solving for the point-wise minimum is easy. We have that:

$$\operatorname{argmin}_c E_{Y|X=x} [(Y - c)^2 \mid X = x] = E[Y \mid X = x].$$

This comes from the fact that an expected squared error is always minimized at a mean. You will prove this on HW1, and you have likely seen it before.

If $E[Y \mid X = x]$ minimizes the point-wise expected prediction error at every x , then the function $\hat{f}(x) = E[Y \mid X = x]$ is the function that minimizes the overall expected prediction error. Of course, we do not know the joint distribution of X and Y , and so we cannot simply set $\hat{f}(x)$ to be this conditional expectation. But we can develop methods that attempt to approximate $\hat{f}(x) = E[Y \mid X = x]$ under various sets of assumptions! As we will see today, both KNN and linear regression try to approximate $E[Y \mid X]$.

2.4 How close do KNN and Linear Regression get to this ideal function?

Let's review a really simple case. We have a single numerical response variable Y and a single numerical predictor variable X . We observe 100 datapoints, so $n = 100$ and $p = 1$. Our goal is to use our dataset $\mathbf{x} = (x_1, \dots, x_n)$, $\mathbf{y} = (y_1, \dots, y_n)$ to come up with a function \hat{f} such that, on new, unseen data, $Y \approx \hat{f}(X)$. Today, we will consider two methods for coming up with \hat{f} .

2.4.1 Simple linear regression

We restrict our attention to \hat{f} that have the form $b_0 + b_1x$ for $b_0, b_1 \in \mathbb{R}$. This makes our problem of finding the best \hat{f} easier, because we limited the complexity of the model class that we are considering.

Based on our training data, we let:

$$\hat{\beta}_0, \hat{\beta}_1 = \operatorname{argmin}_{b_0, b_1} \sum_{i=1}^n (y_i - b_0 - b_1 x_i)^2. \quad (2)$$

We then let $\hat{f}(X) = \hat{\beta}_0 + \hat{\beta}_1 X$. Under the assumption that $E[Y] = \beta_0 + \beta_1 X$ for some true, unknown β_0 and β_1 , this solution directly approximates the best possible function $E[Y \mid X]$.

However, if this assumption does *not* hold, then the linear model is too limiting. The result of limiting our model class so much is that we might have a lot of bias. If our linearity assumption does not hold, then no matter how much data we have $E[\hat{f}(x)]$ just cannot be that close to f . This is why simple models can have a lot of bias— we did not give them enough complexity to be able to capture the true function!

You are all experts in linear regression! So we will not spend too much more space in the notes on it.

2.4.2 K-nearest neighbors

KNN is at the opposite end of the spectrum from linear regression. The premise is quite simple. KNN lets

$$\hat{f}(x) = \frac{1}{k} \sum_{x_i \in N_k(x)} y_i, \quad (3)$$

where $N_k(x)$ is a function that returns the k points in the training set that are closest to the input point x according to some distance metric (for us, Euclidean distance). So, avoiding equations: KNN makes predictions by finding the k training observations with x_i closest to x , and predicts that the response for x will be the average of the responses for those k points.

The hyper-parameter k is chosen by the user. When $k = n$, KNN just predicts $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$ for all observations, regardless of x . As k decreases, the functions returned by KNN get increasingly wiggly and flexible. When $k = 1$, the KNN function is a step function that can be arbitrarily wiggly, and that always has 0 training error.

Note that KNN directly approximates $E[Y | X = x]$ *locally*, making no assumptions on the structure of $E[Y | X]$. When k is small, the estimation of $E[Y | X = x]$ uses only data that are really close to x . This lets our prediction function get arbitrarily wiggly— we will not run into an issue of bias. However, with small k , there is a lot of variance in each individual prediction, since it is made using very few datapoints. When k is large, there is less variance in the predictions, but we do not let our predictions be as wiggly, and so we introduce more bias. Thus, KNN is a great place to see the bias-variance tradeoff at work.

A few notes on KNN:

- KNN is non-parametric; we cannot write down the function \hat{f} without storing basically the entire dataset. The number of effective parameters grows with the sample size n .
- Linear regression in theory takes some time to train, but it is nearly instantaneous to make new predictions. KNN involves no training step, but it could be computationally expensive to obtain new predictions.

2.5 Time for an R demo!

The RMarkdown document that we go over in class will be posted on GLOW!

3 Thursday, Feb 13: adding more predictor variables!

Recall our typical regression setting. We assume that our data are i.i.d. realizations of random variables (X, Y) , where

$$Y = f(X) + \epsilon.$$

We assume that $E[\epsilon] = 0$ and $\epsilon \perp\!\!\!\perp X$, but the function $f()$ is unknown. Our goal is to find a function \hat{f} such that $Y \approx \hat{f}(X)$.

More specifically, we would like to find \hat{f} that minimizes the expected squared error loss for a new, unseen datapoint (X, Y) . So ideally, we are looking for

$$\operatorname{argmin}_{\hat{f} \in \mathcal{F}} E_{X,Y} \left[\left(Y - \hat{f}(X) \right)^2 \right], \quad (4)$$

where \mathcal{F} is some class of functions. On your homework, you will argue that if \mathcal{F} were totally unconstrained, we would want to set $\hat{f}(x) = E[Y | X = x]$. This argument was also in the Lecture 2 notes, but we skipped it.

This is where we hit practical issues. We do not know the joint distribution of X and Y , and so we cannot pick \hat{f} to be this ideal function $E[Y | X = x]$. And we really cannot search efficiently over all possible real-valued functions \mathcal{F} to find a good choice for \hat{f} . So, a statistical learning algorithm typically restricts the class \mathcal{F} to make this task doable. So far in this class, we have discussed two possible methods for picking \hat{f} . It turns out that both of these can be viewed as approximating $E[Y | X = x]$; they just do this using different sets of assumptions.

Up until now, we have been assuming that we just have one predictor variable, and so X is a scalar. Today, we will let X be a vector, meaning that we have p predictor variables X_1, \dots, X_p . This is going to introduce a lot more nuance to the comparison between linear regression and KNN!

Agenda :

1. Linear regression in high dimensions.
2. KNN in high dimensions.
3. R demo, and overall comparison of linear regression vs. KNN, without preprocessing.

3.1 Linear regression in high dimensions

We restrict our model class \mathcal{F} to

$$\mathcal{F} = \{f : f(x) = b_0 + b^T x, b_0 \in \mathbb{R}, b \in \mathbb{R}^p\},$$

such that (4) becomes

$$\operatorname{argmin}_{b_0 \in \mathbb{R}, b \in \mathbb{R}^p} E_{X,Y} \left[\left(Y - b_0 - b^T X \right)^2 \right]. \quad (5)$$

As you know from Stat 346, it is convenient to append a column of 1s to our predictor matrix \mathbf{X} so that we are just optimizing over a single $b \in \mathbb{R}^{p+1}$. That way, we don't need to keep writing the intercept separately. Adopting this convention, with linear regression we approximate (5) by letting

$$\hat{\beta} = \arg \min_{b \in \mathbb{R}^{p+1}} \sum_{i=1}^n (y_i - b^T x_i)^2. \quad (6)$$

You know 1,000 things about this estimator from Stat 346. For example, as long as $n > p$, the solution to (6) has a closed-form: $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T y$. Furthermore, under the assumption that $E[Y | X] = \beta^T X$ for some true unknown β (i.e. under the assumption that the true model is linear), this estimator is BLUE (best unbiased linear estimator), where “best” here means lowest variance.² This is the *Gauss-Markov Theorem*, and I am assuming that you saw it in Stat 346.

Here are a few pros and cons of linear regression to keep in mind. We will discuss these a lot, and also revisit them in our R demo.

- **Pro: efficiency:** If $n > p$, we have a closed-form solution. This means that the function is efficient to train. It is also really efficient to generate new predictions.
- **Con: identifiability:** If $n < p$, we cannot solve for $\hat{\beta}$ because there is no unique solution; the model is not identifiable.
- **Con: bias:** If the true model is not linear, we will have bias! Because the linear model might simply not be flexible enough to model the true relationship between the predictors and the response.
- **Pro/Con: variance:** Because it is not super flexible, linear regression should be a low-variance method. Unfortunately, when we start adding a lot of irrelevant or redundant predictors to the model, the variance can get very high.
- **Pro: interpretability:** Linear regression is interpretable! We can look at the estimated function \hat{f} and say what variables are important, and in what direction they are contributing to our predictions.
- **Pro: usability:** Linear regression is easy to use “out-of-the-box” for a non-expert. There isn't anything to *tune* if we just regress y on all the variables (which we can do as long as $p < n$). Any refinement that the user wants to do after that is pretty easy to explain/understand.

3.2 KNN in high dimensions

We know that the best solution to (4) is $E[Y | X = x]$. So, we assume only that $E[Y | X = x]$ is smooth enough that:

$$E[Y | X = x] \approx E[Y | X \text{ is in a neighborhood of } x].$$

We then approximate this function directly by letting:

$$\hat{f}(x) = \frac{1}{k} \sum_{i \in N_k(x)} y_i,$$

where $N_k(x)$ is the k training set datapoints that are closest to the desired test point x .

We didn't emphasize this point when we only had a single predictor variable, but the notion of “closest” datapoints requires a distance metric! Our distances are now measured in p -dimensional predictor space, and we actually have many different distance metrics that we could use. Unless otherwise specified, assume that we are using Euclidean distance. This means that:

$$N_1(x^{\text{test}}) = \arg \min_{i=1, \dots, n} \sqrt{\sum_{j=1}^p (x_j^{\text{test}} - x_{ij})^2} = \arg \min_{i=1, \dots, n} \|x^{\text{test}} - x_i\|_2.$$

Importantly, if your different predictors x are measured with different units, this neighbor function might not even make sense! This neighbor function treats all features the same. If one of our features is “price of house” and another feature is “number of bedrooms in house”, one of these has values in the hundreds of thousands and the other has values that are likely below 10. In this setting, the euclidean distance between points will be totally dominated by price, and bedrooms will be basically ignored. Because of this, you should almost certainly scale your features before applying KNN: usually we would do this by dividing every feature by its standard deviation. This creates a distance

²Do you remember what it means to call this a *linear* estimator? A hint is that it would still be linear even if we used polynomial features like X^2 .

function where all features contribute equally. Below, we will talk about how this can lead to its own issues.

There are some pros and cons of KNN that are relevant regardless of the dimension p .

- **Con: usability:** we need to choose k . This means that KNN cannot be used directly “out of the box” - we probably want to choose k using cross-validation.
- **Pro/Con: bias/variance tradeoff:**
 - With small k , we can approximate functions that are arbitrarily wiggly. So we don’t need to worry that our function class \mathcal{F} is not sufficiently flexible.
 - When k is small, each prediction is generated with very few datapoints, which means that we have high variance.
 - When k is large, a prediction might be generated with points that are actually far away from the test point x^{test} , which can introduce bias.
 - If we have enough data points n such that our predictor space is densely populated with training points, we can pick a pretty large k and still have it be the case that every training point in $N_k(x)$ is actually close to x . This means that we won’t have much bias, but we will also have low variance since k is large.
 - For a given problem, we may or may not find a “sweet-spot” k that makes KNN work really well!
- **Pro / Con: efficiency:**
 - Computational cost of training is essentially free. Just store the training dataset.
 - At testing, we need to compute distances between the new test point and all points in the training set. This could be slow if you implement it naively, but people have cool algorithms for finding nearest neighbors efficiently.

There are some additional cons of KNN that come up when the dimension p is large.

- **Impact of irrelevant features:** Suppose that we have p features in our dataset, but only a small subset of these features actually impact the response Y . All of these features are used in computing the neighbors of x ! So, we are letting totally irrelevant features contribute to our distance metric. This could introduce either bias or variance. The bias would be because I am doing a bad job selecting the meaningful neighbors. The variance would be because my prediction can vary based on extra random noise in some random irrelevant feature.
- **Lack of interpretability:** Unlike linear regression, running KNN on a high dimensional dataset tells you nothing at all about what features are most associated with Y . If you wanted to remove irrelevant features to improve performance, KNN provides no built in guidance for doing this. This is in contrast to linear regression.
- **Computational cost:** Storing the training set is actually quite expensive when n and p are large. So even though “training the model” just involves “storing the training set” - this is not free! Neither is computing lots of pairwise distances in high dimensions, or searching through high dimensional space for neighbors! We will need to come up with more clever algorithms for KNN to get around this.
- **Curse of dimensionality:** It turns out that, in high dimensions, points just become far from one another! So selecting neighbors in high dimensions is just a really hard task! The entire notion of neighbors hardly makes sense anymore, because our high-dimensional feature space will not be densely populated unless n is really large compared to p . So even though we can theoretically do KNN when $p > n$ (unlike for linear regression), it is going to work very poorly in this setting.

3.3 Comparing linear regression and KNN, without preprocessing

See the R Demo from class.

We ran out of time before we could talk about what to do about this! We will talk about that next time!

3.4 Miscellaneous

People asked good questions in class, so here are some extra points that I wanted to add!

- Another big advantage of linear regression over KNN is that we have theory that lets us do inference (e.g. prediction intervals!) to rigorously quantify our uncertainty for linear regression.
 - This inference depends on some number of assumptions about the underlying data generating mechanism. For example, the typical inference that we do assumes homoskedasticity (equal variance of our errors), but robust/sandwich standard errors let us relax this assumption (at the possible cost of statistical power).
 - We like that KNN requires no assumptions whatsoever: we don’t need to believe that our model is truly linear, etc. But a downside of this lack of assumptions is that we don’t have many tools available to make

prediction intervals! This is a typical tradeoff: we would love to avoid assumptions, but usually avoiding assumptions also means we can't say much about our final answer!

- Off the top of my head, I don't know much about inference for KNN. But I included some random thoughts below!
- For KNN with $K = 3$, each prediction is based on three datapoints! Without additional assumptions, there is going to be a LOT of uncertainty associated with this prediction.
- Later in the semester, we will talk about conformal inference, which lets us make prediction intervals for any machine learning model. This could be applied to KNN!
- To make a prediction interval, we need both an estimate but also a measure of sampling uncertainty associated with this estimate. Maybe we could use bootstrapping to estimate the SE of a KNN prediction and make an interval based on that? When I googled "inference for KNN", I found a bunch of suggestions about bootstrapping! It would be interesting to try this out and see if it works!

* A sketch of the idea: Get B different KNN predictions for a point x^{test} by using different random bootstrap subsamples of the training dataset each time. Then, compute the standard deviation of these B predictions. Then, get your main prediction \hat{y}^{test} using all of the data. Let $\hat{y}^{\text{test}} \pm 2\hat{SD}(\hat{y}^{\text{test}})$ be your prediction interval. Would this cover the true y^{test} 95% of the time? I am not sure because KNN is strange and bootstrapping does have theory/requirements. Let's try it out in a future R demo during lecture.

- We mentioned that KNN requires storing the entire dataset to make test predictions, and that it is inefficient. This is not strictly speaking true, because people have made smart algorithms for getting around this problem. If you google "KNN with KD tree or Ball tree" you will find a lot of resources on fast algorithms. These are the ones that I learned about in my ML class, but I think a lot of others exist too!
- Someone asked if you need to store the entire training dataset to fit a linear regression model. It turns out that you don't! I will put a reading on your HW about "online learning" for linear regression. The idea is that, instead of storing our training dataset X, y all at once and computing $\hat{\beta} = (X^T X)^{-1} X^T y$, we can instead write $\hat{\beta}$ as:

$$\hat{\beta} = \arg \min_b \sum_{i=1}^n (y_i - x_i^T b)^2$$

and assume that we only have access to a single example x_i, y_i at a time. We have algorithms that can step along the gradient of this loss function according to one example x_i, y_i at a time in order to find the minimum. So the full data never needs to be stored!

- I'm sure there are things I forgot. Keep asking good questions!

4 Monday, 2/17: Feature selection, feature engineering, and regularization!

We ended last class by talking about how both KNN and Linear Regression can do badly when the number of variables p is very large! There are two distinct problems.

- For linear regression with no preprocessing, increasing p just increases the model complexity of a linear regression model. This allows us to overfit (or memorize) our training set, which leads to very high variance and poor test error.
- For KNN, p doesn't really have anything to do with model complexity, but we do poorly because KNN works best when we have training points that are dense in our feature space (so that the k nearest neighbors to a test point x are actually close to x). As p grows, it is really really hard for points to be dense (n would need to be HUGE), and so we are making predictions with neighbors that aren't really that close (bias) and who our neighbors actually are is affected by noise in so many different predictors (variance).

Real data is high dimensional! So we really need a way around these issues! Today, we will talk about two strategies for avoiding issues with high dimensions.

- **Strategy 1:** Preprocess the data to reduce the dimension before we apply the algorithm. Within strategy 1, we have two sub-strategies.
 - **Feature selection:** You have all used this for linear regression; it's a really natural fit. It definitely *could* be used for KNN, but you would need to pick some sort of selection algorithm that might not have anything to do with KNN.
 - **Feature extraction:** Can do this before KNN, linear regression, or any other algorithm.
- **Strategy 2:** Modify the algorithm using **regularization** to automatically reduce the variance.

- We are going to talk about this in the context of linear regression, not KNN.
- The concept (Regularization!!) will be relevant to any algorithm we use this semester that gets fit by minimizing a loss function. KNN is one of the only algorithms that we will use this semester that doesn't have this loss-function form, and so regularization is not directly applicable.
- It will turn out that Lasso regularization ends up looking a lot like variable selection! But they arise from slightly different concepts!

Let's talk about each of these! I am going to zoom through material today. Chapter 6 of ISL has a really nice treatment of feature selection, feature extraction, and regularization for linear regression. Please read this on your own!

4.1 Feature selection

The idea is that if we have p predictors $\{X_j : j \in \{1, \dots, p\}\}$ but we don't think that all of them are relevant, we should select a subset $\mathcal{S} \in \{1, \dots, p\}$. We should then do our statistical learning algorithm using only $\{X_j : j \in \mathcal{S}\}$. If we do a good job with selection, we should improve our performance, because our variance will be smaller (for linear regression) and we will have less curse of dimensionality (for KNN).

4.1.1 Guess and check

You have all done variable selection for linear regression via what I might call “guess and check” before. You have all fit a model with a bunch of variables, and then looked at the p-values and decided only to include variables that seem significant! You have also checked for things like multicollinearity, and removed variables that you are worried are redundant. All of these strategies are great, and show a big advantage of linear regression: we have interpretable tools for built-in variable selection.

4.1.2 Best subset regression

A more automated way to do variable selection would be to try every single subset $\mathcal{S} \in \{1, \dots, p\}$ and choose the subset that leads to the lowest test set MSE, the lowest AIC or BIC, or is the best according to some other metric.

If p is small, we could just try out fitting least squares models to all of the different subsets $\mathcal{S} \subset \{1, \dots, p\}$. We could then directly compare our metrics for every possible subset, and if we pick the subset that leads to the lowest value we have our solution! This is best-subset regression, and I think you should have heard about it in Stat 346! Best-subset regression is kind of silly, because it is computationally infeasible when p is big, which is exactly the situation in which we need it!

4.1.3 Forward stepwise regression

Since best subset regression is computationally infeasible, one idea is to try to approximate the solution to best subset regression with a more reasonable search strategy. With best subset regression, we would need to try out 2^p different models. With the forward stepwise algorithm suggested below, we need to try out at most $p + (p - 1) + (p - 2) + \dots$ models.

The idea of forward stepwise regression is simple:

- Start with an empty model that only includes an intercept
- Until a stopping criteria is met:
 - Look through all possible predictors that are not yet in the model. Add the predictor that most improves the model at this moment!

We get to decide what we mean by “most improves the model”. We could, at each step, add the most significant predictor, or the one that most improves R^2 , etc. We could use a stopping criteria such as: “until the BIC stops improving” or “until no variable that could be added has a p-value less than 0.05”. If we do this, then the size of our final model is determined for us, using only the training set.

We could also use no stopping criteria, and just go until we run out of predictors, or until the number of predictors is equal to the number of datapoints. If we do this, we get a sequence of nested models, where the variables were added in a greedy order. We could select our final model size by seeing which of these nested models minimizes the test set error, for example. You did this on your homework. This idea of getting a whole sequence of models, and then selecting the one that minimizes test set error, will look a bit like regularization!

Note that backwards stepwise regression is also a thing that you may have learned about in Stat 346. I think it is unsatisfying when we are discussing high dimensional regression, since it cannot be used for $p > n$ (since you cannot fit the initial model to step backwards from).

I think you all know a bit about feature selection for linear regression from Stat 346! So I will not talk too much more about it.

4.1.4 Feature selection for KNN??

KNN does not lend itself to a built-in way to do feature selection, as far as I know. We could “try out fitting KNN with different features included or removed”, and compare test MSE for different options. This is a lot like best-subset selection; it is computationally infeasible to try out all possible options. We could also run something like stepwise linear regression as a preprocessing step to KNN. This would be a strange thing to do, but the idea would be that we think we need wiggly functions (hence KNN), but we first want to have some efficient way to select variables that seem associated with Y . People also use a preprocessing method called marginal screening, which just computes the marginal correlation between Y and X_j for $j = 1, \dots, p$ and only keeps variables j that seem correlated with Y . You could try this out with KNN! A fear with any of these methods is that removing some X s can alter who is neighbors with who in your data: which might be good but it also might be bad!

4.1.5 Overview:

In general, what does feature selection get us, and what are the risks?

- **Interpretability:** selecting a small number of features makes our final model more interpretable.
- **Usability:** A user needs to tune a hyper-parameter that controls “how many variables to keep.” This makes everything a bit more complicated, but at least the tuning parameter is interpretable, and automated procedures exist.
- **Bias:** If we fail to include an important variable, we could introduce bias. In other words, we could accidentally select a model that is too simple.
- **Variance:** The whole point of feature selection is to reduce variance! It definitely does this.
- **Computational efficiency:** Some methods like best subset selection and stepwise regression could be slow. So the actual selection step can be slow. But in general, if we select only some features from our data and use these for the rest of our analyses, we have less data to store and we will make downstream tasks faster.
- **Inference:** If you were hoping that, after variable selection, you could get nice p-values from `lm()` for each of your selected variables, you are wrong! This is the problem of *post-selection inference*, which is my research area! We cannot use the same data more model selection and model inference. If we want to do inference after model selection, we need to refit the selected model on a totally held-out test set. Ask me about this sometime, or do this topic for your final project! There are papers that do fancy math to do inference after stepwise regression!

4.2 Feature extraction

The idea is that if we have p predictors $\{X_j : j \in \{1, \dots, p\}\}$ but we think that a lot of them are redundant, maybe we can compress the information from the p features into a smaller number of features $\{\tilde{X}_k = f_k(\mathbf{X}) : k \in \{1, \dots, S\}\}$. Each of the S new features can contain information from all of the old features.

Consider image classification. Our high-dimensional set of predictors is every pixel in an image. Feature selection will work terribly here: we can’t just select some of the pixels for every image and expect to do a good job. But there are probably low-dimensional concepts hidden in the images that can capture all of the information that we need, without storing every single pixel. Thus, image classification is a setting where feature extraction is really helpful but where feature selection makes no sense.

Here are a few examples:

- The new features \tilde{X}_k for $k \in \{1, \dots, S\}$ could be the first S principal components of the original feature matrix \mathbf{X} . We still need to choose S : this is now a tuning parameter.
- It turns out that you can randomly project your original p -dimensional feature matrix \mathbf{X} into a lower dimensional subspace to get $\tilde{\mathbf{X}}$. Magical theorems say that the distances between observations are approximately preserved under random projections, and so this can work pretty well as a preprocessing step to KNN.
- PCA and random projections are linear embeddings. You could use something like UMAP or tSNE as a non-linear embedding. These are popular in genomics! They put your points onto a low-dimensional manifold.
- Autoencoders are a really cool way to learn a low dimensional representation of a high dimensional object!

PCA is a really classic example, and I think that you are all familiar with PCA from Stat 346. I included a little sidebar on PCA below, just in case we have time to go over it or just in case you are curious. But how PCA actually works is not really our topic for today.

The general idea of any of these is that we can avoid the curse of dimensionality if we can capture most of the information about our \mathbf{X} variables in fewer dimensions. This is going to work best when we have a lot of redundant predictors.

4.2.1 PCA

In the simplest case, if we assume that our feature matrix X has been centered and scaled so that the mean of each variable is 0 and the standard deviation of each variable is 1, and we also assume that $n > p$, then we can take the singular value decomposition of X , and write it as

$$X = UDV^T,$$

where $U \in \mathbb{R}^{n \times p}$ is a matrix whose columns are orthogonal unit vectors, $D \in \mathbb{R}^{p \times p}$ is a diagonal matrix, and $V \in \mathbb{R}^{p \times p}$ is an orthonormal matrix.

The columns of V define a new set of axes in our predictor space. V_1 represents the direction that contains as much of the variance in X as possible. V_2 represents the direction orthogonal to V_1 that explains as much of the leftover variance as possible, and so on. In PCA-speak, these are called the loadings. They correspond to the eigenvectors of the correlation matrix of the data X , and are ordered in such a way that V_1 corresponds to the biggest eigenvalue, V_2 to the second biggest eigenvalue, etc.

If the columns of V are the axes, then the columns of UD store the position of each datapoint along these axes. We call these the scores.

Let U_r denote the first k columns of u , let D_r denote the first r rows and columns of D , and let V_r denote the first r columns of V . Then, the matrix

$$U_r D_r V_r^T,$$

is the “best” (according to mean-squared-error or L2-norm) rank- r approximation to X .

What this means is that if we let $\tilde{X} = U_r D_r \in \mathbb{R}^{n \times r}$, we have made ourselves r new variables that store “as much of the variation in X as possible” (from a MSE perspective, and among linear embeddings).

The new variables are also independent of one another, so there is no multicollinearity left. We can use this \tilde{X} in our downstream task. We probably only want to do this if it stores a large proportion of the total variance in X : we can plot this proportion of variance vs. r to help decide how many principal components to keep!

4.2.2 Overview

In general, what does feature extraction get us, and what are the risks?

- **Interpretability:** Often, feature extraction can make a final model less interpretable. Occasionally, you can get lucky, and identify your top few principal components as recognizable concepts.
- **Usability:** You need to figure out how to extract features, how many PCs to keep, etc. This adds a tuning task that I think is less straightforward than the selection task.
- **Bias:** If you don’t retain enough good information about \mathbf{X} , you could introduce bias.
- **Variance:** The point is to reduce variance!
- **Computational efficiency:** Story is the same as for feature selection. Something like UMAP is computationally hard to run, but it saves you time and space down the road because you don’t need to store your entire high dimensional dataset anymore.

One more nice thing is that something like PCA has nothing to do with linear regression. We can use PCA before KNN and it works really well: you will do this on your homework! So the techniques we use for feature extraction tend to be very general; even if Chapter 6 of your ISL textbook talks about PCA in the context of linear regression only.

4.3 Regularization

The idea of regularization is to reduce our variance by adding a penalty on model complexity directly into our loss function!

4.3.1 Best subset regression through L0-regularization!

We can revisit the idea of best subset regression, but case it as a regularization problem. We can do linear regression, but we can let

$$\hat{\beta} = \operatorname{argmin}_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - x_i^T b)^2 + \lambda \|b\|_0, \quad (7)$$

where

$$\|b\|_0 = \{\#i : b_i \neq 0\}.$$

This is just saying that we want to fit a least squares regression, but we want to prefer a *sparse* solution.

The cool thing about writing the objective function in this way is that the parameter λ directly controls the bias-variance tradeoff. If we set λ to be very large, we will select a model with very few non-zero coefficients. This model will have a lot of bias (even if the *true* model is linear! because we might have accidentally left out important variables). However, the model will have low variance because it is so simple. On the other hand, as λ approaches 0, we approach our un-regularized high-dimensional regression, whose variance we know is very high.

This is a nice idea! But there is one huge problem! Can we actually solve (7)?

Unfortunately, the answer is no! Not efficiently! The only way to solve (7) exactly is to try out all 2^p possible models, which is infeasible. While this first example of regularization is not actually something that we do in practice, the idea turns out to be really useful!

4.3.2 Ridge regression through L2-regularization!

One reason that we cannot solve (7) is that our favorite way to find minimums is to take a derivative, and $\|b\|_0$ is not differentiable. So, what if we just modified (7) to make it differentiable? Ridge regression solves the following optimization problem:

$$\hat{\beta} = \operatorname{argmin}_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - x_i^T b)^2 + \lambda \|b\|_2^2, \quad (8)$$

where

$$\|b\|_2^2 = \sum_{i=1}^p b_i^2.$$

Unlike (7), we can solve (12) in closed-form. You will do this on HW2! But ... what does it get us?

If, in truth, $y_i = x_i^T \beta + \epsilon$, then the solution to (12) is biased for β . The amount of bias will grow as λ grows. However, we can prove that, as λ grows, the variance of the solution to (12) always goes down. So, while the least squares solution ($\lambda = 0$) is known to be the Best Linear Unbiased Estimator, it turns out that some solutions to (12) (i.e. solutions for some other values of λ) could have lower expected prediction error than the least squares solution—even though they are biased! This is really cool- but how do we find these λ s?

We have not officially talked about K-fold cross-validation yet in this class. While it is a really general technique that has nothing to do specifically with ridge regression, I am going to talk about it right now! Because it is the way we will choose a value of λ for ridge regression.

Here is how we would use K-fold cross-validation in this context!

- Divide our n datapoints into K non-overlapping folds of data.
- For k in $1, \dots, K$:
 - Let the k th fold be the test set. Let the other $K - 1$ folds be the training set.
 - For λ' in a big grid of possible values of λ :
 - * Solve (12) with $\lambda = \lambda'$, using the training set.
 - * Compute the prediction error on the test set. Save this error for this fold and this value λ' .
- For each λ' , add up the test set errors across all K folds.
- Select the λ^* that minimizes the total test set error. (Or use the 1-SE rule to select a slightly different λ^* . We will come back to this at some point on a HW.)
- Refit (12), using all n datapoints and λ^* . Let this be your final model.

Ok, so we know that we can solve (12). And we also know that we can use cross validation to pick a λ^* that hopefully hits a sweet spot of bias and variance. There are a few practical notes to talk about.

- It is good to center and scale your X variables before you apply ridge. The reason you should center them is that we want our model to spiritually include an intercept, but we don't want to put the penalty λ on the intercept. It is good to scale them because we don't want a variable with different units to get unfairly penalized by the penalty term- we want all of the coordinates of b_i to be on the same scale.
- The solution to (12) is not sparse, so is not necessarily interpretable. Sometimes, people look at the output of (12) and look at which variables got their coefficients shrunk very small. They then decide to remove those, and refit a final model using least squares to only the remaining variables. This makes the final model more interpretable. We should note that standard “lm()” based inference is invalid after doing this (it's the same issue as doing inference after variable selection).
- The amount by which the ridge solution shrinks each element of $\hat{\beta}$ actually depends on the singular vectors of X ! Which is cool! More redundancy in the columns of X will lead to more shrinkage. This makes sense, because we also know that more redundancy in the columns of X means more variance in a least squares solution, so more shrinkage is needed!
- There are a lot of cool things we could study for ridge. It's mathematically beautiful! For example, there is a way to view ridge regression as a continuous analog of regression in principal components! That is so cool. A few cool things about ridge will be on your HW or in the textbook. For now, we will move on!

4.3.3 Lasso regression through L1 regularization

It is sort of a bummer that ridge regression does not give us sparse solutions. It does reduce the variance of the fitted model, but just by looking at the fitted model it is hard to tell that we have reduced its complexity. That is where lasso comes in. Magically, it turns out that the solution to

$$\hat{\beta} = \operatorname{argmin}_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - x_i^T b)^2 + \lambda \|b\|_1, \quad (9)$$

where

$$\|b\|_1 = \sum_{i=1}^p |b_i|,$$

tends to be sparse! We call this Lasso regression!

A lot of things are the same as with ridge. The parameter λ governs a bias-variance tradeoff. You should be sure to center and scale your variables before you use Lasso. You should use cross validation to select the best value of λ .

By sparse, we mean that, for large enough λ , all of the coefficients will be shrunk towards 0, but several will be exactly 0. This means that the Lasso actually performs variable selection! This is such a simple change compared to (12). Why would the solution suddenly be sparse? This is a fun optimization fact that is often explained using Figure 1.

Of course, an issue is that $\|b\|_1$ is no longer differentiable. In fact, this is a non-convex optimization problem. So, while Ridge had a closed-form solution, Lasso does not. We need to solve (13) with an iterative algorithm called coordinate descent.

In a setting where all variables are actually related to the response variable, Ridge regression tends to have slightly better test set performance than Lasso (at their respective optimal values of λ , which do not need to be the same). However, in a setting where the true solution is sparse, Lasso can outperform ridge. So, in practice, both are useful tools, and you might want to compare them for a given problem using cross validation. Personally, I love the Lasso because I love the sparse property for interpretability. But some problems are not sparse! (Remember our image example!).

Note that, if you are using Lasso for variable selection, you might want to handle categorical dummy variables in a special way. You might want to make sure that groups of β s representing categories of the same variable are either selected or not selected together. You can do this with the group-lasso.

4.3.4 Overview

In general, what does regularization get us, and what are the risks?

- We will work with so many algorithms this semester that can be written as: “minimize this loss function over a training set”. Any algorithm that has this form has the risk of overfitting to the training set! But adding regularization on model complexity can always help us avoid overfitting! This is beautiful and general!
- **Usability:** We need to do cross validation to choose λ . Luckily, for things like Ridge and Lasso, this is really easy and is just built into the software.




Figure 1: A figure taken from ISL that illustrates why the Lasso prefers sparse solutions while Ridge does not. If we are optimizing the least squares objective (oval) subject to the constraint that either the L1 or L2 norm of β is not too large, the our solution will lie in the place where the contours of our least squares objective first hit our constraint set. For the L2 constraint (circle), this can happen anywhere. For the L1 solution (diamond), this is more likely to happen on one of our 4 pointy corners. For more intuition on why, you should see the .Rmd document that I posted on GLOW, which will let you draw this picture for yourself for a particular dataset.

- **Bias:** Regularization introduces bias.
- **Variance:** Regularization reduces variance. That is the whole point!
- **Computational efficiency vs. interpretability:** From an interpretability standpoint, we most would have wanted to do L0 regularization. This encourages sparse solutions, but doesn't actually shrink the values of coefficients that are in our model. However, L0 regularization is computationally infeasible. On the other hand, L2 regularization is really fast and efficient, but the solution is not sparse and so isn't as interpretable. L1 regularization is a nice happy medium! It does not have a closed form solution but we can still solve it, and we get a sparse solution (with coefficients that also got shrunk).

4.4 Wrap up

This was a whirlwind intro to regularization, ridge, and lasso, which are REALLY important topics. We will do some R demos next class. We will also discuss Figure 1 next class: we really have not yet done it justice.

In general, even if Ridge and Lasso feel rushed in this class, I hope that you will remember them forever! They are really fundamental concepts! We have ridge/lasso logistic regression too- it's not only a linear regression thing! It's a really beautiful example of the bias-variance tradeoff and the interplay between optimization and statistics!

5 Thursday, 2/20: More on regularization

We did not end up getting to classification at all! Because there is a LOT to say about Ridge and Lasso, so we spent a lot of time on them!

5.1 Recap from last time on Regularization

If we have n training datapoints $(x_1, y_1), \dots, (x_n, y_n)$, then ordinary least squares solves the following objective function on the training set:

$$\hat{\beta}_{OLS} = \arg \min_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - b^T x_i)^2 \quad (10)$$

When our true model is linear, this is unbiased for the “true β ”. But, it can have high variance, especially if the number of predictors p is large and many of those predictors are irrelevant.

Last class, we mentioned three regularization methods. The point of all three is to reduce the variance of (10), but all three do this at the expense of possibly introducing bias. The three methods were:

- **Best subset:**

$$\hat{\beta}_{subset,\lambda} = \arg \min_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - b^T x_i)^2 + \lambda \|b\|_0 \quad (11)$$

- **Ridge:**

$$\hat{\beta}_{ridge,\lambda} = \arg \min_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - b^T x_i)^2 + \lambda \|b\|_2^2 \quad (12)$$

- **Lasso:**

$$\hat{\beta}_{lasso,\lambda} = \arg \min_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - b^T x_i)^2 + \lambda \|b\|_1 \quad (13)$$

If we had to summarize an overall comparison of pros and cons very quickly, I think we would want to say:

- Best subset regression is interpretable! And, if we select the “true” subset of important variables, we have unbiased coefficients for the “true” non-zero elements of β ! Unfortunately, it is really computationally expensive!
- Ridge regression has a closed form solution so it is really efficient! But it doesn’t lead to an interpretable, low-dimensional model. It just reduces variance.
- Lasso regression is interpretable and can be solved more efficiently than best subset! We do still shrink the non-zero coefficients, so even if we select the “true” subset of important variables, we will have bias in the coefficients.

We visualized the bias and variance of Ridge and Lasso as a function of λ in an R demo that is posted on GLOW. We compared to linear model and stepwise regression. We went through the document kind of fast, but there is a lot of good stuff in the document: please read and make sure you understand!

There are a few more things that we need to discuss.

- What the heck is going on in Figure 1? More specifically, what is the connection between regularization and constrained optimization?
- Can we say precisely whether best subset or ridge or lasso is more or less wiggly?

These are covered below.

5.2 Connection between regularization and constrained optimization

To understand Figure 1, it is important to know (14), (12), and (13) can all be written as constrained optimization problems. This is discussed on Page 243 of ISL.

Let’s focus on Lasso for simplicity. It turns out that (13) is equivalent to:

$$\hat{\beta}_{lasso,\lambda} = \arg \min_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - b^T x_i)^2 \quad \text{subject to: } \|b\|_1 < s.$$

For some s , which depends on λ and also the values of x_1, \dots, x_n and y_1, \dots, y_n .

This tells us that our original idea of penalized regression is equivalent to a constrained optimization problem: the problem of solving the least squares problem, but subject to the constraint that $\|b\|_1 < s$ for some budget s . This formulation is useful because there are smart people who know a lot about constrained optimization!

In class, people asked really good questions about WHY we can write a loss+penalty problem as a constrained optimization problem, and also about the connection between s and λ . This has to do with the theory of the Lagrangian and Lagrange multipliers for constrained optimization! I put a document on GLOW with a lot of notes about this. Please take a look!

How does this idea relate to Figure 1?

- Figure 1 draws the contours of the least squares objective function for a given dataset. The minimum of the least squares objective function is the OLS solution, which is marked as $\hat{\beta}$ in the figure. The contours are drawn as a function of b_1 and b_2 , which we are optimizing over.
- Since we can see ridge and lasso as solving this objective subject to a budget constraint, the diamond and the circle draw the Lasso and Ridge budget constraints for a particular s .

- The theory of Lagrange multipliers and constrained optimization tells us that our constrained solution occurs where the contours of the OLS function are tangent to the constraint set. This I think you should be able to understand.
- It turns out that, because the lasso constraint is "pointy", solutions are more likely to occur on the axis: i.e. at sparse solutions where one element of β is exactly 0. This I do not think we really have quite enough theory to understand: we would need to study convex optimization and the KKT conditions and such. But I hope the picture gives you a little bit of intuition!

The .Rmd document that I posted on GLOW gives you code to actually make Figure 1 for any dataset that you generate, and plot the Ridge and Lasso solutions. I hope that you will play around with this code until you understand what is being plotted, and believe me that Lasso tends to give sparse solutions!

5.3 Other

Some random thoughts and notes.

5.3.1 Modern comparison!

This paper was published in 2020, which is crazy recent for something about such fundamental topics: <https://www.jstor.org/stable/pdf/26997932.pdf>.

The big idea is that recent computational advances have actually made best-subset regression more doable! So ... should we just do best-subset? Is it the gold-standard over lasso? Apparently not! Best subset and lasso can each outperform each-other in certain settings, and the overall winner is actually the "relaxed lasso". I recommend you read the paper!

5.3.2 Using the lasso for variable selection in practice, and inference after lasso

I was not planning to cover this in class, but then people asked a lot of questions related to it!

Suppose that you are using lasso for variable selection.

You can run lasso (with CV to choose λ on your data). This will return a p -dimensional coefficient vector with a lot of 0s.

You could just start interpreting the non-zero regression coefficients in the way you would interpret multiple regression coefficients. But the slightly strange thing is that, not only did lasso set a bunch of coefficients to 0, but it also shrunk the non-zero coefficients! While it turns out that this shrinking is mathematically good for the bias/variance tradeoff, it isn't good for interpretation!

So, a common thing to do is to re-fit your model using `lm()` (unpenalized regression) after running lasso, but using only the variables that lasso selected.

This is nice! Under the assumption that lasso selected the "true" set of non-zero variables, these coefficients are unbiased for the "true" coefficients and you can interpret them and do inference. This is a very strong assumption. Among variables that lasso selected that turn out to not truly be important, if you trust the `lm()` p-values you will get a hugely inflated Type 1 error rate. This is what I mean when I say that "inference after lasso is invalid".

This is my passion and my research area so if you want to know more about this come ask me! The stepwise regression problem on HW3 also gets at this. But mainly I did not spend much class time on this because inference isn't necessarily the focus of this class.

Also, it's good to know about this "re-fit the model without a penalty" idea, because it seems that they use this technique in the "modern comparison" paper above when they do the "relaxed lasso", which they find works really well!

5.3.3 ElasticNet

We have a method that combines the strength and weaknesses of Ridge and Lasso.

$$\hat{\beta}_{\text{elasticnet}, \lambda, \alpha} = \arg \min_{b \in \mathbb{R}^p} \sum_{i=1}^n (y_i - b^T x_i)^2 + \lambda \left(\alpha \|b\|_1 + \frac{(1-\alpha)}{2} \|b\|_2^2 \right) \quad (14)$$

You could go explore and see if you think this actually works better than Lasso or Ridge individually! And explore its properties! You can also see the shape of the constraint function in ESL Page 73. This section of ESL also talks more about the connection between regularization and Bayesian regression with different priors. How cool and neat!!

6 Monday, 2/24: Introduction to classification

If we rush through this material in class, I recommend that you go revisit Chapter 2 of ISL. It introduces the basics of classification at the same time as the basics of regression. After reviewing this, we will move onto Chapter 4, which discusses methods for classification.

6.1 What is classification

We are switching gears! We assume that we still have a setup with n observations, p predictors in X , and one response variable y . But now y is a categorical variable that can take on one of G possible values³

We still believe that there is some *true* data generating process that generates y from X , with some noise. However, it may not make sense anymore to write $y = f(X) + \epsilon$. Because, if y is a category, what form does the noise ϵ take on? Instead, we assume that y can take on values g_1, \dots, g_G . And that there is some true but known distribution for $Pr(Y = g_k | X)$ that we want to model.

In the special case that Y is binary (there are only two classes), then it must be the case that:

$$Y | X \sim \text{Bernoulli}(f(X)),$$

and so we just need to model $f(X)$. This is an easier case, so we will focus on this a lot. In the case where we have multiple categories, we just replace Bernoulli with Multinomial: we are still going to model probabilities of belonging to certain classes.

But in either case, our goal is to make predictions \hat{Y} using the predictors X ! And we want our predictions to work well for unseen data. And we do not know the true function $f()$, so we need to estimate it.

It is going to turn out that a lot of things are the same as what we discussed in the regression case! For example, the following aspects will all still be important when we evaluate models:

- Bias, which will improve when models are more complex.
- Variance, which will get worse when models are more complex.
- Interpretability, which often gets worse when models are more complex,
- Computational efficiency
- How easy would it be for a non-expert to use it?

6.2 What is our new goal, and what is the ideal model?

Instead of squared error loss, we will use 0/1 loss. We let:

$$L(Y, \hat{f}(X)) = \begin{cases} 0 & \text{if } Y = \hat{Y}(X) \\ 1 & \text{if } Y \neq \hat{Y}(X) \end{cases}.$$

We want to minimize the expected value of this 0/1 loss over possible new datapoints that we might see. Using the law of total expectation twice, if the possible values of Y are $\{g_1, \dots, g_G\}$, this means minimizing:

$$\begin{aligned} E_{X,Y} [L(Y, \hat{Y}(X))] &= E_X [E_{Y|X} [L(Y, \hat{Y}(X)) | X]] \\ &= E_X \left[\sum_{k=1}^G L(Y, \hat{Y}(X)) Pr(Y = g_k | X) \right] \\ &= E_X \left[\sum_{k=1}^G 1\{\hat{Y}(X) \neq g_k\} Pr(Y = g_k | X) \right] \end{aligned}$$

Think about the sum over the G categories. This sum will always have one term that is 0, and the rest of the $G - 1$ terms will have value $Pr(Y = g_k | X)$. To minimize this quantity, we should always let $\hat{Y}(X)$ be the category g_k that maximizes $Pr(Y = g_k | X)$. That way, we are zero-ing out the largest of the G terms $Pr(Y = g_k | X)$ in our sum, and are thus making the sum as small as possible.

So, for regression, the very best possible predictor \hat{Y} if we knew the data generating mechanism was to let $\hat{Y}(x) = E[Y | X = x]$. For classification, the very best possible prediction \hat{Y} if we knew the data generating mechanism is to

³A lot of people use K for number of classes. I think this is sort of horrible, because we also have KNN and K -fold cross-validation. We need a new letter!

let

$$\hat{Y}(x) = \arg \max_{k \in 1, \dots, G} Pr(Y = g_k | X = x).$$

This is called the **Bayes classifier**. The error rate of this classifier is called the **Bayes error rate**. The Bayes error rate is like the irreducible error that we had before for linear regression, where the irreducible error was $Var(\epsilon)$ and it related to variation in Y that could not be explained by X . The Bayes error rate is the same: no matter how good our statistical learning model is, we can never beat the Bayes error rate, and we can never beat the model that knows $Pr(Y = g_k | X = x)$ and always predicts the Y that maximizes this probability.

Instead, we will think about methods that try to approximate the Bayes classifier using various techniques! Today we will cover three: KNN, logistic regression, and LDA.

6.3 KNN classification

For a test point x^{test} , we find the k nearest neighbors in the training set. We let our prediction for x^{test} be the majority class of these neighbors. The majority class among the K -nearest neighbors is a direct estimate of $\arg \max_{k \in 1, \dots, G} Pr(Y = g_k | X = x^{\text{test}})$, if we assume that $Pr(Y = g_k | X = x^{\text{test}})$ is well-approximated by $Pr(Y = g_k | X \text{ is a neighbor of } x^{\text{test}})$.

This makes very few assumptions on the function form of $Pr(Y = g_k | X = x)$: it only assumes that neighbors are similar. But we will have a curse of dimensionality problem again. And a bias/variance tradeoff as we change k . The story is basically the same as it was for regression!

6.4 Logistic regression

Suppose for now that there are only two classes! So Y can take on values 0 or 1. The notion of the Bayes classifier tells us that we should try to model $Pr(Y = 1 | X = x)$. Since we are doing binary classification, this is all we need: we will predict $\hat{Y} = 1$ if our estimated probability is greater than 0.5, and we will predict $\hat{Y} = 0$ otherwise.

If we would like to fit a parametric model for $Pr(Y = 1 | X = x)$, a simple first idea would be to assume that

$$Pr(Y = 1 | X = x) = \beta^T X.$$

We could estimate this with linear regression! But the issue is that we will end up getting predicted probabilities outside of 0 and 1.

The solution, which you have all seen before, is to use a *link function*. The link function will let us use the basic idea of a linear model, but constrain our predictions to be between 0 and 1 so that they can be interpreted as probabilities. This is logistic regression: confusingly named, because it is a classification method, but it also has deep ties to linear regression and other generalized linear models. We assume that:

$$Pr(Y = 1 | X = x) = \frac{e^{\beta^T X}}{1 + e^{\beta^T X}}.$$

To solve this, we find the vector $\hat{\beta}$ that minimizes the negative log likelihood of the training data.

- This model ends up being pretty interpretable (recall from Stat 346 that we can interpret our coefficients on a log-odds scale).
- Furthermore, if our model assumption holds (the log-odds truly are linear in the X s), then theorems from Stat 360 tell us that the maximum likelihood estimator is asymptotically unbiased and achieves the asymptotically smallest possible variance among all unbiased estimators. This is just like our BLUE fact for linear models. Basically, logistic regression should perform well if our assumptions are met.
- However, if we have a ton of predictors p , then our variance can be quite high, and we might actually do better with a biased estimate. Since this optimization problem fits the general framework of looking for a $\hat{\beta}$ that minimizes a training dataset *loss function*, we can reduce our variance with by adding a ridge or lasso penalty to our loss function! And everything will work like it did for linear models! We still use `cv.glmnet`! We can make our model less wiggly with regularization.
- It is also still true that we could have bias if our log odds are not truly linear. So we could add polynomial terms or something to make our model more wiggly.

6.5 LDA: discussion lead by Alessa!

The Bayes classifier tells us that we should model $Pr(Y = g_k | X)$. Logistic regression tried to model this directly. But, what if we first apply Bayes rule to note that:

$$Pr(Y = g_k | X) \propto Pr(X | Y = g_k)pr(Y = g_k).$$

This is an interesting idea! If we imagine that our predictors X are *generated* from one distribution when $Y = 1$ and another distribution when $Y = 0$, then using this formulation to model the distribution of X separately for the two classes kind of makes sense!

LDA, or linear discriminant analysis, decides that we should model $Pr(X | Y = g_k)pr(Y = g_k)$ directly. It does this under the assumption that

$$Pr(X | Y = g_k) \sim N_p(\mu_k, \Sigma),$$

where the mean vectors μ_k for each class and the common covariance matrix Σ are unknown and must be estimated from the training data. LDA estimates these parameters, and then assigns $\hat{Y}(X)$ to be the class that maximizes $\widehat{Pr}(X | Y = g_k)\widehat{Pr}(Y = g_k)$.

Why would a method like this be called “linear” discriminant analysis?

Well, under the assumption of multivariate normality, the *decision boundary* between predicting $\hat{Y}(X) = g_1$ and $\hat{Y}(X) = g_2$ is the set of points where $\widehat{Pr}(X | Y = g_1)\widehat{Pr}(Y = g_1) = \widehat{Pr}(X | Y = g_2)\widehat{Pr}(Y = g_2)$. Let’s solve for this decision boundary under this normality assumption.

Let $\hat{\mu}_1$ be our estimated mean for class g_1 . Let $\hat{\mu}_2$ be our estimated mean for class g_2 . Let $\hat{\Sigma}$ be our shared estimated covariance. Let $\hat{\pi}_1$ and $\hat{\pi}_2$ denote the estimated probabilities of falling in class 1 or class 2. These do not depend on X .

Let’s solve for x such that:

$$(2\pi)^{-k/2}|\hat{\Sigma}|^{-1/2} \exp\left(-\frac{1}{2}(x - \hat{\mu}_1)^T \hat{\Sigma}^{-1}(x - \hat{\mu}_1)\right) \hat{\pi}_1 = (2\pi)^{-k/2}|\hat{\Sigma}|^{-1/2} \exp\left(-\frac{1}{2}(x - \hat{\mu}_2)^T \hat{\Sigma}^{-1}(x - \hat{\mu}_2)\right) \hat{\pi}_2$$

This means solving for x such that

$$-\frac{1}{2}(x - \hat{\mu}_1)^T \hat{\Sigma}^{-1}(x - \hat{\mu}_1) + \log(\hat{\pi}_1) = -\frac{1}{2}(x - \hat{\mu}_2)^T \hat{\Sigma}^{-1}(x - \hat{\mu}_2) + \log(\hat{\pi}_2).$$

Simplifying further,

$$-x^T \hat{\Sigma}^{-1}(\hat{\mu}_1 + \hat{\mu}_2) + \frac{1}{2}\hat{\mu}_1^T \hat{\Sigma}^{-1}\hat{\mu}_1 + \log\left(\frac{\hat{\pi}_1}{\hat{\pi}_2}\right) = \frac{1}{2}\hat{\mu}_2^T \hat{\Sigma}^{-1}\hat{\mu}_2.$$

Note that this is a linear function of the vector x : the quadratic terms cancelled out. This tells us that the boundary between the regions where we predict class 1 and the region where we predict class 2 is a straight line or a plane in X -space. So, LDA tries to draw straight lines or planes in X -space to separate classes. It turns out that this will work well if the classes are truly linearly separable, and will not work well otherwise.

I am sure that Alessa will tell us more about the pros and cons of LDA and how it compares to other methods. These are just initial notes. I bet you can already start making some guesses about bias, variance, interpretability, etc.

On your own time, convince yourself that if you drop the assumption that Σ is shared across classes, you find that the decision boundary is quadratic! This is QDA instead of LDA!

6.6 Asymmetric Loss and ROC Curves

The premise of today’s class was that we want to minimize the expected 0/1 loss. But you all know that, in real life, sometimes false positives are more severe than false negatives, or vis versa.

In such a setting, we might not want to pick the classifier that minimizes the expected 0/1 loss. We might want to pick the classifier that minimizes the false positive rate while maintaining a reasonable false negative rate, or something like that!

This leads us to the notion of an ROC curve, which you will see on your homework 3!

7 Thursday, Feb 27: making models more complex

7.1 Where are we?

We have been talking about ways to predict Y from predictors X_1, \dots, X_p . We have now covered cases where Y is numerical and cases where Y is categorical: and it turns out that many of the common themes hold in both cases.

In basically all cases, we are trying to predict Y using X . We think there is some true function $f()$ that would help us make predictions. But, since we don't know $f()$, we will use our training data to come up with an approximation $\hat{f}(X)$. We usually pick the $\hat{f}(X)$ that does the "best job" on our training data (i.e. $\hat{f}(X) \approx y$), but then we use a test set or cross-validation to make sure that we are not *overfitting*.

It is up to us to decide how complex we want to allow $\hat{f}(X)$ to be. Do we want to restrict ourselves to something simple like a linear model? Or do we want to allow $\hat{f}(X)$ to be very wiggly? This choice always involves a tradeoff between bias and variance.

Last week, we talked about a setting where p is large, and where even a simple linear model has variance that is too high. Thus, we talked about feature selection, dimension reduction, and regularization. These are all ways to take the linear model:

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \dots + \hat{\beta}_p X_p$$

and make this linear model LESS wiggly, so as to reduce the variance.

Today, we are going to talk about ways to take this simple linear model and make it MORE wiggly, so as to reduce the bias. For simplicity, to start, we will return to the case where $p = 1$. The model:

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X_1$$

is SO simple. We might be able to see just from looking at a plot that it is not expressive enough to model our data. What do we do about this? This is the topic of Divij's talk. But first, let's spend 5 minutes on Degrees of Freedom, since we skipped it a few classes ago.

7.2 Degrees of Freedom and Effective parameters!

We have been talking about how the bias/variance tradeoff is a function of model complexity. I have been calling this "how wiggly a function is"- or how capable the model is of overfitting to the training set.

One way to formally measure the complexity of a model is the *degrees of freedom*. I am sure that you have all heard of degrees of freedom before, but I am not sure that you have ever seen it defined formally. As far as I can tell, ISL avoids defining this formally: they just talk about it being the effective number of parameters in a model.

I feel like we should mention the official definition at least once. Assume that x_1, \dots, x_n are fixed. Let $g(y) = \hat{y}$, where g is a "model fitting procedure", that takes in the observed values of y as input and trains a prediction model and returns the training set predictions. We are not writing this as \hat{f} because it does not refer to a specific fitted model: it refers to the procedure. And we are not writing it as a function of X because the X s are just hiding in the background since they are not random. According to this paper, <https://www.stat.berkeley.edu/~ryantibs/papers/sdf.pdf>,

$$df(g) = \frac{1}{\sigma^2} \sum_{i=1}^n Cov(Y_i, g(Y)_i),$$

where σ is $Var(Y_i)$, which is assumed to be constant across i .

Take a few moments to let this definition sink in. What is it telling us? What is the degrees of freedom of a procedure that memorizes the training set (and therefore gets 0 training error)?

$$df(\text{memorizer}) = \frac{1}{\sigma^2} \sum_{i=1}^n Cov(Y_i, Y_i) = \frac{1}{\sigma^2} \sum_{i=1}^n \sigma^2 = \frac{1}{\sigma^2} n \sigma^2 = n.$$

This tells us that the degrees of freedom of 1NN, for example, which we know to be extremely wiggly, is the size of the training set n !

Calculating degrees of freedom is straightforward for simple procedures like linear regression. For more complicated procedures, actually computing this is really complicated! So we might need to estimate or approximate some degrees of freedom.

The concept of degrees of freedom is very related to the concept of "optimism". These concepts are covered in Chapter 7 of **ESL** (not ISL). I have been sprinkling these concepts in throughout the course, but next week I will be sure to

put a question specifically about model complexity on the HW!

I don't think we are going to spend too much more time on this. ISL says it is "outside the scope of the book". But I think we should mention DF or approximate DF for some models we have seen so far.

- Linear regression with p coefficients: p degrees of freedom.
- KNN: we are making around n/k unique predictions, so we have around n/k degrees of freedom. In fact, n/k is an unbiased estimator for the degrees of freedom apparently. But the true DF depends on the distribution of the X s: how many of your "regions" overlap affects the number of effective parameters.
- Ridge: You can prove directly from the definition that ridge regression with p coefficients has $< p$ degrees of freedom, and the DF shrinks as λ grows. This is because the penalty makes it strictly less wiggly than unpenalized regression on all p predictors.
- Lasso: If you apply Lasso with p predictors but end up with p_0 predictors, it turns out that p_0 is an unbiased estimate for the degrees of freedom of Lasso! This is magic! We are considering many more than p_0 possible models, which would make you think it is bigger than p_0 . But, we are also ending up with a p_0 -dimensional linear regression model and applying shrinkage to this model, which would make you think it is smaller than p_0 . These two factors apparently average out to p_0 !
- Stepwise regression or subset regression: If you start with p predictors but end up with p_0 predictors in the final model, your DF should be larger than p_0 . Because, even though your final model uses p_0 , you also gave yourself a bunch of options for models to try out, which allows for a bit more potential for overfitting compared to just lm on a fixed set of p_0 predictors.

TLDR: Ridge and Lasso are both LESS WIGGLY than linear regression with p predictors. They have fewer than p "effective parameters" because of the regularization.

Note that, in this group, KNN is the only model whose degrees of freedom grows with the size of the training set n . This is in fact what it means for KNN to be *nonparametric*.

Returning to our example for today, suppose that we only have one predictor X and we are going to fit the model:

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X_1.$$

This is sort of a bummer because it only has two degrees of freedom! This might not be enough! We all already know that we could increase the complexity (df) of this model by instead fitting something like

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_1^2 + \hat{\beta}_3 X_1^3.$$

Today we will discuss even better ways to do this!

7.3 Splines

There are actually two distinct topics here: regression splines and smoothing splines. At first glance, they do not seem that similar! But a beautiful theorem tells us why they are related, and therefore why they both get the name "spline".

7.3.1 Regression splines

These start from a pretty straightforward motivation.

We already know about polynomial regression, We can fit the model

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_1^2 + \hat{\beta}_3 X_1^3$$

using "lm()" in R, by just passing in our squared terms and cubic terms as new predictors.

But we don't need to limit ourselves to polynomials. Many of you have likely also used models of the form

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 1(X_1 < c),$$

which produces a step function. This is what we do when we decide that maybe our model is better if we convert a numerical predictor variable X_1 to a categorical predictor variable. Once again, we can just fit this model with "lm". More generally, we can use any set of fixed/known basis function to turn X_1 into an arbitrarily wiggly and complicated function. As long as we use less than n basis functions, we can still fit the model with 'lm()' in R.

People have proposed regression splines as a really nice way to make basis functions. These are piecewise polynomials. The idea is that, to make the model $\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_1^2 + \hat{\beta}_3 X_1^3$ more complex, we could add a term of degree 4

or degree 5. But high degree polynomials tend to have really crazy behavior. What if, instead, we limited ourselves to cubic, but we added a “knot”, and we let the coefficients be different on either side of the knot. This could be nice, but this could also lead to discontinuous functions which we don’t want! The spline basis representation gives a way to make continuous piecewise-constant polynomials with smooth derivatives. A *natural spline* is additionally linear outside of the range of the training X s, which keeps predictions from blowing up to crazy values if we extrapolate beyond our training data.

I don’t really care if you know the formulas for a spline basis representation: I just care that you understand how the concept a spline fits into our framework of the bias/variance tradeoff. We could decide how many knots to give our spline using cross validation! I also care that you understand that fitting a spline just means using “lm()” on a new set of features. You will do this on your homework!

7.3.2 Smoothing splines

These start from a totally different, non-parametric motivation. Still assume that we only have one predictor. Suppose that we wanted to find the function $g()$ that minimizes

$$\sum_{i=1}^n (y_i - g(x_i))^2$$

on the training set, but we did not want to specify a particular form for $g()$. For example, we don’t want to choose in advance to make it linear, or a cubic spline, or a sin curve, etc. What we do know is that a function that is too wiggly will have high variance and will not generalize well. So, we decide that we will add a penalty term for “non-smooth” functions $g()$.

This sounds like regularization! Indeed, we will try to minimize:

$$\sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt.$$

Integrating over the second derivative just means summarizing how much $g'(t)$ changes over the entire range of t . If we fit a linear model, then $g''(t)$ is 0, and so there is no penalty on a linear model. The parameter λ is a tuning parameter that tells us how much we should penalize non-linear or wiggly functions.

The function $g()$ that minimizes this equation is called a smoothing spline. Why is it called a spline? See the beautiful theorem.

7.3.3 Beautiful theorem

It turns out that, out of all functions $g()$, the one that minimizes

$$\sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt.$$

is actually a natural cubic spline with knots at every training data point x_1, \dots, x_n . Woah!!!

Does this mean that, to get an optimal smoothing spline, we can just use “lm()” in R and regress y on the natural spline basis function with x_1, \dots, x_n knots?

- No
- The first reason we cannot do this is that the natural spline basis function with x_1, \dots, x_n knots has more than n terms. This means that we would be fitting a linear model with n datapoints onto more than n predictors. R cannot solve this for us, since $(X^T X)$ cannot be inverted.
- Also, if we COULD solve the least squares solution for regressing y on the natural spline basis function with x_1, \dots, x_n knots, it would get 0 training error and would memorize the training set. This is not actually what we want!
- The penalty term λ ends up enforcing that we also shrink all of our coefficients towards 0: as we would with ridge regression.
- So, we fit this using a different R package. I think it might be the default in “geom_smooth()”! But it is really cool to know about the connection with a natural cubic spline.
- QUESTION: from the penalty perspective, why does it make perfect sense that $g()$ is a NATURAL cubic spline, meaning that it is linear outside of the range of the training points?

7.4 GAMs

7.4.1 What is a GAM?

Let's build on the spline idea, but now let's suppose that we actually have more than one predictor. We have X_1, \dots, X_p , and we want to let the function $\hat{f}(X)$ be wiggly in all of the predictors (if needed).

We could try to do a spline basis expansion for every single variable. But this would start to get kind of crazy! We don't necessarily have a larger n than we did before. If we add 12 coefficients for every predictor variable, we are going to have the number of coefficients exceed the number of variables pretty quickly: and then we can no longer fit with `lm`!

Let's restrict our attention to models where we don't allow the predictors to interact with each-other. This is limiting in some ways. But, if we want to start letting variables interact, the complexity of our problem blows up even more quickly.

So, let's model:

$$Y = \beta_0 + \beta_1 f_1(X_1) + \beta_2 f_2(X_2) + \beta_3 f_3(X_3) + \dots$$

Each $f_j()$ could be any basis function representation of X_j . It could simply be a linear model if we think that Y is linear in X_j , but it could also be a full on smoothing spline for a different predictor. Note that, because f_j can be whatever we want, it can also absorb the coefficient β_j . So we will actually just write this as:

$$Y = \beta_0 + f_1(X_1) + f_2(X_2) + f_3(X_3) + \dots$$

A nice thing about the additivity is that I can visualize one variable at a time, and decide how wiggly I think Y needs to be in that particular variable. This does not take into account the idea of "holding all other variables constant", but can be a nice first approximation.

GAMS are really powerful. Suppose that we really want to be able to *interpret* the effect of X_1 on Y , but we want to make sure we have already sucked up all possible variation that could be explained by other components X_2, \dots, X_p . We could let $f_j()$ be a smoothing spline for all $j > 1$, but force $f_1()$ to be a linear function. This would let us interpret the linear effect of X_1 on the part of Y that would still be noise even if we fit a many-dimensional `geom_smooth()` to all other predictors. That is cool!

7.4.2 How do we fit a GAM?

Unless all components f_j are pre-specified with a parametric basis function, we cannot fit a GAM using least squares. Instead, we take advantage of the additive structure to do an iterative back-fitting procedure. The function `GAM` in R implements this for you, and by default it uses a natural spline for each f_j .

To learn about the back-fitting algorithm, see page 297 of ESL. The idea is that, since everything is additive, we can iteratively fit a model that predicts $Y - \beta_0 - \sum_{j \neq j^*} \hat{f}_j(X_j)$ using X_{j^*} . This will let us see the function that "right now" best predicts Y using X_{j^*} , after taking into account all variation currently explained by other predictors. We start with random guesses for all of the \hat{f}_j terms, and we iterate through $j = 1, \dots, p$ many times until convergence.

What I want you to know about back-fitting is: the fact that we can do one variable at a time makes this relatively easy! With neural networks next week, we will not be able to go one variable at a time, because we will have interactions! This will necessitate a more advanced algorithm.

7.5 What do we think about splines and GAMs?

- Bias: they have less bias than linear regression if the true data generating mechanism is not linear.
- Variance: they can have high variance if we let them be too wiggly.
- Interpretability: they are not black boxes. they are parametric, and we can examine them to figure out which predictors impact our predictions. but they can get kind of complex and kind of hard to interpret! Especially without drawing a graph.
- Computational efficiency: regression splines can be fit very efficiently. GAMs require an iterative algorithm, but the fact that they are additive still makes this algorithm reasonably efficient.
- Anything else I am missing?

You could add a spline basis function to a logistic regression. You can also definitely use GAMs for classification. So ... none of this was specific to regression! The concepts of wiggly and less wiggly apply to classification too!

The main idea of class next Monday will be: what happens when we drop the additivity requirement?! It is going to turn out that, when we let models be non-linear AND we allow for interactions, we get crazy complex models! Such

as neural networks!

8 Monday, March 3: Introduction to a simple neural network

Main references: ISL 10.1, 10.2, 10.6, 10.7, and 10.8. I also think that ESL Chapter 11 has good context, but some of it is beyond the scope of this class.

8.1 A statistical motivation

8.1.1 GAMs, but with feature engineering?

Last week, we discussed GAMs. These are functions where we let

$$\hat{Y} = \hat{\beta}_0 + \hat{f}_1(X_1) + \hat{f}_2(X_2) + \dots + \hat{f}_p(X_p).$$

We let the functions $\hat{f}_j()$ be extremely non-linear functions, such as smoothing splines. Even if we let $\hat{f}_j()$ be completely non-parametric and wiggly, fitting a GAM is not too bad. The additive structure means that we can iterate through one variable at a time, and just fit a one-dimensional model of $Y - \sum_{j \neq j^*} \hat{f}_j(X)$ on X_{j^*} .

But, the additive structure of the model will always introduce some bias if the true data generating mechanism is not additive! Consider image classification. Why would the image's class be determined by a sum of functions of every individual pixel? It is really the patterns between the pixels that matters. And we don't need the ability to interpret the effect of each individual pixel: this is going to be the type of data that neural networks are really good for!

We know that we could fit a GAM to the principal components of X . We could also manually add interaction terms. All of that would sort of help our GAM problem. But let's try to do this in an automated way.

The idea for today is that maybe we want to let:

$$\hat{Y} = \beta_0 + \sum_{k=1}^K g_k(w^T X).$$

I am getting this particular form of this equation from ESL section 11.1. This is a GAM, but it is a GAM applied to K features, each of which is a linear combination of ALL of the X s. And we haven't yet determined the weights w that will make up the linear combinations. In projection pursuit regression (ESL, 11.1), the $g_k()$ are smoothers, as in additive models. Today, we will have them be something simpler.

As noted in ESL, project pursuit regression never became widely used. This is because, at the time of its introduction, you would need really powerful computers to actually fit the model. However, the reason that it is exciting for us is that it was a completely statistical motivation for a topic that was simultaneously becoming popular in the field of AI. We will discuss the AI motivation later.

8.1.2 The neural network model

ISL introduces its first neural network model on page 404. The idea is that we will let $\hat{Y} = \hat{f}(X)$, where $\hat{f}(X)$ must come from a class of functions that can be written as:

$$\begin{aligned} f(X) &= \beta_0 + \sum_{k=1}^K \beta_k h_k(X) \\ &= \beta_0 + \sum_{k=1}^K \beta_k g \left(w_{k0} + \sum_{j=1}^p w_{kj} X_j \right). \end{aligned} \tag{15}$$

The function $g()$ is a non-linear activation function. The idea is that we model Y as a linear combination of K "hidden units" or "activations". K is a parameter that we get to pick. Each of these is a non-linear transformation of a simple linear combination of the X variables. Note that if we chose $g()$ to be a linear function, this would just collapse down into a linear model with no interaction terms.

So ... this is our neural network! We need to use our training data to come up with estimates for $\beta_0, \beta_1, \dots, \beta_K$ and w_{k0}, \dots, w_{kp} for $k = 1, \dots, K$. This is a total of $(K+1) + (p+1) * K = (p+2) * K + 1$ parameters. We want to come up with guesses such that $Y \approx \hat{f}(X)$, hopefully on both our training data but also on unseen test data.

Because of the fact that there are so many parameters in (15), we often draw this model as a picture so as to keep track of everything. The picture is shown in Figure 2. The idea is that every line in the picture corresponds to a parameter that we learn. Once we have drawn this as a picture, we can call the setup of our model the “architecture”. The architecture is up to us! If we wanted to have hidden node A_1 not connect to X_1 , that would be totally fine: we would just have one less arrow in our picture and one less parameter to estimate.

There is one note in Figure 2 that is not reflected in (15). In (15), I was imagining a regression problem: so once we take our linear combination of linear combinations, we have $\hat{Y} = \hat{f}(X)$. If we are doing classification, we might want a final step where we convert the output of our network into a probability (using a sigmoid or softmax function), and then maybe convert it to a discrete prediction. Multi-class classification is seamless: we just have one output node per class, and we later convert to probabilities with a softmax function.

10.1 Single Layer Neural Networks 405




Figure 2: A visualization of (15). The arrows between the inputs X and the hidden later each represent a parameter w_{kj} , where the k tells us which hidden layer we connect to and the j tells us which input variable to connect to. Sometimes we draw an extra input layer node to represent the intercept. The arrows from the hidden layer to the prediction layer are the parameters β_k . All of these parameters must get estimated.

While we won’t be able to fit models like the one shown in Figure 2 directly with least squares, we will be able to fit it by optimizing over a loss function: something that you all have already seen many times! So ... a neural network really isn’t all that different from a linear model! This is a very statistical perspective on neural networks.

8.1.3 Deep learning

The really key idea in Figure 2 is that: **linear combinations of non-linear transformations of linear combinations** can be arbitrarily complicated and wiggly! In fact, there is a theorem that says that any continuous function $f()$ can be expressed in the form (15) as long as $g()$ is non-polynomial and K is big enough. Another idea, if we don’t want to make K huge, is to add more layers! And more levels of non-linearity! This is shown in Figure 3, which is also taken from ISL.

The idea of approximating really complicated functions with these networks has really taken off in the field of deep learning! Models like the one in Figure 3 have become very state of the art in image classification, computer vision, LLMs, and basically any other big-tech application that you can possibly think of. Whatever big data you see out there in the real world- it is probably making use of deep neural networks! It almost makes me feel like there is no need to teach the other methods that we have been discussing or will be discussing in this class. But only almost! Deep neural networks are not appropriate for every circumstance, and they certainly have a lot of downsides!

8.2 Historical context and a less statistical motivation

When neural networks were first developed in the 1940s, there was a biological interpretation. Scientists wanted to model the living brain! An early reference is McCulloch and Pitts (1943).

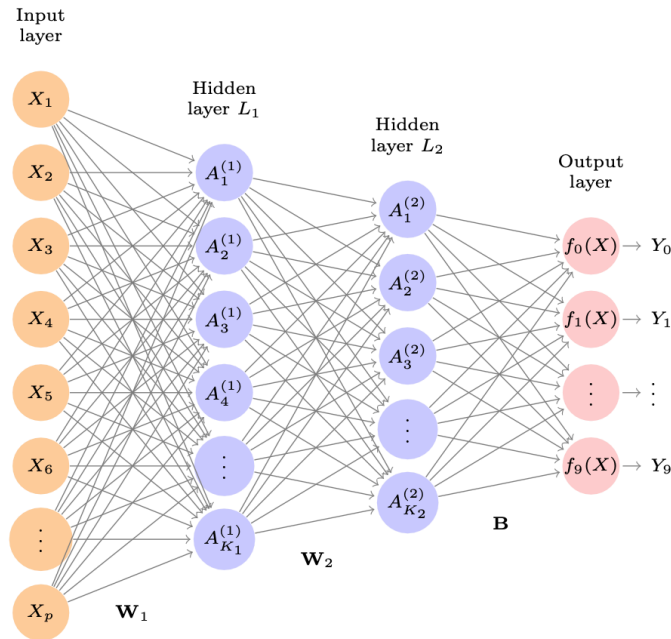


FIGURE 10.4. Neural network diagram with two hidden layers and multiple outputs, suitable for the MNIST handwritten-digit problem. The input layer has $p = 784$ units, the two hidden layers $K_1 = 256$ and $K_2 = 128$ units respectively, and the output layer 10 units. Along with intercepts (referred to as biases in the deep-learning community) this network has 235,146 parameters (referred to as weights).

Figure 3: A multilayer neural net. Note that here we have multiple outputs Y because this is for a categorical output with 10 categories. We represent this as a vector of length 10, where $Y_k = 1$ if $Y = k$ and $Y_k = 0$ otherwise.

The nodes in the hidden layer are supposed to represent neurons. A neuron only fires if the amount of signal being passed into it exceeds a certain threshold. So, the activation function $g()$ was typically a step function: either the node (neuron) is turned on, or turned off. However, the step function was later replaced with differentiable alternatives for the purposes of actually fitting the model. But, the idea that the activation function should often be near 0 until it eventually gets “activated” or “fires” remained.

Do you think that neural networks work well because they model the human brain (which works well), or do you think they work well because they are very expressive linear models that do automatic feature engineering? It is up to you to decide on your interpretation! I think the latter, but I also do not know much about the real human brain!

A few more historical notes:

- The ideas of neural networks are so old and were developed independently by different fields!
- But ... when did they really take off?
- Up until 2010 or so, they were not popular in computer vision. Only Yann LeCun and Geoff Hinton were really pursuing these. After some students of Geoff Hinton won an ImageNet competition (used a neural network for an image classification task and got the best performance), then people started to pay attention.
- By the way, my history knowledge was really rusty and anecdotal. I found this slide deck from a course at Univ. Wisconsin https://sebastianraschka.com/pdf/lecture-notes/stat479ss19/L02_dl-history_slides.pdf and I liked the way the material was presented: he has screenshots from a lot of cool books. Check these out and check out the related sources if you want to know more!
- Things I learned: the ReLU activation function was an idea of Hinton+coauthor, and it was in 2010. And this change in activation function must have been a huge difference maker!
- You can also check out Section 1.2 of the free deep learning book: <https://www.deeplearningbook.org>.

8.3 What do I pick for my activation function $g()$?

- In projection pursuit regression, the idea was that this could be a smoothing spline or something! But projection pursuit regression was always going to have like $K = 5$ and one hidden layer. If we want to do BIG neural networks, we need something simple. And hopefully differentiable, because we will fit with gradient descent.
- From the “neurons in the brain” view of neural networks, a 0/1 step function made sense. A neuron is either turned on or turned off. But, since a step function is not differentiable, the sigmoid/logistic: $g(z) = \frac{1}{1+e^{-z}}$ was a nice alternative. Or tanh.
- But, nowadays, a really common one is: $g(z) = \max(0, z)$. This is called the ReLU. This keeps the motivation of a neuron being “turned on” vs. “turned off,” but now allows really strong signals to be carried forward by having larger magnitude. One issue with this activation function is the “dead ReLU” problem, which we will discuss.

8.4 Themes

People have the tendency to treat neural networks and deep learning as magical! I hope to convince you that these are not magical, and that they in fact relate to many things that you have already seen in this class: e.g. feature engineering, basis expansion. And I hope you can start to see how neural networks relate a bit to what I talked about on the very first day of class: the biggest difference between modern machine learning and classical statistics has to do with “how much is pre-specified”. Neural network relates to feature engineering or basis expansions, but the form of the new features or the basis functions was not pre-specified.

Because NNs are not so different than everything else we can see in this class, we can discuss our favorite themes.

- Bias:
 - We can approximate ANY continuous function f with a neural network with one hidden layer and a non-polynomial activation function. We might need a really huge (but finite) value of K . But still! This is so cool. It means that we are not limited by modeling assumptions for a neural network: we should be able to model any true f really well.
 - One note: this theorem says that any continuous function f can be expressed using a picture like Figure 2 and certain values of the weights. It does not say that we will be able to estimate the weights from our training data!
 - But still. Neural networks have really low bias! No bias if we use enough layers or nodes.
 - They work really well!
- Variance
 - We know that variance depends on degrees of freedom, which is the effective number of parameters that we are estimating. If we make a lot of hidden layers, or if we make K really big, then surely we will have more parameters than we do training observations n . Isn't this really bad for variance? Won't we memorize our training set and overfit?
 - Short answer: statisticians were very skeptical of deep learning for these reasons for many years. Over-parameterized models should score poorly for variance! And they definitely do when you don't have a really big value of n .
 - We should regularize to help keep the variance under control. More on this when we discuss double descent! We should use a validation set to stop training before we overfit.
 - But sometimes, when you have enough data, NNs perform really well on test sets, even though they are over-parameterized. This also isn't magic: we need to distinguish between number of parameters and number of effective parameters
- Interpretability
 - This is our first true black-box model. It is really hard to understand where the predictions of a neural network model are coming from!
 - Do not pick deep learning if you have a scientific application where you need to interpret and explain your results!
 - Since the model itself is a black-box, we will need to use clever explainability techniques on top of the model to try to understand what is going on: this is a hot area of research that we will discuss after spring break.
- Usability (is the method “off the shelf”)?
 - Neural networks tend to require a lot more tuning than something like a random forest. They are NOT very “off-the-shelf”. This is a reason why they took a while to become popular. They were hiding in the background for many years before deep learning really took off.

- Evidence: I will not make you do a HW problem where you actually use neural nets, because the R packages are finicky!
- There is flexibility, which is nice: you can modify the architecture really to your liking. If you know what you are doing, this is great! But this flexibility does make it harder to make all-purpose, useable software.
- Computational efficiency
 - We are going to actually go over gradient descent and backpropagation on Thursday! So you will learn more about fitting.
 - We need to use a slow iterative algorithm to fit: and we always need to be worried that maybe we did not converge in our fitting!
 - However, there is also something very parallel and distributed about our computations, which makes huge deep neural networks actually feasible to fit.

9 Thursday, March 6: Fitting a neural network: backpropagation and gradient descent

Recall from Monday that our neural network model with one hidden layer is a model that makes predictions of the form:

$$\hat{y} = f(X) = \beta_0 + \sum_{k=1}^K \beta_k g \left(w_{k0} + \sum_{j=1}^p w_{kj} x_j \right). \quad (16)$$

Because this is a lot of parameters to keep track of, we sometimes draw the model as a picture. But ultimately, given our neural network architecture (our picture, which should tell us K and should tell us if the network is “fully connected”), we need to use our training data to FIT the network.

This just means coming up with values $\hat{\theta} = (\hat{w}_{01}, \dots, \hat{w}_{Kp}, \hat{\beta}_0, \dots, \hat{\beta}_K)$ that minimize the error over our training set:

$$\sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

Of course eventually we will also care about overfitting and test error. But to start, let’s just figure out how we might even minimize the loss when \hat{Y}_i has the form in (16). It is not going to be simple! We are going to need to use an optimization technique called gradient descent. The key insight for neural networks is that we can use the chain rule to do gradient descent in a very organized way, called backpropagation.

9.1 Gradient descent

Let θ be a long vector that stores all of our parameters in our network. So our goal is to minimize, with respect to θ ,

$$L(\theta) = \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

Unfortunately, for neural networks, this is too complicated to just take the derivative with respect to every single parameter and set them equal to 0. This would be a massive system of equations.

So instead of solving for a minimum directly, let’s try to do it iteratively.

Let’s start with θ^0 : a set of random guesses for every parameter in the network. These values will produce estimates \hat{Y}_i for every element in our training set: they probably will just not be very good estimates!

Then, let’s take steps along the negative gradient (with respect to θ) of the loss function. The gradient of $L(\theta)$ with respect to θ is denoted $\nabla L(\cdot)$. The gradient is simply the vector of partial derivatives with respect to each element of θ . Once we have taken this step, we will update our values of all of our parameters, so θ^0 becomes θ^1 . This new vector then yields new predictions \hat{Y} for every training point.

The idea is that, if we take steps in the direction of the negative gradient of the loss function, we always move a current guess θ^t to a new guess θ^{t+1} in a way that will make the loss smaller the next time we compute it. What a nice idea!

While we take a step in the “direction” of the gradient, we don’t want to take a step that is the SIZE of the gradient.

We want to take really small steps so that we don't "overshoot" a minimum. So we let:

$$\theta^{t+1} = \theta^t - \rho \nabla L(\theta^t),$$

where ρ is a small learning rate that we get to pick. Picking it can actually be really finicky: if it is TOO small you will never find a good value of θ . But if it is too big you can "overshoot" a minimum.

Note: if our step size ρ is going to be approximately the same for all elements of θ , we probably want the elements of θ to all be in the same "units". This means that we probably want to normalize our input variables X ! This is not changed from other algorithms.

We keep taking small steps in the direction of the negative gradient until our "guesses" stop changing. At this point, we have landed in a region of θ -space where the derivative is 0! This hopefully means that we found a minimum of the loss function!

Unfortunately, there is no guarantee that it is a global minimum. For really complicated non-convex loss functions, we can get stuck in suboptimal local minima. We will discuss strategies for avoid it later! For now, see Figure 4 for an illustration of gradient descent in a setting where θ is one-dimensional.

Gradient descent is a really general optimization technique! It is used in contexts that have nothing to do with fitting neural networks!





FIGURE 10.17. Illustration of gradient descent for one-dimensional θ . The objective function $R(\theta)$ is not convex, and has two minima, one at $\theta = -0.46$ (local), the other at $\theta = 1.02$ (global). Starting at some value θ^0 (typically randomly chosen), each step in θ moves downhill — against the gradient — until it cannot go down any further. Here gradient descent reached the global minimum in 7 steps.

Figure 4: An illustration from ISL of gradient descent for a very simple one-dimensional θ . In this case, we found the global minimum. But that is not guaranteed!

9.2 Backpropagation and the Chain Rule

Let's study gradient descent for a one-layer feed-forward neural network in the following very simple case. This is a network with one hidden layer and three hidden nodes ($K = 3$). There are also only three inputs ($p = 3$). For simplicity, I did not draw the intercept (bias) terms only the diagram.



In this case, suppose that we are doing regression, and so the function $f()$ at the very end is just the identity function. Recall that when we are doing classification, we need to turn our final output into a probability or a prediction using a link function. But for regression we are all set on our original scale.

For given values of w_{10}, \dots, w_{pK} and β_0, \dots, β_K , each input example x_i gets turned into a \hat{y}_i . Our squared error loss for our training set is:

$$\begin{aligned} L(\theta) &= \sum_{i=1}^n L_i(\theta) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{k=1}^3 \beta_k A_k \right)^2 \\ &= \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{k=1}^3 \beta_k g \left(w_{k0} + \sum_{j=1}^p w_{kj} x_{ij} \right) \right)^2. \end{aligned}$$

I wrote this at multiple levels of granularity. Sometimes it is nice to explicitly see that the loss function depends on every w and every β . Other times, it is kind of nice to hide this detail.

Suppose that we begin with random guesses for every β and every w . Then, we send our entire training set “forward” through the neural network and compute this loss function given the current predictions. Then, we compute the gradient. We step along the gradient, and we update our guesses of w and β accordingly.

For our neural network loss function, let’s figure out the steps along the gradient. The key idea is that we can chain rule!

In this example, since $f()$ is just the identity function, we have that:

$$\begin{aligned} \frac{dL}{d\beta_k} &= -2 \sum_{i=1}^n (y_i - \hat{y}_i) \frac{d\hat{y}_i}{d\beta_k} \\ &= -2 \sum_{i=1}^n (y_i - \hat{y}_i) A_k \\ &= -2 \sum_{i=1}^n (y_i - \hat{y}_i) g \left(w_{k0} + \sum_{j=1}^p w_{kj} x_{ij} \right) \end{aligned}$$

And then:

$$\begin{aligned}\frac{dL}{dw_{kj}} &= -2 \sum_{i=1}^n (y_i - \hat{y}_i) \frac{d\hat{y}_i}{dw_{kj}} \\ &= -2 \sum_{i=1}^n (y_i - \hat{y}_i) \beta_k g' \left(w_{k0} + \sum_{j=1}^p w_{kj} x_{ij} \right) x_{ij},\end{aligned}$$

where the specific form depends on the activation function $g()$. But hopefully we choose a differentiable $g()$ where this isn't too complicated.

Note that both of these partial derivatives imply that the direction of our step depends on our residuals $y_i - \hat{y}_i$: this is a good thing! We want to make the residuals smaller!

The reason that we call this special version of gradient descent “backpropagation” is that we can think of these applications of the chain rule as stepping backward through our network, and passing our residuals/errors along as we go. As we add more layers to the network, we add more chain rules! But nothing gets too difficult.

Recall that gradient descent lets:

$$\theta^{t+1} = \theta^t - \rho \nabla L(\theta^t).$$

So, the size of the step depends on the learning rate ρ , the size of the residuals $(y_i - \hat{y}_i)$, the magnitude of the gradient evaluated at θ^t , and the actual inputs x_{ij} . This is a lot of things! Making sure that we take the right sized steps can be very finicky.

9.3 Stochastic Gradient Descent

Instead of computing the loss function and the gradient over all n training instances at once, we typically go instance-by-instance or batch-by-batch. This is related to *online learning*, which you read about for HW2. A batch is just a subset of the n training observations. The motivation is that we will speed up our learning about θ by updating θ^t more often. If our training set is huge, we don't want to bother processing a massive training set with a bad guess θ^t . We should take advantage of the fact that we see a lot of errors in the first few training examples, and we should update θ^t to θ^{t+1} right away.

If we are using SGD and we end up iterating through the entire training set 10 times, this means that we used 10 epochs. But, this might mean that we used $T = 100$ updates of θ , if our batch size was 1/10 of the training set. We need to decide how many epochs we want, etc.

This is stochastic because our gradient computed on a small number of training observations should be a good approximation of our overall gradient. But the exact step we take depends on the specific training observation! It also turns out that the extra noise added with SGD can help us avoid local minima! That is cool too! We bounce around θ -space and explore a bit more.

Very strangely: it turns out that stochastic gradient descent enforces its own form of approximately quadratic regularization. I don't think that this is at all obvious. I learned from preparing for this class that the reason is because stochastic gradient descent gives updates for θ that remain in the row space of the inputs. Out of all solutions that would give 0 training error (in a setting where this is possible), stochastic gradient descent is supposed to find the one with the smallest L2 norm. “Supposed to” because I think it still might not due to initial conditions and step size. But anyway, I don't think this was the motivation for people coming up with SGD, but it is a recently studied property, and it explains double descent!! So cool.

9.4 Complications

While (stochastic) gradient descent is a very general idea that can be used in many contexts, it is ultimately only a way to find an approximate minimizer of a loss function. There are a lot of complications.

- We need to choose good initial guesses for the parameters. If we choose bad ones, we can get stuck at local minima.
- We need to choose a good step size! To converge in a reasonable amount of time without jumping over minima.
- We have to worry about “vanishing gradients” / the dead ReLU: where once a ReLU unit is outputting a 0 for every training example, its weights never get updated again.
- We need to choose batch size for SGD.
- We need to choose how many iterations or epochs to train for. In theory, I guess we should go until the θ^t stop changing (convergence). But ... what if we have used up all of our computational resources and we have not yet

converged?

- We also need to pick our model architecture! How many hidden nodes do we want? How many hidden layers? What activation function should we use?

Previously, to pick a tuning parameter, you used cross validation. Doing cross validation for ALL of these choices and coming up with the best possible parameter sets is a full time job. And, literally, this is a full time job that you could all go get! TLDR: you basically need to be an expert to fit a neural network. It is not an “off-the-shelf” algorithm.

9.5 Overfitting and model complexity

We know that a neural network model with K hidden units in one layer and p inputs needs to learn around $(p + 1) * (K + 1)$ parameters. If we add more hidden layers, this just gets bigger.

We know that, if we have n training datapoints, a model with more than n parameters should be able to memorize the training set. And we know that this is bad, because we will overfit, and will have poor generalization error.

In modern deep neural networks, we often have more parameters than datapoints! So, how can we prevent overfitting?

There are several strategies that people were already doing:

- Add explicit regularization to our loss function! Like λ times the L1 or L2 norm of our w s and our β s. This affects every step of our gradient descent calculation, and could make it more complicated. But it will keep the variance of our model down, as we know.
- Random dropout: A possibly simpler version of regularization. We can just randomly turn off some hidden nodes when we pass some training observations through the network. This forces other nodes to pick up the slack of these turned-off nodes. This means that a single node cannot be used to memorize a single training point— because it has to perform multiple tasks. What a super cool idea!
- Early stopping via a validation set. Suppose that we want to encourage our weights to be small. We can start with small initial values. Then, every time we finish a batch or an epoch, we can check our error on a validation set. Presumably at first, this will be decreasing as we update our weights. If, at some point it starts increasing, we can assume that we started overfitting, and we can just stop training (even if we haven’t converged). This is cool and should also work well!

Note that any type of regularization adds to our training complications. How to we pick our penalty parameter? How do we pick how much dropout to have, or how MUCH the validation error needs to increase in order for us to think we should stop? There is already so much going on! This is not automated!

Ok, but even with all of these strategies, don’t we think that these massive models will overfit? And perform badly in terms of the bias-variance tradeoff? Especially when we don’t have that many training observations?

- Mostly, yes I do think this!
- But empirically, we have seen a strange phenomenon called double descent.
- It turns out it is not actually that strange or mysterious! Let’s explore.

9.6 Double descent, and its relation to stochastic gradient descent

This is one of my favorite topics! See ISL 10.8 for a beautiful explanation. You will also explore this concept on your homework, and I may fill in these notes in more detail after Thursday’s class.

The story of double descent is shown in Figure 5. We know that as the degrees of freedom in a model approaches n , the training set size, the training error goes to 0. This is the interpolation threshold; or the point at which our model is simply memorizing the training data.

Classical statistical theory says that the test set should be really bad at this point.

But, in practice, people noticed that for deep learning, you could keep making the model MORE complex (more layers, more nodes) AFTER the interpolation threshold, and while the training error remains 0, the test error starts to decrease again. This is called double descent: and some people thought it was magic! They thought it totally changed our interpretation of ALL of statistics!! What happened to the bias-variance tradeoff?

It turns out that the issue with this plot is what we draw on the X-axis. If we draw the X-axis as “number of parameters”, double descent does indeed happen. But if we correctly make the axis the degrees of freedom, we will not see the double descent phenomenon. See this awesome recent paper: https://proceedings.neurips.cc/paper_files/paper/2023/hash/aec5e2847c5ae90f939ab786774856cc-Abstract-Conference.html.

On your homework, you will see how this plays out for linear regression onto natural splines. It turns out that, if we continue to increase the degrees of freedom of our spline basis past the interpolation threshold, we end up doing some automatic regularization that we didn’t (necessarily) realize we were doing. This means that we aren’t actually just




Figure 5: A classic illustration of double descent.

increasing model complexity without bound! We are fitting a regularized model, and so number of parameters is no longer equal to model complexity.

For deep learning, it turns out that stochastic gradient descent actually does implicit regularization. If there are many possible parameter vectors that would all give the same interpolating solution, SGD finds the minimum-norm solution. This is regularization! So ... there is no magic. Deep learning performs well because it has no bias and is ALSO regularized to have low variance.

9.7 Extensions of neural networks

We have barely scratched the surface. There are so many cool topics we could cover. Such as multitask learning, auto-encoders, convolutional neural networks, or recurrent neural networks. Auto-encoders also relate to semi-supervised learning. We are not going to cover these! But hopefully you now have a little bit of foundational knowledge to understand these better in the future. And you could always do one of these for your final project!

10 Monday, March 10: Classification and Regression Trees

This is one of my favorite topics! At first, this algorithm might seem like a totally “new idea”. But later, I hope you will see a lot of connections with algorithms we have already studied!

One note before we start: there are actually a lot of algorithms out there for building classification and regression trees! I will get to this in “historical context.” Whenever I don’t say otherwise, if I am talking about a tree algorithm, assume that I am talking about the CART framework, which was popularized by Breiman et al. in 1984.

10.1 The main idea of the algorithm

Let’s start with some really basic motivation for a regression tree. At first glance, it might seem a bit different than the other methods we have seen in this class.

The motivation begins with the left panel of Figure 6. We have two covariates, X_1 and X_2 . And then we have a numerical response variable Y . We begin the algorithm with the simplest possible model. This model just ignores the covariates and predicts $\hat{y} = \bar{y}$ for all observations. In this case, $\bar{y} = -2.6$. This is the “intercept only” model. It has MSE given by: $\sum_{i=1}^n (y_i - \bar{y})^2$.

The algorithm then proceeds in a way that might remind you of forward stepwise regression. The algorithm says: at this moment in time, what is the binary split in my covariate space that will most improve the MSE, if I now let each sub-region be summarized by its own sample mean. This is illustrated in the second panel of Figure 6. In this particular dataset, it turns out that the best way to chop our space into two is to draw a vertical line at $X_1 = -0.89$. Once we do this, any observation to the left of the split gets $\hat{y} = 10$, and every observation to the right of the split gets $\hat{y} = -7$. This cutoff of $X_1 = -0.89$ was chosen greedily after considering all possible splits.

Let’s write this down more rigorously. As of step 1 in the algorithm, our model has a single “region” in it. This region is $R_0 = \mathbb{R}^p$: the whole covariate space. The set of possible splits are indexed by $j \in 1, \dots, p$ and $s \in 1, \dots, n$, where

$x_{j,(s)}$ denotes the s th order statistic of the j th covariate⁴. At the first level of the tree, we search exhaustively for:

$$j^*, s^* = \arg \max_{j \in 1, \dots, p, s \in 1, \dots, n} \text{Gain}(R_0, j, s) = \sum_{i \in R_0} (y_i - \bar{y}_{R_0})^2 - \left(\sum_{i \in R_{L(j,s)}} (y_i - \bar{y}_{R_{L(j,s)}})^2 + \sum_{i \in R_{R(j,s)}} (y_i - \bar{y}_{R_{R(j,s)}})^2 \right), \quad (17)$$

where $R_{L(j,s)} = \{i \in R_0 : x_{j,i} \leq x_{j,(s)}\}$ and $R_{R(j,s)} = \{i \in R_0 : x_{j,i} > x_{j,(s)}\}$. We refer to these as the regions to the left and the right of the split. Note also that $\bar{y}_R = \frac{1}{|i \in R|} \sum_{i \in R} y_i$ just denotes the sample mean of y within a certain region. This was a lot of notation, but it can be worth it to make sure that you really understand what the algorithm is doing! But the idea is simple: we choose the split that most improves the MSE of our model right now. Alternate Gain functions are possible, but this MSE one is the one used by Breiman et al. (1984) in their very widely used CART algorithm.

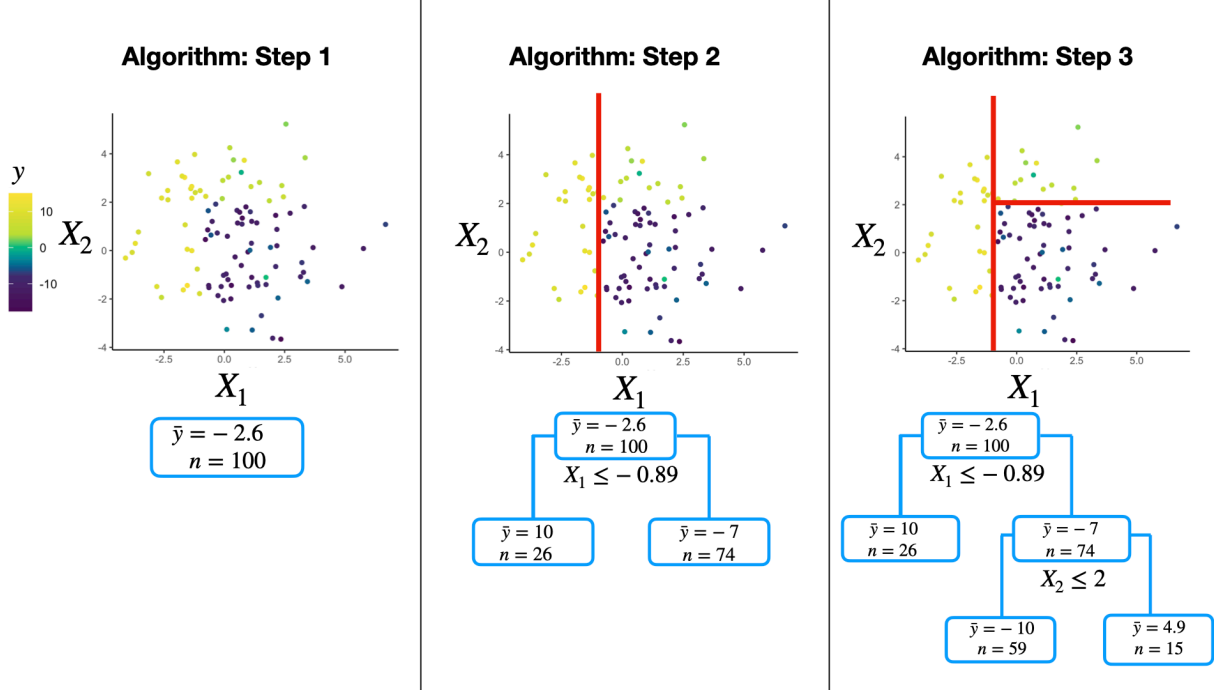


Figure 6: A little animation of the first three steps of a regression tree algorithm. Here, we have two covariates: X_1 and X_2 . And then we have a numerical response variable Y . In this case, it seems like the data is generated from a “true tree” structure, where there are rectangles in covariance space that define the average value of Y . The CART algorithm uses binary recursive partitioning to greedily search for the best possible rectangles, according to MSE!

Finally, the right panel of Figure 6 shows what happens next. Regression trees are *recursive partitioning algorithms*. So, the greedy search for the best possible splits, but now R_0 is replaced by the already-selected regions $R_{L(j^*,s^*)}$ and $R_{R(j^*,s^*)}$. We proceed until a stopping criteria is met.

The final model is a set of nested rectangles in covariate space. The prediction in each rectangle is just the sample mean of the training observations in that rectangle. We draw the model as a tree; as shown in Figure 6. We write this model as:

$$\hat{y} = \hat{f}(X) = \sum_{R \in \text{TREE}} \bar{y}_R \mathbf{1}\{X \in R\}.$$

Sometimes the final regions in the tree are called leaves, and the series of splits that lead to them are called branches. This really leans into the tree analogy. I will just as often refer to them as terminal regions or terminal nodes, etc.

We really have now gone over the main idea of trees! Now we will talk about some considerations and extensions.

⁴In the case of non-ordinal categorical X_j , we assign orders to the categories by sorting them by their average value of Y in the training set. Thus, the order statistics are still defined; we just have a lot of ties in the order statistics: meaning that $x_{j,(1)} = x_{j,(2)} = x_{j,(3)}$ if all three observations belong to the same category.

10.2 Considerations, extensions, and historical context

10.2.1 Stopping criteria

How do we know how big to build our tree? As you may have guessed, this is the main knob that controls our bias-variance tradeoff for trees.

The biggest possible tree would keep splitting until we have one training observation per leaf region. This tree would have worst-case depth n , but likely depth closer to $\log(n)$ if the tree remains more balanced. I think you all know that a tree like this will overfit severely to the training data, and we probably do not want to build this tree! We should stop before this!

Some simple ideas for when to stop are to pre-specify the maximum depth of our tree, or the minimum node size in our tree (i.e. stop splitting when the region has less than 10 training observations in it). However, these are stopping criteria that do not adapt to the amount of signal in the data. The true “right-sized” tree probably depends on the signal in our data! So, there are some more popular choices.

Consider (17). What if we decided to not choose ANY split if this gain does not exceed some threshold? For example, if the MSE does not improve by more than 5% when pick the best possible way to split this region, do not split this region at all. Such an idea is very similar to doing forward stepwise selection with an AIC or BIC stopping criteria. We recognize that ANY split will improve the training MSE *some*, but we want to make sure that the split seems “worth it”. In the `rpart` R package, this is controlled by the parameter `cp`. From the `rpart` documentation: “any split that does not decrease the overall lack of fit by a factor of `cp` is not attempted. For instance, with anova splitting, this means that the overall R-squared must increase by `cp` at each step. The main role of this parameter is to save computing time by avoiding splits that are obviously not worthwhile.” The default in `rpart` is 0.01, which is pretty small. This is an example of an *adaptive stopping criteria*.

Your textbook (ISL) calls the adaptive stopping approach above “short-sided”. The idea is that, due to the greedy nature of CART, it could be the case that no split will exceed the MSE threshold “right now”, but it could help us uncover an important interaction in the next level of the tree. So, it would be good to avoid stopping too early!

Another idea is to grow a tree that is purposely too large, and then to **prune** splits away from the tree. More formally, let:

$$L_\lambda(T) = \sum_{R \in T} \sum_{i \in R} (y_i - \bar{y}_R)^2 + \lambda |T|.$$

This is a loss function that incorporates both tree MSE as well as a penalty term that penalizes larger trees: $|T|$ denotes the number of leaves in a tree. It turns out that nice properties of trees tell us that, if we start by building a large tree T^{big} , then there is a simple nested sequence of sub-trees that corresponds to the best tree for different values of λ . And we can find this sequence of trees easily: starting from T^{big} we prune the “weakest link” in the tree one at a time. This gives us the nested sequence of trees. This is a “solution path”- like we had for Ridge or Lasso.

To be really complete, we would want to use cross-validation to pick the best value of λ . And then we would want to re-fit a tree to the full training set using this value of λ . If you are not yet comfortable with the idea of cross validation + refitting to entire training set, please ask questions before the midterm! It is an important idea!

10.2.2 Categorical predictors

CART can really seamlessly handle categorical predictors. We don’t need to turn them into dummy variables!

The Breiman et al. CART algorithm always makes BINARY splits; even for categorical variables. You might encounter other decision tree algorithms some day that would make a three-legged-split for a categorical variable with three categories.

Ordinal categorical variables still have order statistics, so we split in the same manner than we did above in (17). For unordered categorical variables with k categories, we do not need to consider 2^k possible ways to split the variable into two groups. We just order the categories by \bar{y} on the training set, and then only let ourselves make a binary split that respects the ordering of these categories.

NOTE: some critics of CART do not like the following fact. A numerical covariate has $O(n)$ chances to be chosen as the winning split. A binary categorical covariate has only 1 chance to be chosen as the winning split! If the binary covariate is truly important, but is associated with the numerical one, the numerical one will often appear in our tree! Just due to the extra random chance that it gets to be selected as a winner. This is too bad! There are modifications to the algorithm that try to get around this bias.

10.2.3 Categorical response

CART is really easily used for either regression or classification.

For classification trees, all we do is modify our gain function. Instead of using MSE, we split based on either Gini Index or Entropy. With either of these, we are choosing a split that makes the resulting child nodes as pure as possible: meaning homogenous with respect to y .

While we use Gini Index or Entropy to choose our splits greedily, we might still do cross validation using simple 0/1 classification error loss.

10.2.4 History

I got some nice historical notes from this paper: <https://pages.stat.wisc.edu/~loh/treeprogs/guide/LohISI14.pdf>.

The first classification tree algorithm was in 1963, and it was published under the name “automatic interaction detection (AID)”. This should already tell you something about why trees have been so popular- people absolutely love this feature that they can identify interactions without those interactions being pre-specified. At first, AID did not attract much attention from statisticians. People were worried about overfitting. People were also worried that, in the presence of correlated predictors, the conclusions could be spurious. Probably only one of the pair of correlated predictors will be selected for the tree, and the other will be totally left out. It might be dangerous to over-interpret the tree as signaling variable importance in that case⁵. At the same time, however, computer scientists were making their own decision trees for “concept learning.” It seems that an early reference is Hunt, Marin, and Stone (1966). The algorithm that I learned about in CS class is ID3, which was published by Quinlan in a series of work in the 1980s. This is again a setting where great ideas were coming out concurrently in multiple disciplines!

The credit for “popularizing” trees, at least in statistics, goes to Breiman et al. in 1984. I am pretty sure that one of the main innovations by Breiman et al. that really helped people take up trees was the introduction of efficient cost-complexity pruning.

Since trees first became popular, there has been a lot more work on them! Some statisticians (including me) do not like that the splits in a CART tree lack a notion of statistical significance⁶. One family of competing algorithms is called CTree; splits are chosen based on statistical significance, rather than based on an improvement in MSE. There are also models that make splits on linear combinations of variables, instead of single variables. Or models that fit an entire regression model to each leaf node, rather than choosing a piecewise constant model. And I am sure I am missing a lot of innovations!

10.3 What do we think about trees?

10.3.1 Bias

An extremely large tree has almost no bias. We can approximate basically any function with a big combination of step functions. So if n is big and our tree is big, trees are flexible. It is nice that trees can capture interactions between variables and other sorts of non-linear relationships without us needing to pre-specify.

However, a tree of reasonable size is very limited. Consider the case where $Y = 5 * X_1$. To approximate this well by a sequence of binary splits on the variable X_1 , we will need a lot of splits! As we add more and more, our approximation will get better and better! But this function will obviously be much easier to approximate with a linear model!

In general, if the true data generating mechanism is not made up of rectangles in covariate space, a simple tree might struggle. But, if you are worried that rectangles in covariate space are REALLY biased, not that a really big tree works a lot like 1-NN, which we know does not have bias. See below for more about the connection to KNN.

10.3.2 Variance

An extremely large tree has a lot of variance because it overfits. But even a small tree can have a lot of variance due to the greedy nature of trees. A small change to the input dataset can totally change our first split, which could then affect our whole tree- especially if we only plan to make a small tree. So trees do not score super well for variance.

10.3.3 Interpretability

Trees are a dream! We can explain our predictions so easily to non-experts! Even things like interactions do not seem scary when they are presented as a tree! However, I have some notes below about a fundamental issue. If we

⁵I agree!

⁶This is very related to my claim, which we never finished going over, that you cannot easily do inference after stepwise regression. Greedily searching for optimal things makes inference hard.

know that our tree is unstable (high variance), what are we really interpreting? This makes formal inference really important. Which is something I worked on in my PhD!

10.3.4 Use-ability

Trees are a dream! They are “off the shelf”. The only tuning parameter is tree size, and tree size is something that we understand. And the nested-tree property of the pruning algorithm makes it really nice.

We don’t need to scale our predictors or worry about whether or not we are including an intercept. Categorical predictors need not be converted to binary. Missing data is also seamless to incorporate using the concept of surrogate splits. You don’t need to start with preprocessing of variable selection. Trees do built in variable selection, and are not TOO impacted by curse of dimensionality.

10.3.5 Computational Efficiency

When I first learned about trees, I thought that the algorithm sounded slow. Search exhaustively for best possible split? But now that I know about things like neural nets, I’m less concerned haha. Trees are pretty easy to implement. The greediness saves us. $O(np)$ things to compute at each level to choose a split. At least no squared terms, right? And the tree depth will be at MOST n , usually more like $\log(n)$ if balanced. So maybe $O(np\log(n))$. This really isn’t bad. The pruning algorithm is also efficient.

10.4 Drawing connections!

Right now, it might seem like decision trees are a random topic we have thrown in that are totally different from the other algorithms that we have seen. This is not true! Let’s draw some connections.

10.4.1 Write as a regression model or optimization problem

Really, a regression tree defines a model class of piecewise constant models on rectangles in covariate space. So we could just abandon the whole tree idea and say that we are looking for a set of regions R_1, \dots, R_T such that we minimize

$$L_\lambda(T) = \sum_{r=1}^T \sum_{i \in R_r} (y_i - \bar{y}_{R_r})^2 + \lambda|T|,$$

and where $\hat{y} = \sum_{r=1}^T \mathbf{1}(X \in R_r) \bar{y}_{R_r}$.

This looks a bit more like an optimization problem or a regression problem. We are regressing onto lots and lots of possible step-function indicator variables, but we are doing variable selection first to decide which ones to include. We can think of trees in this way! And then the innovation is just that this loss function will be really really difficult to minimize exactly. So, we give up on searching fully for the optimal tree. We instead use our greedy, top-down approach. Which finds a pretty good tree, but not necessarily the very best optimal tree. This reminds us of the difference between best-subset regression, which is infeasible, and stepwise regression, which is a greedy approximation.

10.4.2 How is a regression tree like KNN?

At the end of the day, the prediction $\hat{y} = \hat{f}(x^{\text{test}})$ for a new datapoint x^{test} is the sample mean of some training points y that are near x^{test} in covariate space. This sounds a lot like KNN! How is it different?

- In finding the points that are “near” x^{test} , we do not consider all covariates X_1, \dots, X_p . We only consider the ones that were selected for the tree. If we did a good job selecting splits for the tree, this should really help with the curse of dimensionality problem that KNN encountered. Irrelevant variables do not contribute!
- And we selected rectangles that lead to good predictions! So we should have retained important directions while ignoring irrelevant ones.
- We still have the “prototype” interpretation: “you got these predictions because other previous points in your rectangle had this average response.” That is a really nice prediction!
- Overall, regression trees can maybe be seen as something that improves on KNN.

10.4.3 Are regression trees parametric or non-parametric

We just said that regression trees are kind of like stepwise regression (which is very parametric), but we also said that they are kind of like KNN (which is non-parametric). Which is it?

Remember that the definition of non-parametric is that our model complexity grows with our sample size n . As we said, the absolute biggest tree we can grow has n leaves: one per training datapoint.

So, if we are building trees to unconstrained depth, then regression trees are non-parametric and are more like KNN. Or, if we build trees with the restriction that the minimum node size is 10 but do no other pruning, then our model really is quite a bit like 10-NN. If we add more training datapoints, we can make more nodes, and so the complexity grows. Unless of course we run out of possible splits, which could happen if our variables are all categorical with not that many categories.

However, if we build trees to a maximum depth of 3, then this is parametric and is a lot more like stepwise regression. We could enumerate the set of all possible models based on the size of our covariate space, and the models would not get more complex as we increase n ; we would just get more observations per terminal node which reduces variance. Actually, we technically have more possible splits to choose from at each point when n increases (order statistics), so if you want to be really precise about this being parametric imagine that your X s are discrete and can take on only a set number of values.

Overall, I think that the difference between fixed-depth trees and fixed-node-size trees as parametric vs. non-parametric models is interesting! And of course, adding in pruning changes the complexity again. You could study this yourself on a final project!

10.5 Philosophically, can we really call an unstable model interpretable?

I think there is a really important point when we talk about the pros and cons of decision trees. I will try to illustrate this point with my R demo.

Regression trees are supposedly so great because they are interpretable. But ... how do we know when we are seeing the truly important variables vs. seeing noise? We could also perturb our training dataset slightly and get a totally different tree, in the context of correlated predictors or similarly important predictors, etc.

Isn't this really bad? What does it mean to interpret something that is unstable? We should probably add some notion of statistical significance to CART. Or work on making it more stable! This is the topic of a lot of research, including my own. Let me know if you want references!

11 Thursday, March 13: Support Vector Machines

Like Monday's class, today's class might at first seem like we are just randomly jumping to a new algorithm. But I hope that by the end of class, you are once again motivated to think of this instead as a new lens through which to study our important themes from the class. Today's theme is all about feature engineering in a really clever way, and what that gets us. SVMs also have some historical importance!

11.1 Motivation for SVMs

To motivate today's class, we are going to be thinking about a picture that we haven't thought about since the day that we covered LDA. The picture is of two quantitative predictors X_1 and X_2 , and a categorical response y (shown by the colors). See Figure 7.

Consider the left panel of Figure 7. Our goal is to predict Y using X_1 and X_2 . In this case, the task looks almost unbelievably easy. The classes are linearly separable!

We learned during our LDA lecture that logistic regression, surprisingly, does quite badly in this perfectly separable case. If you try to fit a logistic regression in R to this data you will get warnings about convergence and "fitted probabilities of 0 or 1" occurring. Logistic regression is all about modeling probabilities of $Y | X$, and there the estimates for certain regions all end up being 0 or 1, which makes the exact coefficients really uncertain.

You can envision this problem a little bit in the middle panel of Figure 7. All three lines perfectly separate the two classes, but they all have totally different slopes. How do we know which line to choose? You know based on last class that a classification tree will choose the vertical straight up and down line. But ... is that really going to be the line that generalizes best to new examples?

LDA and QDA chose between these different lines by making a very strong assumption. They assume that $X | Y$ is Gaussian. See the right panel of Figure 7. Then, you can draw estimated Gaussian contour lines for each class. The decision boundary chosen by LDA or QDA has to do with when these contours are set equal to one another: when is the red class or the blue class more likely, given X , based on the estimated Gaussian densities?

What if we want a way to pick between all of the lines in the center panel of Figure 7? And we want a way that does not make a Gaussian assumption, or arbitrarily restrict itself to straight lines defined by a single variable? This is the




Figure 7: A figure to motivate SVMs, and their differences with logistic regression, LDA, or QDA.

idea of the maximum margin classifier, which is summarized in Figure 8, which is taken from ISL. The idea is quite simple: let's pick between all of the possible separating lines by choosing the one that is as far as possible from all of the training observations.

11.2 Maximum margin classifier using constrained optimization

How do we actually fit the line in Figure 8? There is some math that relates to how we actually draw these lines onto plots. Also, note that we could have more than 2 X variables, and then we would be using a plane and not a line to separate our classes.

In general, we will write our linear boundary as the line:

$$\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p = 0.$$

The idea is to pick a line so that $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p < 0$ whenever $y = -1$ and $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p > 0$ whenever $y = 1$. For today, we are doing binary classification and we are writing our two classes as -1 and 1 , for simplicity.

Any of the lines in the center panel of Figure 7 actually have this property. So now the idea is to do even better. Let's both have it be the case that $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p < 0$ whenever $y = -1$ and $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p > 0$ whenever $y = 1$, but also have it be the case that $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$ is basically never TOO close to 0. Because when it is near 0, it means that points are close to the line and we have uncertainty. If all of our points are far from the boundary, we have less uncertainty.

One quick note about these hyperplanes: if our class boundary is the line $\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p = 0$, then $c(\beta_0 + \beta_1 X_1 + \dots + \beta_p X_p) = 0$ defines the same class boundary for any constant c . Thus, in defining planes, we restrict our attention to β vectors where $\|\beta\|_2^2 = 1$: to make sure that we have a unique solution.

So, the maximum margin classifier says:

$$\begin{aligned} & \underset{\beta_0, \beta_1, \dots, \beta_p, M}{\text{maximize}} && M \\ & \text{subject to} && \sum_{j=1}^p \beta_j^2 = 1 \\ & && \text{and } y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) \geq M. \end{aligned}$$




FIGURE 9.3. There are two classes of observations, shown in blue and in purple. The maximal margin hyperplane is shown as a solid line. The margin is the distance from the solid line to either of the dashed lines. The two blue points and the purple point that lie on the dashed lines are the support vectors, and the distance from those points to the hyperplane is indicated by arrows. The purple and blue grid indicates the decision rule made by a classifier based on this separating hyperplane.

Figure 8: The maximum margin classifier picks the line that is far as possible from all observations.

The first constraint is just for uniqueness. The second constraint says that every training observation is correctly classified (if M is positive), and the fact that we are maximizing M says that we are trying to make all of our prediction values far from 0 (recall that y_i is just -1 or 1).

For your purposes: know that people are good at convex constrained optimization. So this can be solved if the classes are linearly separable: i.e. if a separating hyperplane exists!

There are actually a few interesting properties of this solution. One interesting property is that the optimal hyperplane actually depends on the data ONLY through the points that lie ON the margin M : the other points do not contribute to the β s at all. In Figure 8, there are only 3 of these “support vectors” that actually impact our classifier. That is interesting— we could add 1,000 blue points to Figure 8, and as long as we add them on the “blue side” of the current hyperplane, our solution does not change at all. This is interesting!! And is certainly different than LDA— in LDA, the overall class proportions affect our decision boundary (via the prior).

You might be worried that this property of SVMs— that they only depend on a few observations— could lead to overfitting. I am certainly worried about that! We usually do not want ONE datapoint to impact our entire classifier that much!

The question you should definitely be asking yourself right now is: what if our classes overlap? Real data never looks like Figure 7. How do we use an SVM in this case? To answer this question, we will discuss two concepts.

- Concept 1: we can rephrase our optimization problem to allow a small number of mistakes. (this will actually also help with the overfitting concern, even when our classes are technically separable).
- Concept 2: if we transform our feature space enough, we can probably make the class linearly separable in new feature space.

We will discuss both of these!

11.3 Support vector classifier

The support vector classifier just takes the maximum margin classifier and says “let’s be okay with a few mistakes” (concept 1).

We once again write a constrained optimization problem, and once again smart people know how to solve it efficiently because it is convex.

$$\begin{aligned} & \underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n, M}{\text{maximize}} && M \\ & \text{subject to} && \sum_{j=1}^p \beta_j^2 = 1 \\ & && \text{and } y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \\ & && \text{and } \epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C. \end{aligned}$$

The only thing that we added here is that a training datapoint is allowed to be on the “inside of the margin” ($0 < \epsilon_i < 1$), or even on the wrong side of the hyperplane ($\epsilon_i > 1$). But, we limit how many of these points are allowed using the total cost C , which is something that we pick. We still make predictions based on whether or not $\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$ is greater than or less than 0: we just know that there are some mistakes in our training set.

We usually pick C with cross validation. A large C leads to a cross validation with more bias but less variance. Because, when C is big, we let lots of individual datapoints have non-zero ϵ , which means we let them “break our rules”. This lessens the dependence of the classifier on the individual observations, which means less overfitting and less variance.

Once again, it turns out that ONLY datapoints with non-zero ϵ affect the final classifier. If you remove other points, or add other points that fall outside of the margin on the correct side, you do not change the classifier! The observations that DO affect the classifier are called the support vectors. When C is big: we have a LOT of support vectors. So we depend on a LOT of the data— low variance! When C is small: there are only a few support vectors — we fit the data really well but we have high variance!

Overall though, we are still only focusing on observations that are near the boundary. Points that are very clearly members of one class or the other do not affect our classifier rule! This makes it more like logistic regression than LDA.

11.4 Support vector machine

If our data are not linearly separable, we could just make our cost C really big until we end up with a valid classifier. But ... sometimes a hyperplane in our original feature space is simply not going to give us a good rule! We now turn to Concept 2.

If the relationship between our predictors and our classes is not linear, we should not try to use a separating hyperplane! But the key insight of a support vector machine is that maybe we can use a separating hyperplane in a new feature space. This is exactly the same idea as just adding polynomial terms to linear regression when we think our function is not linear!

If you just put in X_2^2 as a “new covariate”, then you can still fit a support vector classifier. It will give you an equation that has a $\beta_k X_2^2$ term in it. This is not a hyperplane in the original feature space. But if you drew a new feature space that had an axis for X_2^2 , this would be a hyperplane.

So, we can add as many features as we want. We can actually just keep adding features until our classes are perfectly linearly separable, and then we wouldn't even need a budget C ! Although this is probably a bad idea from an overfitting perspective. We should probably keep C , but also add dimensions if we think we have non-linearity.

You could just add a lot of features and then directly try to fit a support vector classifier. But, you could quickly get overwhelmed by a huge number of features, and the computations would actually get really hard (I told you that smart people know how to do these optimization problems, but that doesn't mean they are trivial).

So, running with this idea, we will learn a little but more about optimization so that we can understand a really magical thing called the kernel trick.

11.5 Optimization

We are not covering optimization very much in this class. We are sprinkling in a few concepts here and there (such as gradient descent), but really optimization is a topic of an entire course! So I will not pretend to do justice to these ideas.

BUT, in order to understand the kernel trick, which is VERY important, we need to understand a little bit about optimization and how we actually solve for our support classifier.

We first note that the support vector optimization problem can actually be rewritten as:

$$\begin{aligned} & \underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{minimize}} && \frac{1}{2} \|\beta\|_2^2 + \lambda \sum_{i=1}^n \epsilon_i \\ & \text{subject to:} && y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) \geq (1 - \epsilon_i) \\ & && \text{and } \epsilon_i \geq 0. \end{aligned}$$

There are two differences between this and what we wrote before. First, recall that we were restricting ourselves to $\|\beta\|_2^2 = 1$ because any scalar multiple of a β vector gives us the same separating hyperplane. But ... that same scalar multiple also changes the meaning of a margin M . So, we might as well look for the SMALLEST $\|\beta\|_2^2$ that gives us a margin of 1. And it removes one thing for us to worry about - we don't actually need M . Second, we just moved the penalty on the size of the ϵ to the objective instead of setting it as a hard budget constraint: we already know from Ridge/Lasso that we can do this. We now have a penalty parameter (chosen via cross validation!) instead of a budget constraint.

Ok. Now that it is in this form, how do we solve this?

Well, remember Lagrange multipliers? Kind of? From Math 150/151, or from Econ 251, or from Amina/Bekah's colloquium, or from the Ridge/Lasso document that I put on GLOW? It's okay if you don't remember the details. But the idea is that we can solve a constrained optimization problem by first writing it in its Lagrangian form. This form introduces many free variables. In this case, the Lagrangian primal is:

$$L_P = \frac{1}{2} \|\beta\|_2^2 + \lambda \sum_{i=1}^n \epsilon_i - \sum_{i=1}^n \alpha_i (y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip}) - 1 + \epsilon_i) - \sum_{i=1}^n \mu_i \epsilon_i.$$

We solve this by taking derivatives with respect to β and ϵ and setting them to 0. These derivatives then set up big system of equations. There is theory from convex optimization that says that, instead of solving the Lagrangian directly, we can instead solve the dual function. Go read about this in the Boyd book (chapter 5) that I linked to the Ridge/Lasso notes! https://web.stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf.

In this case, the Lagrangian dual is:

$$L_D = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j.$$

Convex optimization theory says that we can solve our optimization problem by maximizing this with respect to α .

I still haven't told you how to solve this. Go read the Boyd book: Chapter 5! But the magical fact that you should notice is that this dual depends on the training data covariate vectors x_i and x_j only through the dot product $x_i^T x_j$, which is a scalar. For our n datapoints, we can store all of the values $x_i^T x_j$ in an $n \times n$ matrix: p , which is the dimension of our (possibly expanded) feature space, actually does not matter here, once we have written down the inner products.

11.6 The kernel trick

To make this more concrete: if we start with original features vectors x_i for all n individuals, and then transform them into a higher-dimensional space, we could denote this with $h(x_i)$. This might just be a function that takes a p -dimensional vector (x_1, \dots, x_p) and returns the $2p$ -dimensional vector $h(x) = (x_1, \dots, x_p, x_1^2, \dots, x_p^2)$.

The support vector optimization problem, in dual form, turns out to only depend on $h(x_i)^T h(x_j)$. Let's go one step further. I wrote this as a dot-product, but less expand this to any inner product, which we will denote $K(h(x_i), h(x_j))$. This is just a generalization of the dot-product to non-Euclidian spaces. My main reference is the wikipedia page!

What this means for us: we could choose a high-dimensional transformation $h(x)$, and if we happen to have a convenient way to write down $K(h(x_i), h(x_j))$ from a formula, we might need to never actually explicitly write down the high-dimensional vectors $h(x_i)$. We call $K(h(x_i), h(x_j))$ the kernel function. We can work directly with the $n \times n$ matrix of inner-products, and can actually solve for the hyperplane without even needing our high dimensions!

Common choices of the kernel function k are given on page 424 of ESL. The really cool thing is that the actual classifier can also just be written in term of the kernel function: we never need the p -dimensional β .

We can write our classification for a new input x as:

$$\hat{f}(x) = \sum_{i=1}^N \hat{\alpha}_i K(x, x_i) + \hat{\beta}_0.$$

All we need to do is figure out which vectors are the support vectors. These will have $\hat{\alpha}_i \neq 0$. Then, once we have this, we estimate an intercept, and write our classification rule! So cool! We predict class 1 if $\hat{f}(x) > 0$ and predict -1 otherwise.

Kernels have now been used in a lot of contexts beyond SVMs! You already know from polynomial regression and splines that expanding our feature space is nice. And this is a cool extension about when this is really efficient.

11.7 Reconciling Concept 1 and Concept 2

We started with the maximal margin classifier, which was linear and allowed no mistakes.

We then started allowing mistakes, using a budget C , in case our classes are not linearly separable. But then we enlarged our feature space a lot. In a high-dimensional enough feature space, all datasets become linearly separable. So ... do we throw out the idea of making mistakes?

No! See the figures on page 425 on ESL. Big warning: the penalty C used in ESL is the OPPOSITE of the budget/cost C used in ISL. How confusing! That's why I made our penalty λ in these notes. In the way that I wrote it in these notes for the kernel trick version. The penalty works as we would expect.

A big value of λ places a big penalty on mistakes. This means that the boundary will be VERY wiggly and will overfit the training points: so as to avoid any mistakes. This will also mean very few support vectors, so the boundary is totally determined by a small number of datapoints. A small value of λ encourages small $\|\beta\|_2^2$, which makes our boundary smoother.

11.8 Why are we learning about SVMs: do they work well and are they important?

Check out ISL 9.5, or ESL 12.3.4.

When SVMs came out in the 1990s, they were splashy and new and exciting. They worked really well in terms of accuracy, and seemed sort of mysterious. People actually thought that the kernel trick would help us avoid the curse

of dimensionality. Unfortunately(?), this isn't true. It turns out that SVMs are not so mysterious after all: they relate to all of our themes that we have seen in this course so far. And studying these connections is kind of cool and beautiful!

A few takeaways:

- Chapter 9.5 of ISL explains the close connection between SVMs and logistic regression. SVMs really do just have a loss+penalty form, which we are all used to.
- We still have a curse of dimensionality. If our true classification boundary depends linearly only on X_1 and X_2 , but we use a massive non-linear kernel, our “machine” will struggle to learn the correct rule. Bummer!
- Choosing the cost parameter C or the penalty λ is still SO important. We don't want to overfit! We need cross-validation, which sounds slow. Luckily, like for lasso, it turns out that we can solve efficiently for many values of λ at once.

Also, we should mention that we can use SVMs for multi-class classification and regression: we just only went over binary regression. Second, note that the whole idea of a plane means that the X s cannot be categorical: we need to represent categorical variables as dummy variables, and for dummy variables polynomial transformations don't help us :(.

Some overall takeaways: with a big kernel, SVMs have low bias. But we need to control the bias/variance tradeoff with our penalty parameter. They are pretty computationally efficient, and only a little bit interpretable. If we have few support vectors, we can sort of interpret them which is nice. If we have a lot of support vectors and a big kernel, they are basically a black box. So, don't choose them for their interpretability. For use-ability, we just need to choose a penalty parameter and kernel, which is not so bad.

I almost took SVMs out of this class! Because, now that neural networks exist, IDK if people actually use SVMs. But, seeing a big variety of different types of algorithms is nice. To help you draw connections. Also, seeing a variety of different motivations for algorithms might help you understand how new algorithms are developed! Which is important!

Also: the name “support vector machine” sounds so scary and you might encounter it someday! It's not actually scary: now you know it is a penalized hyperplane classifier!

12 Monday, March 17: Catch-up

Today, we returned to the notes from tree day, since we didn't finish them on tree day.

13 Thursday, March 20: Midterm

In-class midterm, followed by very long spring break.

14 Monday, April 7: Model Validation and Selection

14.1 Some comments on the midterms

14.1.1 Review of common mistakes from in class midterm

- **Problem 2:** Need to mention that deep learning is still subject to the bias-variance tradeoff! For a bit it seemed like double descent was going to challenge this, but recent work has shown that double descent is still governed by the same principles!
- **Problem 3:** Recall the difference between a model and a model-fitting-procedure. A linear regression with 5 covariates is a very simple model. We don't really need to worry about overfitting if we fit linear regression with 5 pre-specified covariates. But, the procedure of “search through 16,000 possible covariates and use stepwise regression to select the best 5” is a very complex model-fitting-procedure. We definitely need to be worried about overfitting in this case! You all saw something really similar to this on a HW problem!
 - When we talk about bias / variance / overfitting: we are usually talking about a model-fitting-procedure: something that takes in a random training set and outputs a model. On the other hand, when we talk about interpretability or inference (coming soon!), we often want to talk about a specific model.
 - When we talk about predictive accuracy: which do we care about? This is actually a tricky question, which does not relate to the in-class midterm but does relate to the idea of a “model” vs. “model fitting procedure”.

- * If we are getting ready to “deploy” a specific model to be used in the real world: we probably want to know the predictive accuracy of this specific model: trained on our entire training set!
- * Unfortunately, this specific accuracy is hard to estimate! We can use cross-validation on the training set to estimate predictive accuracy. But ... this CV uses a bunch of different models fit to different subsets of the training data– not our full final model! Tricky!
- * Looking at the predictive accuracy of our fixed model on a large test set would be nice. But ... if we had such a large test, why wouldn’t we use some of this data as training data to improve our model??
- **Problem 4:** In an image classification, the individual pixels are not meaningful! The patterns between the pixels are meaningful. So variable selection and additive models fundamentally don’t make much sense.
 - Consider identical copies of the same image, but in one image everything has been shifted one pixel to the left. Any sort of model that is additive in the original features will be really bad at recognizing these two images as similar!
 - One implication (relevant to what is coming, not the midterm): we need to think about what we mean by “interpretable” for an image classification problem. Consider a lasso model that selects 10 pixels to be used for classification. The model is simple enough for us to see what coefficients go with which pixels. But, is it actually useful for us to know the coefficient of pixel (4,17)? How do we make sense of that?
- **Problem 5:** Honestly, this went pretty well. There weren’t specific rows that everyone got wrong or anything.
 - One commonly forgotten common principle: if we have a tuning parameter like λ in a ridge penalty or k in KNN that directly controls model complexity, we know exactly what changing it will do to training error. Training error is lower for more complex models: it monotonically decreases as we move along the complexity axis! On the other hand, for test error, this tuning parameter has some “ideal” value that depends on the true data generating mechanism! The test error curve is U-shaped! So we often don’t know if it will go up or down as we move along our complexity axis!
- **Problem 6:** Check derivative algebra! Also, stochastic gradient descent processes training data in batches, so can be used for online learning. So, we might use stochastic gradient descent instead of the closed form solution for ridge regression if we have so much training data that we don’t even want to store it in memory, or if our training data is coming to us over time.

14.1.2 From the takehome midterm

- **Problem 1:** If a proof tells you that something has expected value c , and then you simulate to verify your proof, as you increase the number of repetitions, you should be able to get a result that is equal to c with a margin of error of 0.0000001, and even this should shrink with the number of iterations.
 - Sometimes the proof will also rely on the sample size n being big or something- and then you will also need to make this big in your simulation. But sometimes it does not!
 - The empirical average difference between training error and test error should have converged to exactly $2/n\sigma^2$ as the number of iterations increased. This did not depend on n , p , or σ^2 .
 - The key thing here that was tricky: the p in the expression is the number of coefficients in your regression model. So, unless you specifically told `lm()` to remove the intercept, then your effective p is actually $p + 1$!
 - If you happened to pick a big value of n , then the value of $2/n\sigma^2$ barely changes when you change p to $p + 1$. So this makes this hard to spot. But, for a lot of you, I think the difference was large enough that you should have been able to spot this difference.
- **Problem 2:** I wanted people to think carefully about exactly what they were seeing, and go beyond what they read in textbook! What the heck was elasticnet doing with those correlated variables? Was it “good” or “bad”?
- **Problem 3:** I was happy overall! And, this is basically our topic for today! So, people would do even better after today!

14.2 Back to regularly scheduled programming: review!

Before spring break, we learned about so many different machine learning algorithms. We also learned about so many axis on which to compare these algorithms. Let’s briefly review these, since it has been a while. This table (while really full) is not at all exhaustive! It’s just to remind you of algorithms that we have seen, and some basic properties. You should be able to make a table like this yourself but fill in your own thoughts and ideas!

	When?	Bias?	Variance?	Usability?	Interpretability?	Efficiency?
Linear regression (no bells and whistles)	Regression	Unbiased if true model is linear; otherwise likely biased (too simple)	Likely low because simple. Not low if p is really big.	Easy!	Easy!	Easy!
Linear regression with engineered features (splines, polynomials, etc).	Regression, worried linear isn't good enough.	Unbiased if you pick good features; otherwise likely biased.	Variance increases if you make engineered features really complicated (splines). Decreases if you are reducing dimension or redundancy (PCA).	Pretty easy, but you need to pick your features (how many PCs, degree of spline, etc).	Not amazing but not a black box.	Pretty good!
Stepwise regression.	Regression, p is really big, want to select subset.	Unbiased if true model linear or you added nicely engineered features; biased otherwise. Fewer steps = potentially more bias (miss out on important variables).	Grows if p is bigger or number of steps is bigger. So much greedy searching!	Pretty easy! Need to decide step criteria and stop criteria.	Even better!	Can seem slow, but faster than "best subset" because greedy approximation.
Lasso Regression	Regression, p is really big, want to select subset.	Bias grows with λ , and might just be big if true model not linear.	Variance shrinks with λ . Never THAT bad if p not too big.	I think easy! Just use CV to tune λ .	Pretty good!	Pretty good!
Ridge regression	Regression, p is really big; want to reduce variance.	Same as lasso.	Same as lasso.	Same as lasso.	Less interpretable than lasso.	Even better than lasso! Closed form!
Logistic regression; with engineered feature, stepwise, lasso, or ridge extensions	Classification (mostly binary, but has multi-class extensions)	Fill in same as all the linear regression rows. (but now no-bias means that log-odds are linear in the included features). (when classes are well-separated, hard to estimate log-odds in the "in-between" area. high variance.)				
KNN	Regression OR classification, when you think you have non-linearity, and when p is not too big.	Small if k is small.	Big if k is small.	Easy! Just pick k . Not going to work well in high dimensions, and then will need to consider PCA, etc.	Not a black box. "Prototype" interpretation.	Can be bad in terms of time and space without clever tricks.
LDA	Classification, when you think your classes are linearly separable, or, equivalently, Gaussian with shared variance.	Biased if not linearly separable	Pretty low variance.	Easy, I don't even think there is a tuning parameter.	Pretty good!	Easy!
QDA	Classification when you think your classes are Gaussian with non-equal variance	(Same as LDA basically, just slightly higher variance.)				
GAM	Regression. For classification, use logistic GAM.	Low bias because we let our model be really wiggly. But misses out on interactions!	Potentially high variance if we don't regularize or if we have a lot of features	Not bad, but there are a lot of moving pieces we could tweak.	Could be worse. Additive lets us make nice variable-by-variable plots.	Iterative back-fitting required; not bad.
Neural Network	Classification or regression; good for complex tasks!	SUPER LOW. Even gets all the interactions	Really high without proper regularization.	A lot of things to fiddle with; hard to use.	Black-box.	Slow; use stochastic gradient descent.
Trees	Classification or regression.	Bias shrinks as tree depth increases.	Variance grows as tree depth increases.	Super easy, although you do need to decide on a stopping criteria, and choice can affect results	Small trees are really easy for humans to interpret. But beware instability!	Fast!
Support vector machines	Mainly classification! We did not cover the regression setting.	If you use a "high dimensional" kernel and big budget, really low bias!!	Variance increases as you make higher dimensional or increase budget.	I think pretty easy to use out of the box.	Not that interpretable; I guess you can look at the weights if you are using a linear kernel	A bit slower (you might have noticed this in R!).

14.3 No free lunch

Why is it necessary to introduce so many different statistical learning approaches, rather than just a single best method? There is *no free lunch*⁷ in statistics: no one method dominates all others over all possible data sets. On a particular data set, one specific method may work best, but some other method may work better on a different data set. Hence, it is an important task to decide for any given set of data which method produces the best results. Selecting the best approach can be one of the most challenging parts of performing statistical learning in practice.

14.4 How we pick the best model depends on our goals!

In the Breiman “two cultures” reading that you will all complete for Monday, Breiman claims that there are two possible goals of statistical modeling:

1. **Prediction:** Make accurate predictions for new data.
2. **Information:** Describe or understand something about the universe is generating the data.

This probably is not a surprise to any of you- this is a common way to talk about two goals of predictive modeling. In the Efron reading that you will complete for Monday, Efron expands this classical list of two goals to a list with three possible goals.

1. **Prediction:** Make accurate predictions for new data.
2. **Estimation:** Learn the parameters of a model that can be written down (e.g. a regression model).
3. **Attribution:** Formally determine which variables are important.

I like that Efron distinguishes between estimation and attribution. However, it is not clear exactly how Efron’s three goals map to Breiman’s two goals. To gain valuable information, do we need to have parameters that we are estimating? Can we gain information about the universe without doing formal attribution? Also, while I don’t think it is a surprise to any of you that we will pick a different modeling approach depending on our ultimate goal, do you think that these different goals are at odds with one another? Or do you think that they can be complimentary? I look forward to hearing what you all think when I read your reading responses next week!

For the rest of class today, let’s assume that our primary goal is prediction. And let’s go into a bit more depth about how to select the best model when our goal is predictive accuracy.

14.5 A bit more about picking models based on predictive accuracy

First of all, it is important to note that the *no free lunch* idea applies even if our goal is only predictive accuracy!

Consider the extremely naive model that ignores the training data and always predicts $\hat{y} = 3$. Well, if we happen to live in a universe where $Y = 3 + \epsilon$, and none of the covariates X are important, then this model is optimal in terms of its mean squared error on new data! And the model $\hat{y} = \bar{y}_{train}$ will also do quite well. Since, we never know what the true universe data generating mechanism is, we can never say for sure that this “intercept-only” model $\hat{y} = \bar{y}_{train}$ is worse than deep learning!

You all already know that we can’t just compare models based on their training error. If we compare models based on their training, then we will always prefer the most complex models with the highest degrees of freedom. In the example above, deep learning will certainly have lower training error than $\hat{y} = 3$, because it will fit to the noise in the data. Instead, you all already know that we want to assess which models will perform the best on new, unseen data.

Let’s talk a bit more formally about how we assess expected predictive accuracy on new data. We have already touched on many of these methods/concepts, but I want to close the loop on ESL Chapter 7 because it is such an important chapter!

First, we need to mention that there are two separate things that we might be doing when we try to assess the accuracy of a model on new data.

- **Model Selection:** We want to compare several models or several model-fitting-procedures to decide which one to use as our final model.
- **Model Evaluation:** After choosing a final model, we want to estimate its final accuracy on new, unseen data.

Most data analysis pipelines involve model selection. Thus, we usually actually want to do both of these tasks. And, when we are doing both of these tasks, we need to be careful about how we use our data!

⁷This is the name of an actual theorem!

14.5.1 Train / test / validate

The gold standard way to carry out a data analysis pipeline that involves both selection and evaluation is the **train, validate, test** paradigm. The idea is simple.

- We split our data into three datasets: train, validate, and test.
 - To compare a bunch of models, we can train them on the training set, and then we can compute their error on the validation set. We can select the model that seems to perform best on the validation set.
 - The validation-set error of the selected model is a biased estimate of the true error of this model: we picked it specifically because it was a minimum– it is biased downwards!
 - So, we end by evaluating the performance of our selected model on the test set.

However, sometimes we don't have enough data to want to split it into three parts! So, there are some other options.

14.5.2 Fancy equations to estimate test error from training error

One option! If we can compute or estimate the **degrees of freedom or optimism** of a model, then we know that the expected test error is related to the training error via a simple formula. Thus, we can use an equation to figure out a **bound** on the expected test error, and then select the model that seems to have a low expected test error. This eliminates the need for a validation set. This is the reason why we can reasonably select a model based on AIC, BIC, or Mallow's CP. A limitation of these methods is that they relate only to the in-sample (fixed X) prediction error: they say nothing about what might happen at new values of X .

There are a couple of other ways of measuring model complexity besides degrees of freedom or optimism. One is VC dimension! Have any of you heard of VC dimension? This gives you another way to bound the difference between training error and test error. It is a pretty loose bound and we are not going to cover this!

One reason that we don't always use this option: degrees of freedom or VC dimension are very hard to compute for complex procedures!!! So ... we can't actually use these bounds!

14.5.3 Cross Validation

Cross validation re-uses data cleverly, and so can be more efficient than a train/test/validate split. But it is also a huge can of worms that is more complicated than it seems.

First, note that, if we are doing both model selection and estimate predictive accuracy, we really should still separate your data into a training set and a test set. We can then use cross validation on our training set to select a model. We can then re-fit the model with the selected tuning parameter on the entire dataset, and can then evaluate this model on the test set. You have all done this on your HW, basically every time that you have used `glmnet`: you use `cv.glmnet` to pick λ , but you still then evaluated the prediction error on a test set.

One reason that this is important: the CV error of a model that is selected because it minimizes the CV error is biased downward! Because it is a selected minimum! This is a very important concept.

If we are not doing any selection, and are just doing assessment, then we can get away with doing CV on our entire dataset. But, it is still conceptually tricky to figure out exactly what cross validation error is measuring! Since we fit a different model to each fold, and also each of these models only gets to use $(K - 1)/K \times n$ of the data points for training, it is not exactly the same as evaluating the prediction error of a fixed model fit to n observations. When we are just comparing different procedures to select a tuning parameter, we might not care about this. But when we are trying to evaluate a model, we do care!

Figure 7.8 of ESL, which is copied below as Figure 9 shows why the number of folds could have a big impact on what cross validation is actually measuring. Cross-validation as an estimate of expected prediction error has its own bias-variance tradeoff, which depends on the number of folds used. While there is no exact rule for how to decide how many folds to use for cross validation, a good rule of thumb is 5-10.

You need to be really careful about making sure that model selection and model assessment are both done properly using held out data. Recall your HW problem on the wrong way to do cross validation!! And see ESL 7.10.2. The idea is that we need to use CV to evaluate an entire model-fitting procedure. We are not allowed to do any part of the procedure (i.e. variable selection) using the entire dataset– if we do, then the test set isn't truly "held out".

I think that this concludes our study of ESL Chapter 7, and I think you are all experts now at selecting models with good predictive accuracy! Next week, and on your HW, you will think about other things that we might want to consider when selecting a model!




FIGURE 7.8. Hypothetical learning curve for a classifier on a given task: a plot of $1 - \text{Err}$ versus the size of the training set N . With a dataset of 200 observations, 5-fold cross-validation would use training sets of size 160, which would behave much like the full set. However, with a dataset of 50 observations fivefold cross-validation would use training sets of size 40, and this would result in a considerable overestimate of prediction error.

Figure 9: Figure 7.8 from ESL!

15 Thursday, April 10: Bagging and Random Forests

15.1 Model selection recap

On Monday, we talked about how to select a model if the main thing we care about is predictive accuracy. Roughly speaking, if your main goal is predictive accuracy, then you should be selecting the model with the lowest error on a test set.

- You could use a single train/test split. This lets you compare specific models fit to your specific training set.
- You could use something like 5-fold cross-validation. This makes more efficient use of the data, but you are no longer looking at the error of a specific model fit to the specific training set. You are looking at an average error that results from a model-fitting procedure applied to 80% of your training set.

There is a lot more that I could say about model selection or about cross-validation. But, at a basic level, you all know how to do this already. You all did it on your take-home midterm, for example. The one thing I want to emphasize that I mentioned briefly last time is the idea of the **winner's curse**.

- Suppose that, in reality, there are three models that are equally good in terms of their expected prediction error over all possible realizations of new test sets.
- On the single test set that we observed, one of these models will “win” due to random chance (i.e. it will have the lowest error on our observed test set).
- This test error is biased downwards for the expected test error of this model on a NEW test set. Selecting a winner caused us to overfit to our test set.
- For an unbiased estimate of a selected model's performance on new, unseen data, we actually need three datasets: train / validation / test. We can select the model that achieves minimum error on the validation set, and only at the end do we evaluate the accuracy on the totally unseen test set.

The idea of the winner's curse is really important! Selected models tend to let us down in the future :(. Selected policies in government, econ, etc. tend to underperform in the future: they “won” the first stage of selection partially due to random chance, and now they regress to their mean. This is important to be aware of, any time you are doing model selection!

That's all about model selection for now! If you love model selection, consider a related topic for your final project!

15.2 Intro to Ensemble Methods

Today, let's talk about a totally new and potentially crazy idea: if we only care about predictive accuracy, why do we need to select one model? Why can't we take a bunch of different models that all do pretty well, and average or combine their predictions somehow?

- If the different models tend to make mistakes on slightly different types of datapoints, or if their mistakes are uncorrelated with one another, this could really help us!
- We need to have room to improve; if we are already achieving essentially irreducible error, then this will not

help.

In other words: forget model selection. Let's use an *ensemble model*, that combines the results of several models!

You might be thinking that you are going to fit a KNN, a regression tree, and a lasso regression all to the same dataset, and then average the predictions from all of these different models. We could certainly do that! But, for now, let's not go too crazy with ensemble methods. Let's talk about two relatively simple methods for taking a single model fitting procedure and improving it by fitting it many times in a row.

- The general idea of *bagging* is to take one model that has high variance, and reduce the variance by averaging over several repeated copies of this model. We will talk about this today!
- The general idea of *boosting* is to take one model that has high bias, and reduce the bias by iteratively focusing attention on the examples that we are currently messing up on. We will talk about this next time!

15.3 Bagging (Breiman, 1996)

Suppose we have a single base learner, such as a single deep decision tree, that we could train to our entire dataset. Call this model $\hat{f}()$. There is a fair amount that we like about a deep decision tree (it can uncover interaction terms, it has low bias). However, there is a thing that we do not like about this decision tree. Due to the high variance of decision trees, we know that if we had observed a very slightly different training set, we would have a totally different tree. This high variance compromises our predictive accuracy!

Well, here is a really simple idea that works well if we care about predictions. We can take fit B different deep decision trees to B different bootstrap samples of our dataset. Remember that, to get a bootstrap sample, we draw n observations with replacement from our n original training observations. We now have B different prediction models, $\hat{f}_b()$ for $b = 1, \dots, B$: each fit to n observations. We let our final, bagged model (bagging stands for bootstrap aggregation), be:

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x).$$

And that's it! That is the entire idea of bagging!

15.4 Why does bagging reduce variance?

We know in statistics that averaging is a nice way of reducing variance. But also, bagging shouldn't be some magical way to make our model really good: taking B bootstrap samples from our data doesn't actually increase the total amount of information in our data! So, what is going on and why does bagging work?

Let $Var(\hat{f}_{bag}(x))$ be the variance of a prediction for a single fixed test-point x . The randomness is taken over the training of $\hat{f}_{bag}()$ for many different training sets. Let $Var(\hat{f}(x)) = \sigma^2$. What is $Var(\hat{f}_{bag}(x))$?

Well, based on the form of $\hat{f}_{bag}(x)$, we have that:

$$Var_{bag}(\hat{f}(x)) = Var\left(\frac{1}{B} \sum_{b=1}^B \hat{f}_b(x)\right) = \frac{1}{B^2} \left(\sum_{b=1}^B Var(\hat{f}_b(x)) + \sum_{b=1}^B \sum_{b' \neq b}^B Cov(\hat{f}_b(x), \hat{f}_{b'}(x)) \right).$$

Note that, $\hat{f}_b(x)$ and $\hat{f}_{b'}(x)$ are identically distributed random variables: they both result from applying the same procedure to a random subsample of the same data. Let $Var(\hat{f}_b(x)) = \sigma_b^2$ for any $b = 1, \dots, B$. Also, note that when n is large, σ_b^2 should be the same as σ^2 : fitting a model to the whole training set should be the same as fitting a model to a bootstrap sample from the whole dataset. Next, note that $\hat{f}_b(x)$ and $\hat{f}_{b'}(x)$ are not independent, because they are trained on overlapping subsamples of data. Let $Cov(\hat{f}_b(x), \hat{f}_{b'}(x)) = \rho\sigma_b^2$, where ρ is the correlation induced by the overlapping data. Note that this value is the same for all b, b' from 1 to B .

So,

$$Var(\hat{f}_{bag}(x)) = \frac{1}{B^2} (B\sigma_b^2 + B \times (B-1)\rho\sigma_b^2) = \frac{1}{B^2} (B\sigma_b^2(1-\rho) + B^2\rho\sigma_b^2) = \frac{\sigma_b^2(1-\rho)}{B} + \rho\sigma_b^2.$$

Suppose that we are working with a very stable model fitting procedure. In this case, there probably isn't much change from bootstrap sample to bootstrap sample. This means that there is a lot of correlation between $\hat{f}_b(x)$ and $\hat{f}_{b'}(x)$; $\rho \approx 1$. In this case:

$$Var(\hat{f}_{bag}(x)) \approx \frac{\sigma_b^2(1-\rho)}{B} + \rho\sigma_b^2 = \sigma_b^2 \approx \sigma^2.$$

Uh oh! This means that bagging did not help us reduce our variance! We did a lot of computational work to get the

same variance that we would have had if we had fit just one model to our full dataset. The issue here is the stability—when we get basically the same model from every bootstrap sample ($\rho \approx 1$), bagging really does not help us!

On the other hand, suppose we have an unstable model, and ρ is positive but less than 1⁸. In this case, especially if B is large, the variance of our bagged model is smaller than the variance of an individual bootstrap-sample model: $\hat{f}_b(x)$, which is similar in large samples to the variance of our single model applied to the whole dataset: $\hat{f}(x)$. The smaller that ρ is, the better!

The fact that bagging reduces variance more for unstable models led Breiman, in his original bagging paper, to claim that bagging does not help with linear regression or ridge regression (stable) but helps with stepwise regression or trees (greedy and unstable). This is good to know!

15.5 Does bagging reduce bias?

The average of B linear models is still a linear model. Thus, bagging a bunch of linear regressions does not do anything to the expressiveness of our model, and does not impact the bias. On the other hand, consider bagging shallow decision trees (1 split). Each individual tree is so limited in its complexity; but when we combine many together we get a much more complex additive model. One one-level decision tree gets to use only one variable, whereas a bagged ensemble can take into account many variables. In general, the average of B trees cannot necessarily be written as a tree: so we have increased our model space, and therefore may have reduced bias.

15.6 Random forests (Breiman, 2001)

Before 2001, bagged regression trees were already very popular. Trees are a popular learning algorithm; people like that they can uncover interactions without needing to pre-specify them. However, trees are also notoriously unstable. This makes them a prime candidate to benefit from bagging.

However, in 2001, Breiman combined some existing ideas to make an algorithm that (often) works even better than a bagged ensemble of trees. The idea is simple: we can make the variance reduction of bagging even better by further de-correlating the trees in our ensemble of bagged trees? Or, in other words, by making the individual trees even less stable?

In random forests, we do exactly this. We add additional randomness to every individual tree in an ensemble of bagged trees.

- Recall that, to build a CART regression tree, we greedily select the variable and split point that most improves the goodness of fit of our tree right now. We repeat this process recursively; we keep making new splits by picking the best possible variable.
- For each individual CART tree in a random forest, we make a small modification. At each split, we only allow ourselves to consider a random selection of $m \leq p$ variables as the possible split variables. Thus, we may not choose the overall best split variable right now; we might need to choose something else if the “best” variable is not randomly selected.
- This makes the individual trees in the forest more different from one another (smaller ρ). While it may seem from the derivation above that this *definitely* helps reduce the variance of the bagged ensemble, we need to remember that this also makes $\text{Var}(\hat{f}_b(x)) > \text{Var}(\hat{f}(x))$: an individual tree in the forest has extra noise added compared to a single tree fit to all of the data with all of the variables. However, the gain that comes from the de-correlation tends to help.
- This can also help us with bias. Recall that, due to the greedy nature of a tree, a tree could miss out on potentially important variables or interactions if the greedy process leads it to go down a certain initial path (dominated by one important variable) while missing other paths. Randomization helps the tree explore all paths and find all possible signal or interactions. This can be really helpful— it seems to work really well on real data.
- Of course, selecting a small value of m could also be bad for bias. If there are only a few truly important predictors and we select a small value of m , for many trees in our forest these important predictors will be left out. This is too bad— it will cause us to miss out on true signal!
- In the case of many correlated predictors, a random forests helps us explore different choices of correlated predictors in different trees. This can be helpful.

⁸Negative ρ would be quite strange, given that the bootstrap samples overlap significantly. I don’t think that this would happen.

15.7 More details about random forests

15.7.1 Tuning?

In a random forest, we need to choose:

- How many trees (B) should go in the forest?
- How deep should each tree in the forest be?
- How many variables m should be considered for each split?

Conventional wisdom says that it is okay to build deep trees (high variance!) inside of the forest, because we do not care about interpretability and because bagging will help us reduce the variance. Conventional wisdom also says that you might as well make B big if your computational power allows you to. Increasing B “too much” wastes computation, but cannot cause us to overfit, because it does not cause the model to become more complex, it just causes us to converge to an average. We can keep adding trees to the forest until our OOB error (see below) stops improving.

The hardest variable to tune is m . While some “rule of thumb” out there says to select $m = \sqrt{p}$, there are definitely situations out there for which $m = p$ (no extra randomization) is optimal. The RandomForest package in R has a built-in mechanism for tuning m : it seems really important to check this for your dataset.

15.7.2 OOB error evaluation

To tune a random forest, we don’t actually need to do cross-validation. This is because there is already some sample-splitting going on during tree-building!

Each tree in the forest is built to a bootstrap sample of the data. This bootstrap sample tends to have around 2/3 of the observations “in the bag” and around 1/3 of the observations “out of the bag”. We can compute the MSE for this tree for these “out of the bag” data points! More generally, for any data-point, we can get a current “test error” by getting its prediction from the forest using only the trees for which it was out of the bag. This is nice- this MSE is not affected by overfitting.

We can keep adding trees to our forest until the OOB error stops improving. This will be when we have more trees than necessary: since the OOB MSE doesn’t get to use as many trees per datapoint to get predictions as we would use for new data. But it’s still a useful metric that can tell us when we clearly don’t need more trees.

15.7.3 Permutation importance

While random forests provide a clear way to improve on the predictive accuracy of a single tree, they also lose one of the main benefits of trees: interpretability. However, as we discussed, should we really call an unstable model (a single tree) interpretable? If you would like to use a random forest but would also like to know about which variables are important to your model, people have developed a variety of simple *explainability* techniques that try to get at this. One of them is permutation importance, which is built into the random forest R package. We will cover this next Thursday, on interpretability/explainability day!

15.8 A discussion: let’s go back to no-free-lunch

Varya was working on an R Demo for you all that was supposed to show that a random forest outperformed bagging! However, with either simulated data or fairly simple real datasets, she was having trouble showing this result! Often, bagging (use all features) was still beating random forests! After looking through her code and trying to debug it (there were no bugs), I came to some conclusions.

- Just because your textbook says that random forests should outperform bagging doesn’t mean they will on simple datasets: there is no free lunch! No algorithm outperforms all other algorithms on all datasets!
- For the randomization idea to be useful, there needs to be “room for improvement”. For example, if your data has tricky-to-find effects with lots of important variables, then the extra randomization might really diversify the trees in your forest and really help. Real datasets tend to be super complex, which is likely why people claim random forests work so well in the real world.
- When your data is not super complex, you probably want $m = p$. Setting $m < p$ makes the bias and variance worse for every individual tree: so unless you are really gaining something from this when you average, you probably don’t want to do it!
- In practice, since there is no free lunch, it is always important to tune your models! Nothing is REALLY “off the shelf”.

- Don't just blindly trust what a textbook says! Everything depends on your data!

16 Monday, April 14: Boosting

Last time, we defined *ensemble methods*. We then talked more specifically about bagging, which is a general-purpose technique that is designed to take a single model with low bias but high variance (e.g. deep decision tree) and reduce the variance through averaging. Today, we will talk about boosting, which is a general-purpose technique that is designed to take a single “weak learner” with low variance but high bias and adaptively improve it until it becomes a “strong learner”.

A cool thing about boosting is that there were really key innovations from both computer scientists and statisticians at around the same time, and the innovations got combined to make something really powerful! This perhaps shows that the “two cultures” discussed by Breiman aren't really at odds with one another: they compliment one another!

16.1 AdaBoost (Freund and Schapire, 1996)

AdaBoost is an important early boosting algorithm! It is not covered in ISL, but is covered in depth in ESL!

16.1.1 Algorithm

The ideas for boosting were first born in theoretical computer science, and started with classification. We have n datapoints, and our responses y are either -1 or 1 : these represent two classes.

The idea was to wonder: is it possible to take a weak learner or a base learner (any algorithm that performs slightly better than random guessing) and “boost” it until it becomes an arbitrarily accurate “strong” learner? It turns out that it is possible! The idea involves fitting the weak learner T times during rounds $t = 1, \dots, T$. At each round, we re-weight the data points to focus on the ones that we misclassified previously. The pseudo-code is as follows.

- Start with training set $(x_1, y_1), \dots, (x_n, y_n)$, where all weights are $w_1(i) = 1/n$ for $i = 1, \dots, n$.
- For $t = 1, \dots, T$:
 1. Take a random sample of size n , with replacement, from the training dataset, where datapoint i is probability of being selected is its current weight w_i .⁹ If our weak learner allows the use of “weights”, we can also skip this step and go right to step 2, but run the algorithm on the weighted data.
 2. Train the weak learner on the subsampled or weighted dataset. This yields $\hat{f}_t(x)$. Compute its training error, e_t , on the weighted or subsampled dataset.
 3. Let $\alpha_t = \frac{1}{2} \log \left(\frac{1-e_t}{e_t} \right)$.
 4. Update the weights! Let:

$$\tilde{w}_{t+1}(i) = w_t(i) \exp \left(-\alpha_t y_i \hat{f}_t(x_i) \right)$$

and then let $w_{t+1}(1), \dots, w_{t+1}(n)$ be the $\tilde{w}_{t+1}(i)$ values but normalized to sum to 1 so that we have a probability distribution over our training examples.

- The final learner is:

$$\hat{f}_{boost}(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t \hat{f}_t(x) \right).$$

That is it! That is AdaBoost!

16.1.2 Intuition

We need to talk about exactly what is going on in each step. Even though boosting is now much more general than AdaBoost, AdaBoost is such an important early example that I think it is worth focusing on some intuition for AdaBoost before we go further.

- First let's discuss α_t . This is the weight that each weak learner $\hat{f}_t(x)$ will end up getting towards the final prediction. This is based on the training error e_t of each weak learner. The only thing we assume about e_t is that our weak learner can do better than random guessing on our training set: thus, e_t is between 0 and 0.5 for a binary classification task. On this domain, the weight α_t is large and positive if the training error is near 0,

⁹In round 1, this is just a bootstrap sample. But this will change later!

and is near 0 if the training error is near 0.5. So ... the classifiers that do well on their own training sets get more say in the final prediction! This seems reasonable!

- Now let's talk about the weight updates w_t . Recall that we are assuming we wrote down and stored y as -1 and 1 . And we are assuming that the weak learner $\hat{f}_t(x)$ outputs either -1 or 1 for each datapoint. So, the term $y_i \hat{f}_t(x_i)$ in the weight update term will be positive whenever we correctly classified this point, and will be negative otherwise.
 - So, correctly classified points get $\tilde{w}_{t+1} = w_t(i) \exp(-\alpha_t)$.
 - * If this weak learner was doing a good job overall, then α_t is big and $\exp(-\alpha_t)$ is near 0. So, in this case, correctly classified points get small weights in the next round! We don't need to focus on them again!
 - * On the other hand, if this weak learner was doing a bad job overall, then α_t is near 0 and $\tilde{w}_{t+1} \approx w_t$: we will essentially try again with current weights to see if we can improve.
 - Incorrectly classified points get re-weighted as $\tilde{w}_{t+1} = w_t(i) \exp(\alpha_t)$.
 - * If this weak learner was doing a good job overall, then α_t is big and we will increase the weight of this point a LOT. The idea is like "wow, even for an overall really good learner, we missed this point. Now we need to go focus on it."
 - * If this weak learner was doing a bad job overall, then α_t is near 0 and $\tilde{w}_{t+1} \approx w_t$. So we never make large updates based on a weak learner that isn't even doing well!
- At the end, each weak learner gets to contribute a $+1$ or -1 vote to the classification. Its vote matters more if it had a big α_t .

16.1.3 Practical considerations, or cool properties?

- By design, AdaBoost definitely continually reduces training error. If each $\hat{f}_t(x)$ is a bit better than random guessing on the training set, the more of these that we combine the better that we do. Our training error should actually go to 0 if T is big enough. This property holds for binary classification where we start with anything better than random guessing: but might not be practically super interesting because this says nothing about test error.
- People were at first worried that AdaBoost would certainly overfit if T is too big. Remember that on midterm 3 you derived a formula that said that the difference between the training error and the expected test error depends on the degrees of freedom of an algorithm. Working with a similar concept, people showed that an upper bound on the difference between the training error and the expected test error for AdaBoost grows with T . However, in practice, it seems like we don't hit this upper bound. In practice, AdaBoost can work pretty well, even when T is so big that the training error is 0.
- What is going on?!?! When we have a training error of 0, aren't we overfitting? Why would we keep increasing T past this point?
- Even after AdaBoost has reached a point of 0 training error, we can keep iterating and keep re-fitting with new weights. The consequence of this tends to be that we increase our margin of confidence for our classification: much like an SVM! We know from SVMs that increasing the margin can be good for generalization error; we get more confident in our predictions. So ... that's cool! Whether or not extra iterations will increase the margin or overfit to noise depends on the signal to noise ratio in your data.
- Freund and Schapire certainly note that, while AdaBoost can work really well off-the-shelf with some nice properties, it is clearly dependent on the data and the weak learner. It can fail, and is particularly susceptible to noise.

A note on statistics vs. computer science culture. As you know from Breiman, statistics literature often starts by assuming a data-generating mechanism for the data. On the other hand, a lot of computer science machine learning literature starts from a "no-noise" setting, where we are really doing like pattern-recognition, not statistics. Boosting (like SVMs) was first studied by people in the no-noise setting! In a no-noise setting, over-fitting is less of an issue. This difference colored some early comparisons of bagging vs. boosting. In the Dietterich paper that I posted on GLOW, you can see that he finds that bagged trees outperform boosted trees in high-noise settings.

16.1.4 Comparison to Bagging

Now that we understand a bit more, we can compare AdaBoost to bagging!

- The idea of boosting is to start with a weak learner and improve it sequentially. The focus is on reducing bias by focusing on hard examples.

- The idea of bagging is to start with a learner that has high variance and reduce the variance through averaging.
- Both lose interpretability compared to a single model, but can have much higher accuracy.
- Both are supposedly general purpose and off the shelf, but you might still need to do some tweaking or tuning!
- Bagging can be trained in parallel; boosting cannot be because it is adaptive. For huge data applications, we might care about this!
- Bagging might work better in high-noise situations. And in statistics, we always think there is noise!
- Both have some beautiful theory attached that is beyond the scope of this class!

16.2 Boosting is an additive model? (Friedman, Hastie, and Tibshirani, 2000)

Boosting is far more general than just AdaBoost! It can be extended to multi-class classification and regression. In the process of extending it, people realized it has beautiful connections to other concepts in statistics!

We don't have time to do justice to all of these connections! We will really briefly talk about the ideas from one paper: Friedman, Hastie, and Tibshirani (2000). They connect boosting to additive logistic regression. These insights helped with computational efficiency of boosting, let us study its theoretical properties with a wider array of tools, and also led to further extensions! This paper is summarized really nicely in ESL Chapter 10! The notes here are not intended to provide a lot of detail: just to provide connections to things you have seen already in class!

The key insight of this paper is that AdaBoost can be seen as an iterative forward stage-wise algorithm for fitting an additive logistic regression model, that optimizes a particular loss function. This insights helps them see that AdaBoost can be both simplified but also greatly generalized!

16.2.1 Remember GAMs?

Recall from many weeks ago that a GAM (generalized additive model) has the form:

$$\hat{f}(x) = \sum_{j=1}^p \hat{f}_j(x_j).$$

This could be a linear regression model where we are directly modeling y , or could be a logistic regression model where $\hat{f}(x)$ is modeling the log odds of belonging to a certain class. This model does not allow predictors x_j and x_k to interact, unless we pre-specify the interaction as its own feature, which was a limitation.

Recall that we had an iterative procedure (backfitting) for fitting a GAM that allowed each individual $\hat{f}_j(x_j)$ to be something super complicated like a smoothing spline. The backfitting algorithm starts with a guess for each $\hat{f}_j()$. Then, for one j at a time, it fits a model to predict the current residuals $y - \sum_{k \neq j} \hat{f}_k(x_j)$ using x_j .¹⁰ It uses this model as its new, updated guess for $\hat{f}_j()$. We do this iteratively until the guesses for the $\hat{f}()$ stop changing.

Why are we mentioning GAMs? The idea of iteratively fitting relatively simple models to current residuals to improve the model sounds like boosting!!! There are just two differences.

- This model goes predictor-by-predictor; boosting did not have this limitation.
- This model goes back and updates the guess for $\hat{f}_k(x_j)$ multiple times; boosting does not change a simple model once it fits it once; it just adds more simple models.

16.2.2 GAMs with new basis functions

What happens if we make our additive model a little bit less additive?

Let $b(x, \gamma)$ be a basis function, which depends on parameters γ , that depends on all of the predictors x . We could let:

$$\hat{f}(x) = \sum_{m=1}^M \beta_m b(x, \gamma_m). \quad (18)$$

This model is additive in the basis functions, but in terms of the individual features this is a linear combination of non-linear functions of the individual features.

We have seen a model just like this before! A neural network with one hidden layer can be written in this way!¹¹ The γ_m are the weights in the first layer that connect each x to hidden node m . Then, b is the activation function, and β_m

¹⁰If we are doing logistic regression, we don't quite use residuals. We use working residuals, which are more complicated

¹¹We also talked about projection pursuit regression: an extension of GAMs to more complicated models, that also looks just like this!

is the weight in the second layer that connects hidden node m to the output. So, if b is a non-linear transformation of a linear combination of the x s, then this is a neural network, and we could fit it with backpropagation.

But what if b is a simple regression tree, and γ_m encodes the split variables and split points of this tree? All of the sudden, (1) I am not sure how to fit it with backpropagation, but (2) this looks quite a bit like a model we could fit with AdaBoost!

16.2.3 Adaboost as an algorithm for fitting (18)

Suppose that we fit (18) in a greedy forward stagewise manner. The idea of a forward stagewise procedure is that we first figure out β_1 and γ_1 , then we figure out β_2 and γ_2 , etc. Unlike with backfitting, we do not let ourselves go backwards and update β_1 again later in the process. By “fit” the model, I mean pick the values for β_m and γ_m that most improve a certain loss function between $\hat{f}(x)$ and y right now.

It turns out that the weights and update steps for Adaboost are exactly equivalent to what we would get if we fit the model (18) using a greedy forward stagewise procedure with a particular exponential loss function. The reason that this perspective is so cool is that it immediately tells us that we could come up with different boosting procedures by replacing the exponential loss function with another loss function!

For classification, we could consider a loss function that is more robust to outliers. For regression, we can use something like squared error loss, which makes boosting *very* simple (see the algorithm in ISL where they simply fit regression trees to residuals!).

16.2.4 Adaboost as gradient descent?

Forward stage-wise boosting is a very greedy strategy. We want to take the step that most reduces the loss function right now. Recall that gradient descent updates initial guesses for parameters by moving them in the direction that will most reduce the loss function right now. These seem really similar!

In fact, in ESL Section 10.10.2, the authors discuss the fact that, when we add a new tree to a boosting model, we really want to add a tree that moves in the direction of the current gradient of the loss function. Thus, we should fit our new tree to the current gradient of our loss function: this way, we are moving in a direction that is close to the direction of the negative gradient, but we are moving in a way that is allowed in our model class (we are just adding a tree).

For squared error loss, this just means fitting a regression tree to the current residuals! For other loss functions, we similarly fit a tree to a form of current working residuals. Thus, we can really see boosting as fitting (18) by always taking steps in our model space (i.e. steps that add a new tree) in the direction of the gradient of the loss function.

There have been some other innovations that have come from viewing boosting as gradient descent or an additive model. For example, we know that in gradient descent we should take SMALL steps in the direction of the gradient. We shouldn't be too greedy right now, lest we overshoot our optimum or get stuck in a local optimum. This leads to the idea of a shrinkage parameter for boosting; maybe we shouldn't let ourselves take steps that are too big! So maybe we should only take a SMALL step in the direction of each new tree. We might need more trees in the model, but this seems to really help prevent overfitting.

This was a whirlwind: but now you know that gradient boosting exists!

17 Thursday, April 17: Explainability vs. Interpretability!

We are going to spend some time today covering the boosting content that we missed. We are also going to have a short discussion of the “two cultures” reading. Thus, we won’t have time to cover too much new content today.

For the last few classes, we have thrown interpretation to the wind. We have talked about selecting models based on predictive accuracy, and we have talked about ensembles, which are certainly hard to interpret. We have been imagining for the last few classes that our primary goal is predictive accuracy.

Today, let’s assume that our primary goal is to learn something about the world. We don’t need to make predictions for a new, unseen dataset. We need to understand the data that is in front of us! There are actually two ways that we could accomplish this goal.

- We can fit an **interpretable** model, otherwise known as a *glass box*. Examples include simple or sparse linear models, or small decision trees. With these models, we know exactly what variables are playing a role in our predictions, which presumably helps us understand the world.
- We can fit a *black-box* model, but then **explain** this model after the fact using a separate algorithm. This is the huge and growing field known as explainable AI. The only example we have mentioned so far of an explainability technique is permutation importance for random forests (this was also in your Breiman reading). But we haven’t yet done this topic justice!

One theme of today is that this dichotomy is kind of useful, but that it’s also really tricky to define these things. What is interpretable depends on your audience! And what you are trying to convey!

Today, I want to just briefly touch on possible explainability techniques for determining important variables. They are both very simple, and pretty general-purpose with a lot of room for tweaks. There are many versions of these out there: I am presenting really simple versions. We will use this just to have a language for discussing interpretability vs. explainability!

17.1 Permutation importance (no retraining)

The idea is to figure out what variables are most contributing to your predictions in a given black-box model.

- Fit a model $\hat{f}()$ to predict y using a training set. Measure the loss on a test set: i.e. compare $\hat{f}(x^{test})$ and y^{test} .
- For variables $j = 1, \dots, p$:
 - Randomly permute the values of variable j in the test set to get \tilde{x}^{test} .
 - Measure the loss between $\hat{f}(\tilde{x}^{test})$ and y^{test} .
 - The increase in loss from permuting the j th variable is the importance of the j th variable.

This idea is really simple, and is effective in telling us which variables are contributing to our predictions. But, it has a lot of problems if we try to interpret it as a measure of the association between variable j and y .

Imagine that x_j and x_k are extremely highly correlated, and both associated with y . But imagine that our model decides to only use x_k for its predictions. Well, then according to permutation importance, x_j is not important at all. But that doesn’t mean it is not associated with y —it is just being masked by x_k .

This gets at a particular definition of importance: a variable is not important if we can make good predictions without this variable. The thing that it does not take into account is the fact that maybe, without this covariate, we would have fit a different model! This is something that the next method tries to address.

Note that glass-box models have this issue too though: this just means that attribution is really hard. A linear regression cannot always tell which of two correlated predictors is important! So, this isn’t a unique failing of permutation importance; just something to point out!

17.2 Leave one out importance (with retraining)

This one is actually even more strange! The idea is subtly different, and this is much more computationally expensive than the idea above.

- Fit a model $\hat{f}()$ to predict y using a training set. Measure the loss on a test set: i.e. compare $\hat{f}(x^{test})$ and y^{test} .
- For variables $j = 1, \dots, p$:
 - Leave variable j out of the training set to get x_{-j}^{train} .
 - Refit the model on the training set. Now you have $\hat{f}_{-j}()$.
 - Measure the loss between $\hat{f}_{-j}(x_{-j}^{test})$ and y^{test} .
 - The increase in loss from removing the j th variable is the importance of the j th variable.

In some ways, this feels more thorough to me. But in the case where x_j and x_k are extremely highly correlated, and both associated with y , I think we get even worse results. In the previous example, at least x_k would have popped up as important. Here, when j is left out, k can absorb its effect. But when k is left out, j can absorb its effect. So we might end up reporting neither of these as important! This is not desirable!

This again is something that can happen in interpretable models too! If we include x_j and x_k both in a regression in this setting, the variance inflation factors that come from their correlation could lead us to report neither as significant.

17.3 Model distillation (student/teacher methods)

The idea of model distillation is to first train a black box that achieves high classification accuracy. This is the “teacher” model. The next idea is to take the predicted values from this model, which may be probabilities instead of hard classifications, and try to predict these using the covariates using a simple, glass-box like a tree. The idea is that maybe a tree would not originally have done very well on this data. But we have now distilled the response variable into pieces that can be explained by the covariates, and we are giving the tree hints about which observations we are more or less confident about. So now all of the sudden, the tree can do better at explaining these predictions than it could have done at predicting the original classes. Huh! Cool! I would like to learn more about this some day!

17.4 Other explainability methods?

There are many other methods that Lucy or our guest speaker next week might tell us about! And I will try to post some additional readings! This could probably be an entire class, so it was hard to know what to focus on!

17.5 Do we want interpretability? Or is explainability enough?

17.5.1 Some arguments for interpretability

A quote that I love from the Sankaran reading posted on GLOW is that an interpretable model supplements human decision-making, rather than replacing it. And this is how we want to be using our models! Cynthia Rudin, in reading that you will do for HW7, implies that explanations are not always true to what the model is actually doing, and can get us into trouble. If we really truly want transparency, then interpretable models might be the way to go.

17.5.2 Some arguments for explainability

Breiman strongly makes the point that models that are simple enough to be interpretable are often less accurate. He says that this is bad, because we only want to learn about the world from a model that is accurate enough to explain the world. I actually have some sympathy for this viewpoint: can we really learn something trustworthy about the world from a model with really low predictive accuracy? I am not sure! The economists would say yes!

A point that I take more seriously is the point about stability. We love trees because they are supposedly interpretable. But, we also know that a tiny perturbation to a training set could cause a tree to look totally different and have totally different variables at the topmost splits. Should we really be interpreting something so unstable? With this in mind, fitting many trees to bagged samples and recording which variables get selected over and over again suddenly seems a lot more appealing as a measure of variable importance.

Overall, we like that we can see into glass boxes. But, if the glass boxes don’t work well or are not stable, why do we care? I think these are reasonable points!

17.5.3 More coming!

The reading you do for HW7, as well as the guest lecture next Thursday, should help reinforce these concepts!

18 Monday, April 21: Inference

18.1 What do we mean by inference?

In Stat 202, I give my students the following two definitions.

- *Probability*: given an exact probability distribution, including the value of all of the parameters, figure out how a random variable will behave (probabilities of events, moments, etc).

- *Inference*: given one or more realizations of a random variable, figure out what probability distribution it might have come from. If you already assumed that the random variable(s) are from a certain family of distributions or are related in some way, this means estimating values for certain parameters that mean something to you.

If these were the definitions that we were sticking with, then it would seem like it only makes sense to talk about *inference* in Breiman’s “data modeling” world, and not in his “algorithmic” world. Under the definition above, if we don’t think that there is

However, I also tend to think of *inference* as any method for quantifying uncertainty about something that we are estimating from our data. And the thing that we are estimating from data could just be a single prediction! So, maybe it does make sense to do inference even when we are being “algorithmic”. More on this below!

18.2 Reminder from Stat 346: Confidence Intervals vs. Prediction Intervals

In Stat 346, you probably used the words “confidence interval” and “prediction interval” to refer to different quantities. I think that these names are confusing, but reviewing what they mean in a simple linear regression setup will help us understand different types of inference that we might do.

Consider a setup where we are from the data modeling culture. We believe that

$$Y = \beta_0 + \beta_1 X + \epsilon, \quad \epsilon \stackrel{i.i.d.}{\sim} N(0, \sigma^2).$$

We use our training data to come up with estimates $\hat{\beta}_0$ and $\hat{\beta}_1$. These of course also give us estimates $\hat{y}(x)$ for any new x that we might plug in. We also have an estimate s_e of the unknown σ .

Consider the following three types of intervals.

- A confidence interval for the parameter β_1 , which we form using $\hat{\beta}_1 \pm t^* SE(\hat{\beta}_1) = \hat{\beta}_1 \pm t^* \frac{s_e}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2}}$. This second equality just plugs in the form of $SE(\hat{\beta}_1)$, which depends on the estimated amount of noise in y as well as the spread of x .
- A confidence interval for a prediction $\hat{y}(x)$, which we form using $\hat{\beta}_0 + \hat{\beta}_1 x \pm t^* \left(s_e \sqrt{\frac{1}{n} + \frac{x - \bar{x}}{\sum_{i=1}^n (x_i - \bar{x})^2}} \right)$. This is just $\hat{y}(x)$ plus the estimated standard error of $\hat{y}(x)$, which depends on the same terms as $SE(\hat{\beta}_1)$, but additionally depends on how far this particular point is from the middle of the distribution of the training set x . The idea is that, for 95% of all random samples, this interval should contain the true value of $\beta_0 + \beta_1 x$ at this point x . Therefore, this is really just quantifying uncertainty in the parameters β_0 and β_1 ; plus a little transformation that is specific to where you are along the X-axis. This transformation reflects the fact that small changes in estimated slope are more impactful on your prediction when you are far from the center of the data.
- A prediction interval for a prediction $\hat{y}(x)$, which we form using $\hat{\beta}_0 + \hat{\beta}_1 x \pm t^* \left(s_e \sqrt{1 + \frac{1}{n} + \frac{x - \bar{x}}{\sum_{i=1}^n (x_i - \bar{x})^2}} \right)$. Under our assumptions, this interval is supposed to contain the true value of y when $x = x$ 95% of the time. Note that all we did in this simple setting was add an extra factor of 1, which gets multiplied by s_e . So we are just allowing for variability around our regression line, which has size s_e . The variability around the regression line for a single point doesn’t shrink as our sample size increases: only the uncertainty in the line itself does.

The point in reminding you about the difference between a confidence interval and a prediction interval is that a confidence interval often only makes sense if we are coming from a data-modeling culture and are fitting a parametric model. But a prediction interval clearly extends well beyond the parametric setting! Any time that we are making point-predictions $\hat{y} = \hat{f}(x)$, we might wonder if we can turn this into an interval prediction that will contain the true y for this x with high probability.

I still think that it makes sense to call the third bullet point a confidence interval- it’s just a confidence interval for an individual outcome! So you might sometimes hear the words confidence interval or prediction interval used interchangeably- I don’t think that this is a huge problem. But others might care more about precise language!

18.3 Why do we care about inference?

Suppose that your main goal in fitting a statistical learning model is to make predictions for new, unseen observations. In this case, being able to quantify uncertainty about future predictions is incredibly important! So prediction intervals are vital.

Separately, suppose that our main goal is to understand some underlying process or truth about the real world. In this case, we want to do what Efron calls attribution.

- If we are estimating parameters in a linear model to quantify which variables are associated with the response, you all already know why it is important to do inference. We want to make sure that we aren’t seeing a non-zero

effect of a certain variable just due to random chance in our training sample. We want to feel really sure that, in the population, the effect is non-0. So, in a parametric model, it makes sense to do inference to quantify our uncertainty about the underlying population parameter. This typically means doing inference on a pre-specified parameter.

- Suppose we are fitting a regression tree, and our regression tree selects a split on a certain variable and a certain split point. We might want to know that this split is not specific to our training data- we might want to know that this split exists in the population! So, we might want to do inference here, and the inference is not a prediction interval, it is a confidence interval. But, it is different than the linear model above! Because ... the parameter we end up doing inference on is not pre-specified- it is something that was selected by our regression tree. As it turns out, this can get us into all sorts of trouble.

You all know how to make confidence intervals and prediction intervals for linear regression. And, with appropriate assumptions, confidence intervals and prediction intervals work in pretty much the same way for any parametric model. This was really the topic of Stat 360. In this class, I want to focus on what happens when we are in more non-parametric settings. Can we still do meaningful inference?

I divided the sections below into first covering an example of a confidence interval for a parameter that summarizes a dataset vs. an example of a prediction interval for a single observation.

18.4 Confidence intervals for algorithmic models

In order to make a confidence interval that we *wouldn't* call a prediction interval, we do need *some* notion of a population parameter: even if we did not start out with a parametric data-generating model.

To give an example of how we might end up with one of these, suppose that you fit a single-level regression tree to your data. Out of p total predictors, the regression tree decides to make a split on $X_j < s$ for some $j \in 1, \dots, p$ and some constant s . The model now predicts \bar{y}_L for all points on the left side of the split and \bar{y}_R for all points on the right side of the split, where these \bar{y} values were computed on the training set.

Before you built your regression tree, you didn't really have a data-generating model in mind. But, now that you are looking at your tree, you do sort of have one in mind. The model says that:

$$E[Y] = \begin{cases} \mu_L & \text{if } s_j < s \\ \mu_R & \text{otherwise} \end{cases}.$$

In this model, \bar{y}_L and \bar{y}_R are just sample estimates of population parameters; we would like to use them to do inference on the population parameters μ_L and μ_R . We could also write this as a linear model

$$E[Y | X] = \mu_L + \beta 1_{x_j > s},$$

and could imagine that we are interested in doing inference on β , which parameterizes the difference in means across the split in our tree. In fact, if we fit a least squares regression of our training responses y on the indicator vector $1_{x_j > s}$, our sample intercept will be \bar{y}_L and our sample slope will be $\bar{y}_R - \bar{y}_L$. The reason I am bothering to mention this is so that you can see that, after we have fit a regression tree, we can refit it as a nice linear model. But the difference is that the parameters were not specified in advance: this indicator vector $1_{x_j > s}$ won out over all possible indicator vectors.

Thus, as we have alluded to many times in this class, downstream inference based on the training data is invalid. We cannot use our training data to make a confidence interval for μ_L or μ_R . I decided not to go through it in class today, but you can easily make a small simulation study for yourself in R where you show yourself that, if you do this, your confidence intervals will not have 95% coverage!

There are two options:

- Split your data into a training set and a test set. Fit a tree to the training set to come up with your parameters of interest. Do classical inference on these parameters using the test set. While this is very flexible and should work in most settings, it still requires the notion of a "parameter of interest" that is meaningful on the test set- so I'm not sure what it would look like in the case of a true black box. It makes a lot of sense for trees though!
- Come up with fancy math to do non-classical inference that accounts for the selection event. This is a cool area of research! But it tends to require a lot of assumptions and a relatively simple setup (not an arbitrary black box). You can do it for regression trees under a Gaussian assumption- this was my first PhD project!

Overall, I think that this tree example is a good one to keep in mind in your future endeavors in terms of what we might mean when we say we want to do inference for machine learning models! It can mean all sorts of things, but this is an example where it meant that we wrote down a parameter AFTER seeing our fitted machine learning model.

18.5 Prediction intervals for black-box models

As Efron hinted at in the reading that you all did, prediction is actually easier than estimation or attribution. It turns out that we *can* make finite-sample-valid prediction intervals for black-box models with very few distributional assumptions. An example of a framework for doing this is *conformal inference*. Conformal inference is extremely cool; see the RMarkdown demo on GLOW. I will migrate some notes from the RMarkdown demo into this document at a later date.

19 Thursday, April 24: Safety, Ethics, and Explainability

I have no notes to post for today because we have our guest speakers! I will post their slides on GLOW after their talks.

The main things I hope you got from this lecture are:

- There are many explainability techniques that can be used to understand the predictions of a black box model.
- It is possible to “trick” a black-box model with an adversarial attack. For example, Blake showed us the example where we can trick a self-driving car into thinking that a stop sign is actually a “40mph” sign by imperceptibly changing a few pictures. This type of vulnerability is really bad!
- We need to be able to probe models to look for possible vulnerabilities. This is red-teaming. Who should be responsible for this type of careful auditing for real-world models?

You should all let me know what else you learned from this guest lecture, and what else you think the key takeaways here.

20 Monday, April 28: No class

Sorry about the schedule change!

21 Thursday, May 1: Class discussion on ethics, explainability, and fairness

You all submitted great discussion questions and we had a great discussion! No formal lectures notes for this unit. Here are some things that we discussed that I hope you will all keep in mind in your future statistical careers.

- What can go-wrong when black-box prediction models are employed for high-stakes decision making? One example we saw was recidivism prediction (used to create sentencing decisions). You all also looked up your own examples, and found examples in hiring, college admissions, medical resource allocation, and facial recognition.
- Who should be held accountable when ethical issues arise in the use of prediction models? The creator of the model? The person using the model? The local government wherever it is being used? Outside academics who conduct audits of the models? We saw that a lot can go wrong when no one is specifically held responsible!
- Are their settings where it is “good” to use black-box prediction models because they are so accurate, etc? You all mostly think yes; just maybe not in these high-stakes settings where an individual person is affected and would like to be able to argue with their results, etc.
- How do we define fairness? (This could be a whole class).
- Is explainability enough, or do we need true interpretability? (Rudin thinks explainability is not enough).

Think about any other takeaways you have from your reading!

22 Monday, May 5: Introduction to Unsupervised Learning and Clustering (Day 1)

Announcements:

- **Final projects:** Speed presentations a week from Thursday, papers due over reading period.
- **HW8:** Clustering and cluster validation: posted this week, due next Friday. Short, because I want you to do a good job on your final projects!

22.1 What is unsupervised learning?

So far in this class, statistical learning has meant *supervised* learning. We have always had a response variable y (either numerical or categorical), and we have always been at least partially trying to predict y using explanatory variables x . I say *partially* because sometimes prediction has been our primary goal, whereas other times our primary goal has been understanding/interpretation; but we have still achieved understanding/interpretation by fitting a model $\hat{f}(x)$ such that $y \approx \hat{f}(x)$.

Today, we will start a brief unit on *unsupervised learning*. The key difference is that we have no special response variable y ! For your purposes, we have a large matrix of data $X \in \mathbb{R}^{n \times p}$ (assume that categorical variables are represented with indicators), and our goal is to *find structure* in this data.

I think that a large number of unsupervised tasks can be thought of in the following way. We know that our data $X \in \mathbb{R}^{n \times p}$ has some randomness or noise in it. But, we think that

$$E[X] \approx L\beta^T,$$

where $L \in \mathbb{R}^{n \times k}$ and $\beta \in \mathbb{R}^{k \times p}$, where k is often much smaller than p . We can think of L as latent variables; they are unobserved, but they help us understand the structure in our data. If this setup is not what you think of when you think of unsupervised learning, note that it encompasses the two types of unsupervised learning you have likely heard of, because:

- L could be a matrix of indicator-vectors that represent group membership (clustering!).
- L could be continuous, and could represent coordinates of our data along k new axes (principal components, dimension reduction, etc).

In some ways, clustering is the more mathematically difficult task, because it is non-smooth and uses fewer beautiful tools from linear algebra. However, we will start with clustering today because it is easy to visualize, so gives a first good sense of unsupervised learning.

Before we talk about specific clustering algorithms, let's talk about some contexts where unsupervised learning comes up or is used. Because I think it's important to keep these in mind: there are a lot of under-appreciated differences between different contexts where unsupervised learning comes up.

22.2 Where does unsupervised learning come up?

I am sure I am missing contexts; this is just a list to get you started.

Context 1: We really wish we were doing supervised learning, but our labels (y) are missing.

- Consider image classification; a classic task in supervised learning: we have a big set of pixel-level data for each image, and we want to use the pixel-level data to predict if the image is a cat, dog, horse, etc.
- Getting large training datasets from the internet for image classification tasks is really time intensive. Because the images are quick to download, but a human often needs to go through and add labels (cat, dog, horse) to every image.
- If we are processing a lot of data from online and we do not want to use human annotation to get labels, we suddenly have image classification as an unsupervised task.
- We might train an algorithm to cluster our images into groups. Only after we have our groups will we have a human go through and label the entire group: "this cluster represents the dog images".
- When people are developing clustering algorithms with this goal in mind, they often take a dataset that DOES have gold-standard labels, and evaluate the ability of a clustering algorithm to reconstruct these gold-standard labels. This gives us a notion of "accuracy", which we do not always have for unsupervised learning!

Context 2: We think that the p -dimensional variation in our data can be summarized by a smaller number of latent variables; we know in advance what these latent variables represent and how many there are.

- This is basically the same as Context 1, but there doesn't necessarily need to be a concrete prediction task in mind.
- In the analysis of single-cell RNA sequencing data, researchers often want to characterize variation by "cell type". But cell type is not observed in the sequencing process. Thus, they use clustering to come up with estimated cell types (L), and then proceed with further analysis.
- Now that researchers have published various "Cell Atlas's" detailing recognized human cell types, there is in theory an agreed upon set of cell types that we are looking for, and this is essentially a missing data problem rather than an amorphous exploratory task.

- Context 3:** We think that the p -dimensional variation in our data can be summarized by a smaller number of latent variables; we don't know in advance what these represent or how many there are.
- This is a lot like Context 2, but is much more exploratory.
 - The cell type example can sometimes turn into this, because researchers want to not only detect existing cell types but also detect novel cell subtypes; so the number or meaning of the clusters is not always fixed in advance.
 - I think psychologists do this a lot too, with factor analysis. Given a big set of results to, i.e. a personality survey, they try to explain the variation in their results with a smaller number of latent factors. They can then try to call these “personality types”.
 - An online retailer might look for batches of similar customers, to advertise similarly to them in the future. They don't know in advance what they are looking for, but they end up assigning names to the groups: “sports-enthusiasts”, “budget-hunters”, etc.
 - In image classification, we can learn cool latent concepts that summarize high-dimensional pixel data. Once we can visualize these, we might recognize them (i.e. beak means bird). But we don't know how many concepts we are looking for in advance; we will label them after.
- Context 4:** We want to visualize a complex dataset in low-dimensional space, organize our data, or reduce the dimension of our data for some other reason.
- Dimension reduction, a form of unsupervised learning, can just be really nice for visualizing high dimensional data in low dimensional space! If we can reduce the dimension without losing much (i.e. $X \approx L\beta^T$, we can save on the cost of storing our data! We might also believe that much of what we are losing is just noise, so we are de-noising the data.
- Context 5:** We have X and y , but we want to pre-process X with unsupervised learning
- You have actually all already done this!
 - PCA regression in Stat 346! Or, HW1 in this class, where you used PCA as pre-processing for KNN to help with the curse of dimensionality.
 - Dimension reduction can be a nice preprocessing step for supervised learning!
- Context 6:** Imputation for missing data.
- Suppose we have a matrix X of data with a lot of missing values. Maybe downstream we want to do supervised learning, or maybe not.
 - We need to fill in these missing values. If we believe that $X \approx L\beta^T$, and we have enough data to learn good values for L and β (we will talk about methods), then we can impute the missing X_{ij} with $(L\beta^T)_{ij}$. These imputed values respect the general latent trends in the X data.
 - If this is our goal, we can evaluate how well different methods work by taking non-missing data and randomly holding it out, and checking our ability to reconstruct it.
 - So this is another setting that has sneaky connections to supervised learning: there is a right answer, we just don't have the answer at test-time.
- Context 7:** Recommender systems.
- I no longer think that we will have time to cover this one in this class, but it is a good one to know about. Because it is fun!
 - Suppose that Netflix has data for n customers and p movies. Customers can rate a movie that they have watched, and say if they liked it. But, not all n customers have rated all p movies. So there is TONS of missing data in this matrix.
 - We can use dimension reduction to learn latent types of customers based on the rating that they have provided. We can then use imputation (matrix completion) to fill in hypothetical ratings for additional movies. We use this to make recommendations!
 - So cool! And a really straightforward but clever application of unsupervised learning!
 - If you tried to do movie recommendations with a fully supervised approach, you would run into trouble because you have so much missing-ness! So few customers have seen a particular movie!
 - But this matrix completion idea is so clever. They have nice algorithms that are really good at learning good latent factors even in the presence of a lot of missing-ness.

22.3 What is clustering?

A clustering algorithm takes a matrix $X \in \mathbb{R}^{n \times p}$ and splits the n observations into k discrete groups. The groups are supposed to be homogenous within-group and heterogenous between-group. Clustering algorithms differ along a few broad axes:

- Does the number of groups k need to be specified in advance?
- What metric do we use to define a homogenous group? Is it compactness in Euclidean space? Or connectivity along a neighborhood graph?
- Can an observation belong to more than one cluster? Can it have a certain probability of belonging to a few different clusters? This is the difference between hard-clustering and soft-thresholding.

The three clustering algorithms that we will see in this class will hit on these three different properties.

22.4 A very brief introduction to k-means

- K-means is a very popular and very simple clustering algorithm. Jack will present on the details!
- We define homogenous-clusters using within-cluster mean-squared error (euclidian distance between a point and its estimated cluster centroid).
- We need to pre-specify k , and we use “hard”-clustering (each point gets assigned to exactly 1 cluster).
- The basic idea of the algorithm is simple. We start by randomly choosing k cluster centers. Each datapoint then gets assigned to the cluster center it is closest to. We then update the centers to be the mean of all datapoints currently assigned to that cluster. We iterate until the centers/assignments stop changing.
- It turns out that, under the assumption that $X_i \sim N_p\left(\sum_{k=1}^K 1(i \in C_k)\mu_k, \sigma^2\right)$ K-means is searching for the maximum-likelihood estimates for the cluster assignments (each i must be in exactly one cluster C_k) and the cluster centers μ_k . It is searching using an EM algorithm; it is not guaranteed to find the maximum likelihood estimates, since it could get stuck in a local-maxima of the likelihood function. It is sometimes good to try out K-means with several different random seeds, and choose the “best” (according to within-cluster MSE) final cluster.
- Since K-means inherently assumes Gaussian clusters with shared variance, it creates linear boundaries between clusters. This is not always ideal!
- How do we select K ? A topic for next class!
- K-means is a lot like KNN; subject to curse of dimensionality; should scale our features; irrelevant noise features can hurt us, etc.

22.5 A very brief introduction to hierarchical clustering

- Hierarchical clustering is a good framework to be aware of. It is less of “one algorithm” and more of a way to think about clustering. Amina will present on the details!
- The idea of (agglomerative) hierarchical clustering is that we start by assuming every datapoint is its own unique cluster. We then go step by step and we merge datapoints that are most similar to one another. We ended up with a tree-structure for our data, in which points are connected in the tree to similar points. This creates a natural, hierarchical grouping of points, with main clusters and sub-clusters.
- We don’t need to decide the number of clusters in advance; we can decide by looking at our tree.
- We need a notion of a linkage function. How do we decide how close two clusters are to one another? Do we consider all points? Or only points that are closest to one another? Amina will tell us about a few choices of linkages.
- Hierarchical clustering actually starts from a distance matrix; not from raw data. Thus, we don’t need to use Euclidian distance. We could use something else if we wanted!

23 Thursday, May 8: Clustering Day 2: Spectral Clustering and the problem of Validation

23.1 What is Spectral Clustering, and why might we use it?

Figure 10 shows a classic picture that motivates spectral clustering. It is just a toy example: in real life, our data has more than 2 dimensions and it probably does not look like donuts! But it is still a helpful motivator.

The idea is that K-means and other clustering algorithms that use Euclidean distance as a similarity metric will do really poorly on the donut-shaped data. The donut-shaped data shows that we don't always want to prioritize compactness in defining clusters; sometimes we want to prioritize connectivity. This is where graph-based clustering methods can be really useful. Spectral clustering is an example of a graph-based clustering method.

The biggest and most important ideas behind spectral clustering are things that you could have come up with earlier in the semester. One big idea is that we can use distance-metrics besides Euclidean distance! Another big idea is that we can take data that doesn't initially look "nice," and make it look nice by expressing it in new coordinates (we did this with transformations, SVM kernels, etc.) A final big idea (that was maybe under-emphasized early in the semester) is that linear algebra is really cool and magical.

For spectral clustering, we can make a similarity matrix for our data points that calls two points similar if they are in each-other's K nearest neighbors. This might have Euclidean distance hiding underneath (if we use Euclidean distance to define our K nearest neighbors), but you can tell from looking at the right panel of Figure 10 that this similarity metric will do quite well on the donut-shaped data— it will know that all of the red points are "neighbors by transitive property," but none of the blue points are "connected" to any red points. This similarity matrix is $n \times n$ and symmetric: element i, j stores the similarity between object i and object j . It should be non-negative. We put 0s on the diagonal (element i is not marked as being similar to itself). This similarity matrix is S . There are all sorts of ways to define it: we don't need to use KNN. We could use KNN weighted by Euclidean distance, just Euclidean distance, etc.

Another cool thing is that we don't need to start with numerical data; we can start with something like text data! We can make a similarity matrix between words in our "dictionary" by counting the number of times that the words co-occur in a sentence or document. We can then apply spectral clustering to this similarity matrix: this is really widely used in practice!

This similarity matrix can be seen as the adjacency matrix of a graph representation of our dataset. Our new notion of clusters is cliques in this graph: either perfectly connected components, or else groups of nodes with a lot of edges (or a lot of highly weighted edges) between them that are connected to OTHER groups of nodes by very few edges.

Once we have our similarity matrix S , we make our degree matrix D , which just stores the row sums of S on its diagonal. It expresses overall connectedness of a point i with element i, i . Finally, we let $L = D - S$. This is the Laplacian matrix of our graph representation of our data.

Finally, we take an eigen-decomposition of L . We know that we have one eigenvector of 0, because the rows of L must sum to 0 by construction, and so the matrix does not have full rank (the vector of all 1s is in the null space). The multiplicity of 0 as an eigenvector is the number of perfectly connected components of the graph. The idea is that, to find clusters that are not perfectly connected components, we can run k-means on a few of eigenvectors of L corresponding to the small (but non-zero) eigenvalues.

Why did we do all of this work if we were just going to be applying K-means? Because we are now applying K-means to a transformed version of our data. This is like the kernel trick for SVMs! We took data that looked like donuts, and in these new coordinates (it's just a transformation!), it will look like something that K-means can cluster. This is so cool!

All we did was transform our data, but it turns out that small eigenvectors of the L matrix give an optimal transformation if our goal is to represent connected components as compact components! Linear algebra is so cool! Feel free to play around with the R demo on GLOW!

You might notice that there are still a LOT of choices involved in running spectral clustering. The final choice of K for K-means. The dimension of the spectral embedding (how many eigenvectors to use in the transformed version of the data). The distance metric for the initial similarity matrix. Etc. Since we have no ground truth, it is really hard to know how to pick these things!

So, despite having seen a few different types of clustering at this point, we are still left with one basic question: how do we know which clustering is best among a set of candidate options?

We need to be able to do model selection and model validation for clustering!!!

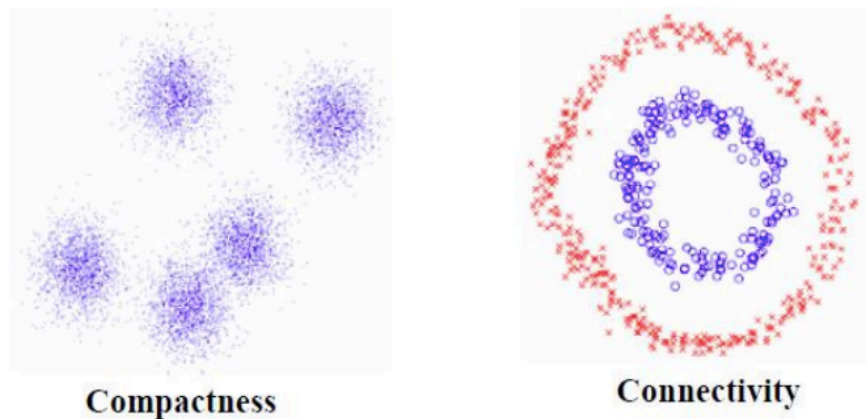


Figure 10: This figure shows two different notions of “good clusters” or “similarity” in two dimensions. On the left, we see 5 clusters that are “good” according to a classic notion of compactness: k-means will be great for identifying these five clusters. Euclidean distance in untransformed space captures the structure in the data! On the right, we see concentric circles. K-means will be terrible at identifying these two clusters, but a graph-based approach (like spectral clustering) with the right similarity metric will do great!

23.2 How can we do model validation for clustering?

Unfortunately, this is really hard. See the slides that I posted on GLOW.

Big picture, we could try to evaluate a clustering using something like within-cluster mean-squared error. This choice of loss function makes a lot more sense in the K-means world, where we are prioritizing compactness.

It should not surprise you all to know that within-cluster MSE will always decrease (on a training set) when we increase k , the number of clusters. This is because of overfitting: a model with 10 clusters can “fit” the data better than a model with 5 clusters; however, it might just be fitting noise, not signal.

In the supervised setting, we avoid issues of overfitting by splitting our data into training and test sets. Unfortunately, this does not work well for clustering! The problem is that, if we cluster a training set, we have nothing to evaluate on the test set! Clustering a training set does not yield an equation that can be easily applied to a test set to obtain cluster labels for a test set, and it does not provide cluster labels for the test set. This is the fundamental difficulty!

Because of this difficulty, in unsupervised learning people often do not strive for nice, U-shaped loss functions that are minimized at the best value of k . Instead, they look for elbow-shaped loss functions, where a visible “elbow” in the loss function curve at k clusters tells us that, past k , we are likely fitting noise and not signal.

This “elbow method” is widely used, but leaves much to be desired! An active area of work over the past 20 years has been trying to formalize the elbow method or else come up with alternate ways to do cross validation that work for unsupervised learning.

24 Monday, May 12: Dimension reduction and recommender systems

We have used principle components or the SVD a few times this semester, but I haven’t really told you very many details about them! Today, we will study the singular value decomposition of a matrix in a little more detail. We will then touch on some really cool applications of the SVD and other matrix decompositions.

The theme for today: unsupervised learning is more than just clustering! There are other task that involve finding structure in unlabeled data!

24.1 What is the SVD?

We can decompose any $X \in \mathbb{R}^{n \times p}$ into

$$X = UDV^T.$$

Without loss of generality, assume that $n \times p$. Then U is an $n \times p$ matrix where the columns are orthogonal unit vectors; D is a diagonal $p \times p$ matrix; V is a $p \times p$ orthonormal matrix. The diagonal elements of D are nonnegative,

and $D_{11} \geq D_{22} \geq D_{33} \dots$. These values in D are called the singular values of X . Because of the properties above, $U^T U = I_p$, $V^T V = I_p$, $V V^T = I_p$, but $U U^T \neq I_n$.

24.2 What are the uses of the SVD?

For the purposes of today, here is the cool thing about the SVD. If you want to approximate X with a rank- k alternative, where $k \ll p$, then the best possible approximation (in terms of SSE) is: $U_k D_k V_k^T$, where U_k represents the first k columns of U , D_k represents the top $k \times k$ block of d , and V_k is the first k columns of V .

We interpret the columns of V as our new axes in a new, transformed space where we need less information to express X . The coordinates of the data points along these new axes are given in UD ; they represent the projection of X into the axes defined by V . It turns out that this is an optimal projection/transformation: $U_k D_k V_k^T$ is as close as we can get to representing X if we only want to store k components instead of p .

This tells us that, if we want to do dimension-reduction and we care about how close our reduced-dimension version of our data is to X , we should do the SVD!

Why would we want to do dimension-reduction?

- As preprocessing for supervised learning (on the predictors), to avoid the curse of dimensionality or redundant predictors.
- As preprocessing for clustering, to avoid the curse of dimensionality or redundant predictors.
- To de-noise our data, or store a cheaper version of our data (i.e. image compression). This is a truly unsupervised task.
- To uncover latent factors in our data. We can interpret the columns of V as defining new variables: these might be interpretable for us! This is a true unsupervised / “finding structure” task.

For any of these applications: deciding k , the latent dimension, is JUST as hard as picking the number of clusters. We have no good cross-validation-ish solution! Because it is unsupervised! We can make things like elbow plots. You did this on your HW!

These are applications that I think you have encountered before. Today, we will at least partially cover matrix completion and missing data, which is an extension that I don't think you have seen before!

24.3 How does this relate to principal components and eigenvalues?

If you take your data matrix X and you center and scale it so that each column has mean 0 and standard deviation 1, then the SVD gives exactly PCA. The columns of V are the loading vectors, and the columns of UD are the score vectors.

If $X = U D V^T$, then $X^T X = V D^2 V^T$. $X^T X$ is a symmetric, square matrix. It turns out that $V D^2 V^T$ is exactly the spectral decomposition of the matrix $X^T X$: The diagonal elements of D^2 are the eigenvalues of $X^T X$ and the columns of V are the eigenvectors of $X^T X$.

This perhaps helps you relate the SVD to previous things that you learned about PCA. If X is centered and scaled, then $X^T X$ is the correlation matrix of X . The leading eigenvectors of the correlation matrix tell you the directions of maximum variation in X ; these are the top principal components. Today, you saw that you can obtain these same directions using the SVD.

I think it's generally good to know that eigenvectors are directions of maximum or minimum variance (depending on if they go with the smallest or the largest eigenvalues), and that singular vectors are the best ways to reconstruct a dataset faithfully in fewer dimensions. I did not retain this after Math 250 at Williams, but I learned it eventually!

24.4 Recommender Systems

This is the application that I want to focus on today, because I think it is unlikely that you have seen it in a previous course, and I think that it is a beautiful application of linear algebra to an important problem.

Suppose that I work for Netflix, and you are a customer of Netflix. You have a viewing/rating history on Netflix, and I also have viewing/rating histories for tons of other users. I want to predict if you will enjoy movie j .

This sounds like a supervised task. I want to make a prediction! But there is a problem with the availability of data. Suppose that I would like to use the predictions from past customers who rated movie j to predict your rating for movie j . Well, I can only include customers who actually watched/rated this particular movie. I also need to know which of these customers is “similar” to you, so I need you to have watched and rated similar previous movies to these

customers. If I try to approach this in a totally supervised manner, I am going to need to really cut down on my dataset size. And ... I will need to do this for every movie and every user separately? This sounds terrible!

While this task sounds fundamentally supervised (predict a movie rating), it turns out to have a really nice solution via unsupervised learning.

Let X be a matrix with n rows for n total Netflix customers and p columns for p total movies on Netflix. Let $X_{ij} = r_{ij}$, a rating, if customer i rated movie j . Otherwise, let the value X_{ij} be missing.

Our goal is to fill in the missing entries in the matrix with our best guesses that respect the structure of the data. We can then use the imputed value for i, j to be the predicted rating for user i and movie j : and use this to decide if we should recommend movie j or not. This is known as matrix completion!

We do this under the assumption that X (if it had no missing values) would be well-approximated with a *low rank* structure, where $X \approx U_k D_k V_k^T$. The latent axes (V) are latent genres of movies (the column space of X) and the coordinates in U define latent cliques of users (the row space of X).

If X had no missing values, we would directly apply the *SVD* to find our latent genres and cliques. Which could be a cool exercise in “understanding the structure in our data,” latent variable modeling, etc. So this is a general way in which the SVD could be interesting! But ... here we are in a very different situation, because probably 99% of the matrix X is missing values (in the case of the Netflix data). Standard computational techniques for computing the SVD cannot be applied!!

The general idea is that we want to solve the optimization problem:

$$\text{minimize}_Z \sum_{i,j \text{ observed}} (X_{ij} - Z_{ij})^2 \quad \text{subject to: } \text{rank}(Z) = k.$$

Unfortunately, this is really computationally hard to solve! But people have come up with nice approximations, algorithms, etc.

ISL, Algorithm 12.1, covers “Hard Impute”, which I think is clever and makes sense and is an awesome demonstration of how learning fundamentals like the SVD can really get you a long way. However, in practice, more complicated algorithms certainly exist too.

You might enjoy reading about the Netflix Prize

(<https://www.thrillist.com/entertainment/nation/the-netflix-prize>) of the early 2010s, which is when this problem and this field really took off!

25 Thursday, May 15: Final presentations

The end! No more lecture notes! Hopefully these lecture notes will improve a lot between now and the next time I teach this class (summer project!).